

## ACTIVITY PACK

# Activity 4: Security Checklist

Activity 4 - Rogue Agent Post-Incident Review

|             |                                      |
|-------------|--------------------------------------|
| Course      | AI Security for Autonomous AI Agents |
| Course code | TGS-2025060473                       |
| Version     | 2.1   20 August 2026                 |

Tertiary Infotech Pte Ltd | UEN 20120096W

## Security Checklist

### Activity 4 - Autonomous Agent Containment Checklist

- ☐ Evaluation targets are synthetic or explicitly authorised and isolated from public infrastructure.
- ☐ Internet access is denied by default and verified from inside the sandbox.
- ☐ Agent, evaluator and target identities are distinct, short-lived and least-privileged.
- ☐ Skills, plugins, MCP servers, tools and dependencies are pinned and allowlisted.
- ☐ The agent cannot discover, install or update capabilities autonomously.
- ☐ Filesystem mounts are minimal and read-only unless the test requires a bounded write path.
- ☐ Outbound network destinations, protocols, data volume and DNS are controlled and logged.
- ☐ Destructive, publishing, identity, credential and cross-boundary actions require human approval.
- ☐ Approval binds the exact command, parameters, target and time window.
- ☐ Rate, cost, step and time budgets stop runaway loops.
- ☐ Behavioural monitoring detects sequences across accounts, repositories and tools.
- ☐ A tested kill switch stops execution and revokes credentials without relying on the agent.
- ☐ Checkpoints support rollback; evidence preserves prompts, plans, tool calls and effects.
- ☐ The incident plan covers third-party notification, legal/privacy assessment and safe disclosure.