

ACTIVITY PACK

Activity 5: Readme

Activity 5 - Agent Governance and the Deployment Gate

Course	AI Security for Autonomous AI Agents
Course code	TGS-2025060473
Version	3.0 20 August 2026

Tertiary Infotech Pte Ltd | UEN 20120096W

Readme

Activity 5 - Agent Governance and the Deployment Gate

Course: AI Security for Autonomous AI Agents **Course Code:** TGS-2025060473 | **TSC:** ICT-INT-0052-1.1 **Block:** Day 2
- Autonomous Agent Security

At a glance

Duration	25 minutes
Format	Small groups of 3 - 4, plenary debrief
Type	Capstone synthetic simulation (SIM)
Learning Outcome	LO3
K/A statements assessed	A1, A2

Purpose

Capstone. Learners act as an AI governance board deciding go/no-go on an autonomous banking agent, reading evaluation data for bias and hallucination, applying PDPA/PDPC and the IMDA agentic framework, and setting autonomy limits and accountability.

Files in this folder

File	Audience	Use
SCENARIO.md	Learner	The case study narrative, artefacts and data
DISCUSSION-QUESTIONS.md	Learner	The five questions to work through in groups
SECURITY-CHECKLIST.md	Learner	Organisational, autonomy, Responsible AI, PDPA and operations deployment gate
Activity-5-*.pdf	Learner	Printable full learner pack including the deployment checklist

How to run it

1. **Set up (5 min).** Form groups of 3 - 4. Issue SCENARIO.md and DISCUSSION-QUESTIONS.md. Trainer debrief notes are supplied separately and are not stored in the learner-facing Activities folder.
2. **Group work (10 minutes).** Groups work the five questions in order. Circulate; the questions are designed to be argued, not looked up.
3. **Presentation (5 minutes).** Each group presents its answer to one nominated question.
4. **Debrief (5 minutes).** Use the separate trainer notes to draw out the teaching points and correct recurring misconceptions.

Learner workflow

1. Reproduce the aggregate and segment calculations from the SIM data and state what the aggregate hides.
2. Map the purpose, notification, protection, retention, transfer and breach-response obligations for each personal-data surface.
3. Classify all ten capabilities as agent alone, human approval required or prohibited, based on impact and reversibility.
4. For each approval, specify the exact action, target, data, destination, evidence and rollback shown to the reviewer.
5. Verify the guardrails, identity scopes, sandbox, egress, memory provenance, logging, monitoring, kill switch and recovery evidence.
6. Decide go, conditional go or no-go, naming the accountable owner, unmet conditions, review date and release evidence.

7. Define the notice, challenge route, human escalation and remedy available to an affected person.

Required evidence

- Reproducible aggregate and segment-risk calculations using only the supplied SIM data.
- A data inventory and autonomy classification for all ten capabilities.
- An approval matrix covering every irreversible or high-impact action.
- A deployment decision whose conditions each have an owner, test, expiry or review date and re-test trigger.
- A human challenge, review and remedy route.

Acceptance criteria

The submission passes when all ten capabilities are classified; every high-impact or irreversible action is prohibited or behind deterministic approval; every release condition can objectively pass or fail; and an affected person can reach a human reviewer and obtain a remedy.

The slide deck remains concept-led. These learner procedures are also reproduced in the Learner Guide.

Alignment

Every question in DISCUSSION-QUESTIONS.md is tagged with the K or A statement it evidences. This activity contributes to **LO3** and is assessed through the Case Study.

This material belongs to Tertiary Infotech Pte Ltd (UEN: 20120096W). All Rights Reserved.