

ACTIVITY PACK

Activity 5 - Draft Your AI Data Governance Policy

Complete learner scenario, practice and security checklist

Course	AI Security for Autonomous AI Agents
Course code	TGS-2025060473
Version	4.0 25 August 2026

Tertiary Infotech Pte Ltd | UEN 20120096W

Readme

Activity 5 - Draft Your AI Data Governance Policy

Purpose: Take a ready-made sample AI data governance policy and adapt it, section by section, to your own team using guided prompts.

What you need

- The **Hermes** agent (or any chatbot your trainer provides)
- The sample policy in this folder: `SAMPLE-AI-DATA-GOVERNANCE-POLICY.md`
- The prompt sheet in this folder: `PROMPTS.md`

Evidence status: SIM - all data is fictional. The sample policy is written for the fictional **Sunset Bay Resort Singapore**. Replace it with your own (also fictional) team details.

Steps

1. Read `SAMPLE-AI-DATA-GOVERNANCE-POLICY.md` once, all the way through. Notice its seven parts: Scope, Roles, Principles, Rules, Audit & Logging, Accountability, Review.
2. Decide the fictional team you will write this for (for example, "front desk", "marketing", "a small clinic").
3. Open the agent and paste the **Scope** section, then use the matching prompt in `PROMPTS.md` to adapt it to your team's AI systems and data.
4. Repeat for **Roles**, **Rules** and **Accountability** - name a real human owner for each responsibility (never the AI).
5. Check the result against the sample: does every AI system have an owner, and does a human sign off on any change the agent makes?
6. Save your adapted policy as your activity output.

What you produce

- A one-to-two page AI data governance policy adapted to your fictional team
- At least one **named human owner** for each key responsibility
- One clear rule stating what the agent may **read**, **write** and **generate**

Reflect (Data Privacy / Job Impact / Ethical Concerns / Cyber Security)

- **Data Privacy:** Which customer data should an agent never be allowed to read? (PDPA in spirit: collect and use only what is needed.)
- **Job Impact:** Who now owns "checking the agent's output" as part of their role?
- **Ethical Concerns:** Your policy says the human is accountable, never the AI. Why does that matter?
- **Cyber Security:** How would you know, weeks later, what the agent read or changed? (Audit & logging.)

© 2026 Tertiary Infotech Pte Ltd | Training use only