

ACTIVITY PACK

Activity 5: Prompts

Activity 5 - Draft Your AI Data Governance Policy

Course	AI Security for Autonomous AI Agents
Course code	TGS-2025060473
Version	4.14 26 August 2026

Tertiary Infotech Pte Ltd | UEN 201200696W

Prompts

Activity 5 - Prompts to Refine Each Section

Evidence status: SIM - all data is fictional. Start with the policy drafted by `ai-data-policy-generator.html` (or the sample in `SAMPLE-AI-DATA-GOVERNANCE-POLICY.md`). Paste one section at a time into the Hermes agent, then send the matching prompt to sharpen it. Keep all names and data fictional.

Set the scene (send once, first)

"I am adapting an AI data governance policy for my fictional team: **[describe your team, e.g. 'a small resort front desk of 5 people']**. We use these AI agents: **[list them, e.g. a WhatsApp guest-support agent and a marketing spreadsheet agent]**. Keep everything fictional and plain-English."

Scope

"Here is the Scope section. Rewrite it for my team. List exactly which AI agents we use and which data assets they touch (bookings, contacts, spreadsheets, generated files). Remove anything that does not apply to us."

Roles

"Here is the Roles section. Assign each role - Data Owner, Agent Owner, Approver, Reviewer - to a job title on my team, and add a one-line responsibility for each. Make sure every role is a human, not the AI."

Principles

"Here is the Principles section. Keep all six principles, but rewrite each one in one short sentence that fits my team's day-to-day work."

Rules

"Here is the Rules section. For my team, list clearly what our agents MAY read, MAY write/change, and MAY generate. State which changes need human approval, and how long we keep files and logs."

Audit & Logging

"Here is the Audit & Logging section. Describe in simple steps how my team would record what each agent did and who would review the logs and how often."

Accountability

"Here is the Accountability section. Rewrite it to name which job title is answerable if an agent produces a wrong output. State plainly that the AI is never accountable."

Review

"Here is the Review section. Set a review cadence for my team and name who owns the review. Add the triggers that would force an early review."

Final check

"Now read my whole adapted policy back and tell me: (1) does every agent have a named human owner? (2) does any change to real data need human approval? Flag anything missing."

© 2026 Tertiary Infotech Pte Ltd | Training use only