

ACTIVITY PACK

Activity 5: Sample Ai Data Governance Policy

Activity 5 - Draft Your AI Data Governance Policy

Course	AI Security for Autonomous AI Agents
Course code	TGS-2025060473
Version	4.0 25 August 2026

Tertiary Infotech Pte Ltd | UEN 20120096W

Sample AI Data Governance Policy - Sunset Bay Resort Singapore

Evidence status: SIM - all data is fictional. A generic sample for training only. It does not describe a real business. Aligned in spirit with the **IMDA Model AI Governance Framework** and the **Singapore Personal Data Protection Act (PDPA)**. No clause numbers are quoted.

Owner of this policy: General Manager, Sunset Bay Resort Singapore **Version:** 1.0 | **Effective:** 25 August 2026 | **Next review:** 25 August 2027

1. Scope

This policy covers every AI system and AI agent used at the resort, and the data they touch.

- **AI systems / agents covered:** the WhatsApp guest-support agent, the marketing-analysis agent (spreadsheets), the slide-building agent, and any chatbot on our website.
- **Data assets covered:** guest booking records, guest contact details, marketing spreadsheets (spend and revenue), generated PowerPoint files and reports, and any training data or examples supplied to an agent.
- **Out of scope:** personal devices and private accounts of staff not used for resort work.

2. Roles

Role	Who (example)	Responsibility
Data Owner	Front Office Manager	Owns guest data; decides what may be shared with an agent.
Agent Owner	Marketing Manager	Owns each AI agent; keeps its settings and skills current.
Approver	General Manager	Approves new agents, new skills, and any change to data rules.
Reviewer	Duty Supervisor	Checks agent output for accuracy before it reaches a guest or manager.

A named human always holds each role. The AI holds none of them.

3. Principles

- **Accuracy** - Output is checked by a human before it is acted upon.
- **Purpose limitation** - Data is used only for the task it was collected for.
- **Protection** - Personal data is kept secure and shared only with approved agents.
- **Provenance** - We know where each piece of data and each agent skill came from.
- **Traceability** - Agent actions can be traced afterwards through logs.
- **Human accountability** - A named person, never the AI, is answerable for outcomes.

4. Rules

What agents MAY read: anonymised or aggregate marketing figures; a single guest's booking only when that guest is being served, and only the fields needed.

What agents MAY write / change: draft replies, draft reports and draft slides. Any change to a live booking, price, or guest record must be **approved by a human** first.

What agents MAY generate: summaries, drafts, slides and analysis - always marked as "AI draft - to be checked."

Approvals for changes: adding a new agent, turning on a new skill, or widening what data an agent can read requires **Approver** sign-off.

Retention: uploaded working files and agent chat logs are kept for 12 months, then deleted, unless needed for a specific record.

5. Audit & Logging

- Every agent keeps a log of what it was asked and what it did.
- Logs record the date, the staff member, the data used and any file produced.
- Logs are reviewed monthly by the Agent Owner and spot-checked by the Approver.

6. Accountability

- Accountability rests with the **named human owner** of each agent and each data asset - **never with the AI**.
- If an agent produces a wrong or harmful output, the Reviewer who released it and the Agent Owner are answerable, and the incident is recorded.

7. Review

- This policy is reviewed **every 12 months**, or sooner if a new agent, skill or data type is introduced, or after any incident.
- The General Manager owns the review and records the changes.

© 2026 Tertiary Infotech Pte Ltd | Training use only