

## ACTIVITY PACK

# Activity 8: Worksheet

Activity 8 - Reflect on Agent Security and Decide Go / No-Go

|             |                                      |
|-------------|--------------------------------------|
| Course      | AI Security for Autonomous AI Agents |
| Course code | TGS-2025060473                       |
| Version     | 4.0   25 August 2026                 |

Tertiary Infotech Pte Ltd | UEN 20120096W

Worksheet

Activity 8 - Go / No-Go Worksheet

Evidence status: SIM - all data is fictional. Fill this in using your Activity 7 findings. Decision is for the fictional Sunset Bay Resort chatbot.

Your name: \_\_\_\_\_ Date: \_\_\_\_\_

Part A - Risk & guardrail table

For each risk, write the real cost if it were live, then the guardrail layer that reduces it.

| Risk (what could leak)      | Real-world cost if live | Guardrail layer that reduces it | Covered? (Y/N) |
|-----------------------------|-------------------------|---------------------------------|----------------|
| All customer bookings       |                         |                                 |                |
| Internal staff memo         |                         |                                 |                |
| Admin password              |                         |                                 |                |
| A guest's phone number      |                         |                                 |                |
| Guest email / booking total |                         |                                 |                |
| (add your own)              |                         |                                 |                |

Guardrail layers to choose from: retrieval filter | input guard | hardened prompt | output guard.

Part B - Safe-rollout checklist

Tick each item only when you can say how it is met.

- ☐ **Scope** - The agent's job and its out-of-bounds tasks are written down: \_\_\_\_\_
- ☐ **Bound** - What it may read / write is defined; risky actions need human approval: \_\_\_\_\_
- ☐ **Test** - The guarded bot was re-tested against the attack prompts and held: \_\_\_\_\_
- ☐ **Approve** - A named human has signed off: \_\_\_\_\_
- ☐ **Pilot** - A small, watched pilot is planned before full rollout: \_\_\_\_\_

Part C - Decision box

Decision (circle one): GO / CONDITIONAL GO / NO-GO

If conditional - the fixes required before launch:

1. \_\_\_\_\_
2. \_\_\_\_\_
3. \_\_\_\_\_

Accountable owner (a named human - never the AI): Name / role:

Kill-switch - how we switch the agent off fast, and who can do it:

Signed: \_\_\_\_\_ Date: \_\_\_\_\_