



Tertiary Infotech Pte Ltd

UEN: 201200696W

LEARNER GUIDE

For

AI Security for Autonomous AI Agents

TGS Ref No: TGS-2025060473

Conducted by

Tertiary Infotech Pte Ltd

UEN: 201200696W

Version 4.15

DOCUMENT VERSION CONTROL RECORD

Version Number	Effective Date of Release	Summary of Included Changes	Author
1.0	16 June 2025	Initial release.	Dr Alfred Ang
2.0	17 August 2026	Retitled and rebuilt for autonomous-agent security; five activities added. TSC K/A text unchanged.	Dr Alfred Ang
2.1	18 August 2026	Added prompt-injection, PDPA, guardrail, approval, skill/plugin and visual coverage. TSC K/A text unchanged.	Dr Alfred Ang
3.0	20 August 2026	Evidence-grounded 207-slide rebuild with product boundaries, cases, controls, activities and source register. TSC K/A text unchanged.	Dr Alfred Ang
4.0	25 August 2026	Rebuilt as a 1-day course with three LO-aligned topics, eight no-code activities and a reflection-based practical assessment. TSC K/A text unchanged.	Dr Alfred Ang
4.1	25 August 2026	Refreshed trainer visuals and embedded the YouTube Online Video; scope, timing and TSC K/A text unchanged.	Dr Alfred Ang
4.2	25 August 2026	Added a coding-harness comparison (Claude Code, Codex, DeepSeek Harness) and a latest-AI-agents reference (OpenClaw, Hermes, Prime Agent, OpenWorker, QM) with official links. TSC K/A text unchanged.	Dr Alfred Ang
4.3	25 August 2026	Beefed up Topic 3 with Singapore case studies (CPA Australia AI adoption vs cybersecurity, CNA gen-AI brand-damage, Tricentis agentic testing, Synapse AgentSea healthcare), a brand-damage discussion and debrief,	Dr Alfred Ang

		a 'Key Challenges of AI to Business' capstone, a visual multi-agent map and a DeepSeek 'everything is a plugin' slide. Visible source citations on all real-life cases. Added tool-agnostic agent-control slides (spot threats; least privilege; approval/validation/shutdown; multi-agent governance; incident response). TSC K/A text unchanged.	
4.4	25 August 2026	Activity 5 rebuilt around the AI Data Policy Generator website (learner-supplied training API key, 7-part IMDA/PDPA-aligned policy framework). Activity walkthrough and governance cheat-sheet aligned to the 7 parts. Added a worked-example slide and cheat-sheet table (Sunset Bay Resort) with concrete Scope, Roles, Rules and Review examples. TSC K/A text unchanged.	Dr Alfred Ang
4.5	25 August 2026	Added Topic 1 slide 'Probabilistic Output — Why Traditional Security Struggles' contrasting deterministic software with probabilistic LLM output and why unpredictable answers defeat signature/rule-based defences. Learner Guide notes renumbered. TSC K/A text unchanged.	Dr Alfred Ang
4.6	25 August 2026	Added Topic 1 slide 'When the Loop Goes Wrong — Misreading the Goal' showing how an open-ended 'remove the malware at all costs' goal can escalate to wiping	Dr Alfred Ang

		every file. Learner Guide notes renumbered. TSC K/A text unchanged.	
4.7	25 August 2026	Added Topic 1 slide 'An Agent on WhatsApp Is Exposed to the Internet' covering the public attack surface of messaging agents — prompt injection and data poisoning leading to confidential-data extraction or server harm. Learner Guide notes renumbered. TSC K/A text unchanged.	Dr Alfred Ang
4.8	26 August 2026	Added verified case study 'Case — A Malicious Skill on ClawHub' (Snyk, Feb 2026) at the end of the Topic 1 Skills and Tools section: a fake Google skill whose SKILL.md instructions tricked users into installing malware. Learner Guide notes renumbered. TSC K/A text unchanged.	Dr Alfred Ang
4.9	26 August 2026	Activity 3 reworked: the agent now creates a 10-slide 'AI Security for AI Agents' deck (saved to the Downloads folder), then redesigns it from an uploaded image template — replacing the Envato-template animated-deck flow. Activity pack, slides and schedule text aligned. TSC K/A text unchanged.	Dr Alfred Ang
4.10	26 August 2026	Activity 2 reworked: learners now upload MOCK-MARKETING-DATA.xlsx to the Hermes + MiniMax agent and prompt it to analyse the data, generate charts and produce a .docx report with insights, analysis and recommendations.	Dr Alfred Ang

		Activity pack, slides and LG aligned. TSC K/A text unchanged.	
4.11	26 August 2026	Added four Topic 1 slides on applications of AI agents in cyber security: threat detection (real-time detection, automated vulnerability analysis, malware classification), network and social-engineering security (intrusion detection, attack classification, phishing defence, log anomaly detection), incident response (patch generation, playbooks, root-cause analysis) and Security Operations Centres (threat hunting, predictive analytics, adaptive decision-making). Learner Guide notes renumbered. TSC K/A text unchanged.	Dr Alfred Ang
4.12	26 August 2026	Added Topic 1 slide 'Agent-to-Agent Messaging - A Security Blind Spot' after the multi-agent worked example: bad or malicious output cascades agent-to-agent, and a growing topology multiplies message paths - reducing visibility and making data-flow tracing and attribution harder. Beefed up Topic 3 security content with ten slides woven into the existing sections: multi-layer bias mitigation; the four structural security gaps of agents (multi-step inputs, tool chaining, opaque execution, untrusted entities); direct vs indirect prompt injection; jailbreaking and jailbreak	Dr Alfred Ang

		persistence; a new 'Advanced Threats — A Look Ahead' section (data poisoning, backdoor attacks, AI supply-chain risk, agent-to-agent attacks, misalignment / Goodhart's Law); and human-in-the-loop at machine speed. Learner Guide notes renumbered. Schedule timing and TSC K/A text unchanged.	
4.13	26 August 2026	Practical instrument renamed from Practical Performance to Case Study across the assessment set, slides, Lesson Plan and Learner Guide. Each Case Study task trimmed to TWO questions: Task 1 keeps the generative-vs-agentic-vs-agent and useful-vs-not-trusted questions; Task 2 keeps the weak-vs-strong prompt and prompt-principles questions; Task 3 keeps the leaky-vs-guarded chatbot and accountability questions. Marks rebalanced (20/24/26, total 70). Schedule timing and TSC K/A text unchanged.	Dr Alfred Ang
4.14	26 August 2026	Learning Outcomes slide now carries the full accredited LO1-LO3 wording. Activity overview slides for Activities 1, 4, 6, 7 and 8 regenerated with activity-specific step flows and 'You'll produce' lines matching the activity packs and Learner Guide. Thank You slide layout corrected. Activity PDFs re-stamped to the current version. Schedule timing and TSC K/A	Dr Alfred Ang

		text unchanged.	
4.15	26 August 2026	Learner Guide and Lesson Plan realigned to the v4.15 deck: Topic 3 job redesign restructured into a narrative sequence (Why Now, The Stakes, Our Operating Model, The Problem, The Solution, The Intervention, The Product, The Worked Example), all slide references renumbered to the 168-slide deck, and Activity 6 extended with the six simulator personas and the A6-Coaching-Sarah worked transcript.	Dr Alfred Ang

TABLE OF CONTENTS

About This Course.....	16
Learning Outcomes.....	16
Knowledge Statements.....	16
Ability Statements.....	16
How to Use This Learner Guide.....	17
Slide 5 — Who Is in the Room?.....	17
Slide 6 — Ground Rules for a Hands-On Day.....	18
Slide 7 — Learning Outcomes.....	19
Slide 8 — One-Day Learning Journey.....	20
Slide 9 — Briefing for Assessment.....	21
Slide 10 — How the Day Is Assessed.....	22
Generative AI, Agentic AI and AI Agents.....	23
Slide 13 — What Topic 1 Will Answer.....	23
Slide 14 — A Very Short History of AI, 2023–2026.....	24
Slide 15 — The Four Waves in Plain Words.....	25
Slide 16 — 2023 — Generative AI Goes Mainstream.....	26
Slide 17 — 2024 — Context Engineering.....	27
Slide 18 — 2025 — Harness Engineering and Agentic AI.....	28
Slide 19 — 2026 — The Year of AI Agents.....	29
How Generative AI Works.....	30
Slide 21 — What 'Generative' Means.....	30
Slide 22 — Two Phases: Training and Inference.....	31
Slide 23 — How an Autoregressive LLM Writes.....	32
Slide 24 — Why the Same Prompt Can Differ.....	33
Slide 25 — Strengths and Limits for Service Work.....	34
Slide 26 — Probabilistic Output — Why Traditional Security Struggles.....	35

Generative AI in the Real World.....	36
Slide 28 — Fashion Models Are Now Generated.....	36
Slide 29 — AI-Generated Video Is Overtaking Whole Industries.....	37
Slide 30 — AI in Marketing and Advertising.....	38
Slide 31 — AI in Film and Movies.....	39
Slide 32 — Customer Service: The Klarna Story.....	40
Slide 33 — Where Service Teams Apply GenAI Today.....	41
Context Engineering.....	42
Slide 35 — Prompt Engineering vs Context Engineering.....	42
Slide 36 — The Six Ingredients of Context.....	43
Slide 37 — How an Agent Gathers Context for a Task.....	44
Slide 38 — A Worked Context Example.....	45
The Agentic Loop and Harness Engineering.....	46
Slide 40 — The Agentic Loop — A Visual Pipeline.....	46
Slide 41 — What a Harness Adds Around the Model.....	47
Slide 42 — Two Harnesses You May Have Heard Of.....	48
Slide 43 — The Coding Harnesses — Claude Code, Codex, DeepSeek.....	49
Slide 44 — DeepSeek Harness — 'Everything Is a Plugin'.....	50
Slide 45 — A Concrete Agentic Loop.....	51
Slide 46 — When the Loop Goes Wrong — Misreading the Goal.....	52
AI Agents: OpenClaw, Hermes and the Latest.....	53
Slide 48 — What Makes Something an 'AI Agent'.....	53
Slide 49 — Peter Steinberger in 2026 — Creator of OpenClaw.....	54
Slide 50 — OpenClaw — A Naming and Adoption Story.....	55
Slide 51 — OpenClaw Runs Where You Already Chat.....	56
Slide 52 — Hermes Agent by Nous Research.....	57

Slide 53 — The Latest AI Agents (2025–2026).....	58
Slide 54 — Applications of AI Agents in Threat Detection.....	59
Slide 55 — AI Agents in Network and Social-Engineering Security.....	60
Slide 56 — Applications of AI Agents in Incident Response.....	61
Slide 57 — Applications of AI in Security Operations Centres (SOCs).....	62
Slide 58 — Try It Today: TIA Support on WhatsApp.....	63
Slide 59 — An Agent on WhatsApp Is Exposed to the Internet.....	64
Skills and Tools.....	65
Slide 61 — Tools vs Skills.....	65
Slide 62 — How an Agent Uses a Tool.....	66
Slide 63 — Why Skills Make Agents Practical.....	67
Slide 64 — Case — A Malicious Skill on ClawHub.....	68
Multi-Agent Systems.....	69
Slide 67 — How a Multi-Agent System Works.....	69
Slide 68 — Why and Why Not Multi-Agent.....	70
Slide 69 — A Service Example — Customer Recovery.....	71
Slide 70 — Agent-to-Agent Messaging — A Security Blind Spot.....	72
Generative AI vs Agentic AI vs AI Agents.....	73
Slide 72 — The Three, Side by Side.....	73
Slide 73 — One-Line Definitions.....	74
Slide 74 — Generation is a step. Agency is a loop. An agent is a system.....	75
Slide 75 — Activity 1 — Talk to an AI Agent, Then Reflect.....	76
Slide 76 — Activity 1 — Group Your Concerns.....	77
Slide 77 — Activity 1 — Debrief.....	78
Prompt Engineering and Post-Training for Autonomous AI Agents.....	79
Slide 79 — What Topic 2 Will Cover.....	79

Slide 80 — The Setup in Three Parts.....	80
Slide 81 — Install the Hermes Desktop App.....	81
Slide 82 — Get a MiniMax Model and API Key.....	82
Slide 83 — Point Hermes at MiniMax M2.7.....	83
Slide 84 — Keep Your Key Safe.....	84
Prompt Engineering Principles.....	85
Slide 86 — Five Principles of a Good Prompt.....	85
Slide 87 — A Bad Prompt vs a Good Prompt.....	86
Slide 88 — Techniques That Reliably Help.....	87
Slide 89 — 'Post-Training' for Agents — Memory and Skills, Not Fine-Tuning.....	88
Slide 90 — Reflect — Does Prompt Engineering Still Matter With Agents?.....	89
Slide 91 — Activity 2 — Analyse Excel Data with the Agent.....	90
Slide 92 — Activity 2 — From Spreadsheet to Report.....	91
Slide 93 — Activity 3 — Create and Redesign a PPT with the Agent.....	92
Slide 94 — Activity 3 — What Good Looks Like.....	93
Slide 95 — Activity 4 — Install Tools and Skills to Do Better.....	94
Slide 97 — Reflect on Working with the Agent.....	95
Slide 98 — Topic 2 Debrief.....	96
Security Risk of Autonomous AI Agents.....	97
Slide 100 — What Topic 3 Will Cover.....	97
Slide 101 — Discussion Point 1 — Who Is Accountable for the Data?.....	98
Slide 102 — AI is never the accountable party. A named human always is.....	99
Slide 103 — Case — Moffatt v Air Canada.....	100
Slide 104 — Traditional Data Governance.....	101
Slide 105 — New Principles for GenAI and Agents.....	102
Slide 106 — Scope of an AI Data Governance Policy.....	103

Slide 107 — The 7-Part AI Data Governance Policy Framework.....	104
Slide 108 — The Sample Policy in Action — Sunset Bay Resort.....	105
Slide 109 — Activity 5 — Draft Your AI Data Governance Policy.....	106
Slide 110 — Discussion Point 2 — Brand Damage from Generative AI.....	107
Slide 111 — Case — Gen-AI Ads Flood Singapore, and Could Backfire.....	108
Slide 112 — Why AI Content Can Damage a Brand.....	109
Slide 113 — Discuss & Debrief — Protecting the Brand.....	110
Slide 114 — Discussion Point 3 — If Agents Take the Work, What Do People Do? ..	111
Slide 115 — Jobs Move — Not Vanish.....	112
Slide 116 — Why Now — 12,000+ Agents in Singapore Healthcare.....	113
Slide 117 — The Stakes — The Numbers Give Perspective.....	114
Slide 118 — Our Operating Model — A Human at Every Consequential Gate.....	115
Slide 119 — The Problem — Adoption Debt When Roles Lag Behind.....	116
Slide 120 — The Solution — An Explicit Human Owner Role.....	116
Slide 121 — Coaching Staff Through the Change.....	116
Slide 122 — A Job-Redesign Framework.....	118
Slide 123 — The Intervention — Use GROW to Turn Resistance into Redesign.....	119
Slide 124 — The Product — The Job-Redesign Coaching Simulator.....	120
Slide 125 — The Worked Example — Sarah's Redesigned Role.....	120
Slide 126 — Activity 6 — Coach a Worried Team Member (Role-Play).....	121
AI Agent Cybersecurity Risks.....	122
Slide 128 — The Singapore Reality — Fast on AI, Slower on Security.....	122
Slide 129 — Discussion Point 4 — An Agent With Too Much Reach.....	123
Slide 130 — Capable Agents, Not 'Rogue' Ones.....	124
Slide 131 — AI Cybersecurity Risks and Mitigation Strategies.....	125
Slide 132 — Going Deeper — Mitigating Bias in AI Systems.....	126

Slide 133 — Spot the Threats Across an Agent Workflow.....	127
Slide 134 — Why Agents Are Hard to Secure — Four Gaps.....	128
Slide 135 — Why Offence Outruns Defence.....	129
Slide 136 — The Response Being Proposed.....	130
Slide 137 — Case — Replit Agent Deletes a Database.....	131
Slide 138 — Prompt Injection and PII Leak — In Plain Terms.....	132
Slide 139 — Prompt Injection — Direct and Indirect.....	133
Slide 140 — Jailbreaking — and Why It Persists.....	134
Slide 141 — Activity 7 — Break a Leaky Chatbot.....	135
Slide 142 — Activity 7 — Compare the Guarded Chatbot.....	136
Slide 144 — Four Guardrail Layers That Stopped the Leak.....	137
Advanced Threats — A Look Ahead.....	138
Slide 145 — Data Poisoning and Backdoor Attacks.....	138
Slide 146 — The AI Supply Chain Is Hard to Secure.....	139
Slide 147 — Agent-to-Agent Attacks.....	140
Slide 148 — Misalignment — Goodhart's Law in Action.....	141
Core Controls for a Live Agent.....	142
Slide 150 — Control 1 — Least Privilege, Allowlists and Credentials.....	142
Slide 151 — Control 2 — Approval, Output Validation and Shutdown.....	143
Slide 152 — Control 3 — Govern Multi-Agent Coordination.....	144
Slide 153 — Control 4 — Investigate and Respond to a Compromised Agent.....	145
Slide 154 — A Framework to Roll Out Agents Safely.....	146
Slide 155 — The Go-Live Gate.....	147
Slide 156 — Human-in-the-Loop for Risky Actions.....	148
Slide 157 — Human-in-the-Loop at Machine Speed.....	149
Slide 158 — Case — Testing AI Before You Trust It.....	150

Slide 159 — Activity 8 — Reflect on Agent Security.....	151
Slide 160 — The Key Challenges of AI to Business.....	152
Slide 161 — Topic 3 Recap.....	153
Slide 162 — Adopt AI boldly — but govern the data, guard the brand, and keep a human accountable.....	154
Slide 163 — Assessment Reminder.....	155
Slide 164 — What Each Assessment Asks.....	156
Slide 166 — Assessment Flow.....	157
Detailed Activity Walkthroughs.....	158
Activity 1 — Talk to an AI Agent, Then Reflect.....	159
Step-by-step.....	159
What you produce.....	159
Activity 2 — Analyse Excel Marketing Data with the Agent.....	160
Step-by-step.....	160
What you produce.....	160
Activity 3 — Create and Redesign a PowerPoint with the Agent.....	161
Step-by-step.....	161
What you produce.....	161
Activity 4 — Install Tools and Skills to Do Better.....	162
Step-by-step.....	162
What you produce.....	162
Activity 5 — Draft Your AI Data Governance Policy.....	163
Step-by-step.....	163
What you produce.....	163
Activity 6 — Coach a Worried Team Member (Role-Play Simulator).....	164
Step-by-step.....	164
What you produce.....	164

Activity 7 — Break a Leaky Chatbot, Then Compare the Guarded One.....	165
Step-by-step.....	165
What you produce.....	165
Activity 8 — Reflect on Agent Security and Decide Go / No-Go.....	166
Step-by-step.....	166
What you produce.....	166
Prompt Engineering Quick Reference.....	167
Good vs Bad Prompts.....	167
Reflection Framework — Four Risk Themes.....	168
AI Data Governance — One-Page Cheat Sheet.....	169
Safe Roll-Out Checklist for AI Agents.....	170
Source Register.....	171

About This Course

AI Security for Autonomous AI Agents is a 1 Day · 8 Hours WSQ course introducing generative AI, agentic AI and autonomous AI agents, and the security, governance and ethical risks they bring to customer-service and hospitality settings. Learners work entirely on ready-made websites, chatbots and AI agents — no coding is required. The day is organised as three topics mapped to the three learning outcomes.

This guide follows the trainer deck in sequence. Each slide is expanded into readable notes, and every one of the eight activities has a full step-by-step, no-code walkthrough.

Learning Outcomes

LO	Learning Outcome
LO1	Demonstrate generative AI concepts and applications relevant to customer service and hospitality management.
LO2	Apply prompt engineering techniques and analyse output variations to improve generative AI performance in service settings.
LO3	Identify ethical risks and analyse bias in AI-generated content used in customer engagement.

Knowledge Statements

Code	Knowledge statement
K1	Importance of data quality, preprocessing, model pipeline and model training (e.g., impact of data bias from training data)
K2	Underlying principles, core concepts and theories governing generative AI
K3	Difference between generative and discriminative models
K4	Impact of prompt engineering on the model outputs of generative AI
K5	Generative AI model workings, including training data, algorithms, and outputs

Ability Statements

Code	Ability statement
A1	Analyse limitations and potential biases in AI-generated content
A2	Identify the ethical implications and societal impact of AI-generated content
A3	Apply understanding of generative AI principles to use cases
A4	Demonstrate the use of generation AI in diverse applications (e.g., summarisation, inference, reasoning, transformation of content, augmentation of content)
A5	Analyse generative AI models' performance metrics and evaluate the influence of prompt variations

How to Use This Learner Guide

This guide follows the trainer deck in sequence across the three topics of the day. It expands each concept into readable notes, then gives you full step-by-step walkthroughs for every activity. You will use ready-made websites, chatbots and AI agents — you do not need to write any code. All data used in the activities is fictional.

Topic	Learning outcome	What you will be able to do
Topic 1	LO1	Explain generative AI, agentic AI and AI agents, and how they differ.
Topic 2	LO2	Apply prompt-engineering techniques and compare output variations.
Topic 3	LO3	Identify ethical, governance and security risks in AI-generated content and agents.

EVIDENCE LABELS

HIST = historical fact · DEF = definition · PROD = product documentation · CASE-V = verified/real dated case · SIM = classroom simulation (all names and numbers fictional) · SYN = teaching synthesis.

SAFE-USE RULE

Never paste real personal data or a production API key into any demo. Use a low-limit training key and delete it after the course.

Slide 5 — Who Is in the Room?

EVIDENCE: ADMIN

Concept or classroom slide; no external factual claim is introduced here.

Read the points below as the few things worth remembering about “Who Is in the Room?”.

Concept	What it means
Role	What service or team do you work in?
AI today	Where do you already meet AI at work?
Goal	What do you want to be able to do by 6pm?

Why it matters at work: In a customer-service or hospitality setting, this shapes how you brief an AI tool and how much you trust its answer before it reaches a guest.

Slide 6 — Ground Rules for a Hands-On Day

EVIDENCE: ADMIN

Concept or classroom slide; no external factual claim is introduced here.

Read the points below as the few things worth remembering about “Ground Rules for a Hands-On Day”.

Concept	What it means
Use the demos	Every activity is a ready-made website or chatbot — no coding
Fictional data only	Never paste real personal or company data into a demo
Ask and share	Try prompts, compare answers, and reflect together

Why it matters at work: In a customer-service or hospitality setting, this shapes how you brief an AI tool and how much you trust its answer before it reaches a guest.

Slide 7 — Learning Outcomes

EVIDENCE: ADMIN

Concept or classroom slide; no external factual claim is introduced here.

Read the points below as the few things worth remembering about “Learning Outcomes”.

Concept	What it means
LO1	Demonstrate generative AI concepts and applications relevant to customer service and hospitality management
LO2	Apply prompt engineering techniques and analyse output variations to improve generative AI performance in service settings
LO3	Identify ethical risks and analyse bias in AI-generated content used in customer engagement

Why it matters at work: In a customer-service or hospitality setting, this shapes how you brief an AI tool and how much you trust its answer before it reaches a guest.

Slide 8 — One-Day Learning Journey

EVIDENCE: ADMIN

Concept or classroom slide; no external factual claim is introduced here.

“One-Day Learning Journey” is best understood as a sequence — the steps below run in order. Each step feeds the next, so a problem early on shows up later in the chain — which is exactly why the order is worth learning. Three topics map one-to-one onto LO1, LO2 and LO3; the assessment closes the day.

1. Topic 1 — GenAI, Agentic AI & Agents
2. Topic 2 — Prompt Engineering
3. Topic 3 — Security & Governance
4. Assessment

NOTE

Three topics map one-to-one onto LO1, LO2 and LO3; the assessment closes the day.

Why it matters at work: In a customer-service or hospitality setting, this shapes how you brief an AI tool and how much you trust its answer before it reaches a guest.

Slide 9 — Briefing for Assessment

EVIDENCE: ADMIN

Concept or classroom slide; no external factual claim is introduced here.

“Briefing for Assessment” is best understood as a sequence — the steps below run in order. Each step feeds the next, so a problem early on shows up later in the chain — which is exactly why the order is worth learning. The briefing for assessment is delivered before the assessment begins.

1. Read each task
2. Draw on today's activities
3. Write from your own evidence
4. Justify your reasoning
5. Submit on the LMS

NOTE

The briefing for assessment is delivered before the assessment begins.

Why it matters at work: In a customer-service or hospitality setting, this shapes how you brief an AI tool and how much you trust its answer before it reaches a guest.

Slide 10 — How the Day Is Assessed

EVIDENCE: ADMIN

Concept or classroom slide; no external factual claim is introduced here.

The table below is easier to recall once you see the pattern behind it.

Instrument	Covers	What you do
Written Assessment (SAQ)	K1–K5	Five short written answers, one per knowledge statement
Case Study	LO1–LO3 (A1–A5)	Three reflection tasks (two questions each) on the activities you completed today
Format	—	Open book · Competent / Not Yet Competent

Why it matters at work: In a customer-service or hospitality setting, this shapes how you brief an AI tool and how much you trust its answer before it reaches a guest.

Generative AI, Agentic AI and AI Agents

[TOPIC 1 · LO1]

From a chat box to systems that plan and act — history, mechanics, use cases and agents.

Slide 13 — What Topic 1 Will Answer

EVIDENCE: SYN

Concept or classroom slide; no external factual claim is introduced here.

“What Topic 1 Will Answer” is best understood as a sequence — the steps below run in order. Each step feeds the next, so a problem early on shows up later in the chain — which is exactly why the order is worth learning. Each question builds the vocabulary you need for prompt engineering and security later.

1. How did we get here?
2. How does GenAI actually work?
3. What is context engineering?
4. What is an agentic loop?
5. What is an AI agent?

NOTE

Each question builds the vocabulary you need for prompt engineering and security later.

Why it matters at work: In a customer-service or hospitality setting, this shapes how you brief an AI tool and how much you trust its answer before it reaches a guest.

Slide 14 — A Very Short History of AI, 2023–2026

EVIDENCE: HIST

Concept or classroom slide; no external factual claim is introduced here.

“A Very Short History of AI, 2023–2026” is best understood as a sequence — the steps below run in order. Each step feeds the next, so a problem early on shows up later in the chain — which is exactly why the order is worth learning. Each year did not replace the last — it added a new layer on top of it.

1. 2023 · Generative AI
2. 2024 · Context Engineering
3. 2025 · Harness Engineering
4. 2026 · AI Agents

NOTE

Each year did not replace the last — it added a new layer on top of it.

Why it matters at work: In a customer-service or hospitality setting, this shapes how you brief an AI tool and how much you trust its answer before it reaches a guest.

Slide 15 — The Four Waves in Plain Words

EVIDENCE: SYN

Concept or classroom slide; no external factual claim is introduced here.

Read the points below as the few things worth remembering about “The Four Waves in Plain Words”.

Concept	What it means
2023 — Generative AI	ChatGPT makes text and image generation mainstream
2024 — Context Engineering	We learn to feed models the right information, not just a prompt
2025 — Harness Engineering	We wrap models in loops that plan, act and verify — agentic AI
2026 — AI Agents	Deployed agents act across our apps, chats and tools

Why it matters at work: In a customer-service or hospitality setting, this shapes how you brief an AI tool and how much you trust its answer before it reaches a guest.

Slide 16 — 2023 — Generative AI Goes Mainstream

EVIDENCE: HIST

Sources: S08 — OpenAI — Introducing ChatGPT (30 Nov 2022)

The points below unpack what “2023 — Generative AI Goes Mainstream” means in practice.

- ChatGPT launched on 30 November 2022 and reached mass adoption through 2023
- One natural-language box exposed writing, coding, translation and analysis to everyone
- Businesses began asking not 'can it chat?' but 'what work can it do?'

Why it matters at work: In a customer-service or hospitality setting, this shapes how you brief an AI tool and how much you trust its answer before it reaches a guest.

Slide 17 — 2024 — Context Engineering

EVIDENCE: HIST

Sources: S10 — Anthropic — Effective context engineering for AI agents (Sep 2025), S11 — Karpathy on 'context engineering' (X, 25 Jun 2025)

“2024 — Context Engineering” is worth a closer look.

- Beyond the promptThe model also needs the right supporting information
- The context windowInstructions, history, memory, tools and retrieved data meet here
- The craftSelect just the right tokens for the model's next step

Why it matters at work: In a customer-service or hospitality setting, this shapes how you brief an AI tool and how much you trust its answer before it reaches a guest.

Slide 18 — 2025 — Harness Engineering and Agentic AI

EVIDENCE: HIST

Sources: S12 — Anthropic — How Claude Code works (agentic harness), S13 — OpenAI — Harness engineering: Codex in an agent-first world

“2025 — Harness Engineering and Agentic AI” is worth a closer look.

- LoopGather context, plan, act and verify until the goal is met
- ToolsThe harness gives the model controlled ways to affect real systems
- SafetyPermissions, sandboxes and verification constrain each action

Why it matters at work: In a customer-service or hospitality setting, this shapes how you brief an AI tool and how much you trust its answer before it reaches a guest.

Slide 19 — 2026 — The Year of AI Agents

EVIDENCE: HIST

Sources: S30 — OpenClaw — Wikipedia (naming and adoption history), S33 — Hermes Agent (Nous Research) — documentation

The points below unpack what “2026 — The Year of AI Agents” means in practice.

- Personal and organisational agents such as OpenClaw and Hermes reached everyday users
- Agents now live inside WhatsApp, Telegram, email and business apps
- The question shifted again: how do we let agents act safely on our behalf?

Why it matters at work: In a customer-service or hospitality setting, this shapes how you brief an AI tool and how much you trust its answer before it reaches a guest.

How Generative AI Works

[TOPIC 1]

Training, inference, and the autoregressive language model.

Slide 21 — What 'Generative' Means

EVIDENCE: DEF

Concept or classroom slide; no external factual claim is introduced here.

Read the points below as the few things worth remembering about “What 'Generative' Means”.

Concept	What it means
Generative AI	Produces new text, image, audio, video or code from patterns it has learned
Discriminative AI	Sorts or scores existing input — spam / not-spam, fraud / not-fraud
Why it matters	Service teams use both: generate a reply, and classify a request

Why it matters at work: In a customer-service or hospitality setting, this shapes how you brief an AI tool and how much you trust its answer before it reaches a guest.

Slide 22 — Two Phases: Training and Inference

EVIDENCE: DEF

Sources: S06 — Brown et al., Language Models are Few-Shot Learners (GPT-3)

“Two Phases: Training and Inference” is best understood as a sequence — the steps below run in order. Each step feeds the next, so a problem early on shows up later in the chain — which is exactly why the order is worth learning. Training happens once and is expensive; inference happens every time you use the model.

1. Gather large text corpora
2. Train — adjust billions of parameters
3. Ship the model
4. Inference — you prompt, it responds

NOTE

Training happens once and is expensive; inference happens every time you use the model.

Why it matters at work: In a customer-service or hospitality setting, this shapes how you brief an AI tool and how much you trust its answer before it reaches a guest.

Slide 23 — How an Autoregressive LLM Writes

EVIDENCE: DEF

Sources: S05 — Vaswani et al., Attention Is All You Need (2017)

“How an Autoregressive LLM Writes” is best understood as a sequence — the steps below run in order. Each step feeds the next, so a problem early on shows up later in the chain — which is exactly why the order is worth learning. A large language model writes one token at a time, each based on everything before it.

1. Read your context as tokens
2. Predict the next token's probabilities
3. Pick one token
4. Append it and repeat

NOTE

A large language model writes one token at a time, each based on everything before it.

Why it matters at work: In a customer-service or hospitality setting, this shapes how you brief an AI tool and how much you trust its answer before it reaches a guest.

Slide 24 — Why the Same Prompt Can Differ

EVIDENCE: DEF

Concept or classroom slide; no external factual claim is introduced here.

Read the points below as the few things worth remembering about “Why the Same Prompt Can Differ”.

Concept	What it means
Probabilities	The model samples from likely next tokens, not a fixed lookup
Temperature	Higher settings add variety; lower settings add consistency
Context	Change the surrounding information and the whole answer shifts

Why it matters at work: In a customer-service or hospitality setting, this shapes how you brief an AI tool and how much you trust its answer before it reaches a guest.

Slide 25 — Strengths and Limits for Service Work

EVIDENCE: DEF

Concept or classroom slide; no external factual claim is introduced here.

Read the points below as the few things worth remembering about “Strengths and Limits for Service Work”.

Concept	What it means
Strong at	Drafting, summarising, translating, reformatting, brainstorming
Weak at	Facts it was never given — it can sound confident yet be wrong
Rule	Treat fluent output as a draft to verify, not as truth

Why it matters at work: In a customer-service or hospitality setting, this shapes how you brief an AI tool and how much you trust its answer before it reaches a guest.

Slide 26 — Probabilistic Output — Why Traditional Security Struggles

EVIDENCE: DEF

Sources: S43 — OWASP Top 10 for LLM Applications

It helps to set the two sides of “Probabilistic Output — Why Traditional Security Struggles” against each other in the table below. The difference between the two columns is the thing you should be able to explain in your own words after the class. You cannot write a rule for every possible answer — so security must also watch what the model says and does, not just what goes in.

Traditional software	A large language model
Deterministic — the same input always produces the same output	Probabilistic — the same prompt can produce a different answer every run
Behaviour can be fully tested before release	The space of possible outputs can never be fully tested
Firewalls and filters match known, fixed patterns	There is no fixed signature — a harmful answer can look new every time
An unexpected output is a bug you can reproduce and patch	An unexpected output is normal behaviour, not a fault

NOTE

You cannot write a rule for every possible answer — so security must also watch what the model says and does, not just what goes in.

Why it matters at work: In a customer-service or hospitality setting, this shapes how you brief an AI tool and how much you trust its answer before it reaches a guest.

Generative AI in the Real World

[TOPIC 1]

Real, dated examples — how generation is already changing work.

Slide 28 — Fashion Models Are Now Generated

EVIDENCE: CASE-V

Sources: S20 — PetaPixel — Mango launches photorealistic AI-generated campaign (Jul 2024), S21 — Process Excellence Network — H&M debuts AI 'digital twins' of models (Mar 2025), S22 — ABC News — AI-generated models in Guess ad in Vogue (Aug 2025)

“Fashion Models Are Now Generated” is worth a closer look.

Why it matters at work: In a customer-service or hospitality setting, this shapes how you brief an AI tool and how much you trust its answer before it reaches a guest.

Slide 29 — AI-Generated Video Is Overtaking Whole Industries

EVIDENCE: SYN

Concept or classroom slide; no external factual claim is introduced here.

“AI-Generated Video Is Overtaking Whole Industries” is worth a closer look. The clip is embedded — click it and press Play in slideshow.

- MusicThis AI-generated singer — no camera, cast or crew — shows how AI is reshaping the music industry
- Three industriesThe next three demos show marketing, film and how fast this is moving
- Made for feedsVertical clips like this are built for TikTok, Reels and YouTube Shorts
- For your teamAsk: where could you use this, and where would you not?

NOTE

The clip is embedded — click it and press Play in slideshow.

Why it matters at work: In a customer-service or hospitality setting, this shapes how you brief an AI tool and how much you trust its answer before it reaches a guest.

Slide 30 — AI in Marketing and Advertising

EVIDENCE: SYN

Concept or classroom slide; no external factual claim is introduced here.

“AI in Marketing and Advertising” is worth a closer look. Embedded clip — click and press Play. Discuss: would you trust an all-AI ad for your brand?

- Campaign in hours Brands now generate ad and campaign video in hours, not weeks
- No production crew No shoot, studio, cast or location — just a prompt and a model
- Coca-Cola Aired an AI-generated 'Holidays Are Coming' ad in 2024 and again in 2025
- Watch for Brand risk, authenticity and disclosure when the whole ad is synthetic

NOTE

Embedded clip — click and press Play. Discuss: would you trust an all-AI ad for your brand?

Why it matters at work: In a customer-service or hospitality setting, this shapes how you brief an AI tool and how much you trust its answer before it reaches a guest.

Slide 31 — AI in Film and Movies

EVIDENCE: SYN

Concept or classroom slide; no external factual claim is introduced here.

“AI in Film and Movies” is worth a closer look. Embedded clip — click and press Play.
Discuss: what does this mean for creative jobs?

- Into film production AI now generates visuals, scenes and effects for film
- Faster and cheaper Shots that needed a crew and budget can be generated
- Still maturing Consistency, control and rights are the open questions
- The takeaway Whole creative industries are being reshaped — fast

NOTE

Embedded clip — click and press Play. Discuss: what does this mean for creative jobs?

Why it matters at work: In a customer-service or hospitality setting, this shapes how you brief an AI tool and how much you trust its answer before it reaches a guest.

Slide 32 — Customer Service: The Klarna Story

EVIDENCE: CASE-V

Sources: S25 — Klarna — AI assistant handles two-thirds of chats in first month (27 Feb 2024), S26 — Entrepreneur — Klarna CEO reverses course, hiring more humans not AI (May 2025)

“Customer Service: The Klarna Story” is best understood as a sequence — the steps below run in order. Each step feeds the next, so a problem early on shows up later in the chain — which is exactly why the order is worth learning. A real cautionary tale: automate the routine, but keep humans for judgement and empathy.

1. Feb 2024 — AI assistant handles 2/3 of chats
2. Work of ~700 agents in month one
3. Resolution time 11 min to under 2 min
4. May 2025 — re-hires humans for quality

NOTE

A real cautionary tale: automate the routine, but keep humans for judgement and empathy.

Why it matters at work: In a customer-service or hospitality setting, this shapes how you brief an AI tool and how much you trust its answer before it reaches a guest.

Slide 33 — Where Service Teams Apply GenAI Today

EVIDENCE: SYN

Concept or classroom slide; no external factual claim is introduced here.

Read the points below as the few things worth remembering about “Where Service Teams Apply GenAI Today”.

Concept	What it means
Front desk	Draft replies, translate, summarise long threads
Back office	Turn notes into reports; extract data from documents
Marketing	Generate copy, images and short video for campaigns

Why it matters at work: In a customer-service or hospitality setting, this shapes how you brief an AI tool and how much you trust its answer before it reaches a guest.

Context Engineering

[TOPIC 1]

The model is only as good as the context you give it.

Slide 35 — Prompt Engineering vs Context Engineering

EVIDENCE: DEF

Sources: S10 — Anthropic — Effective context engineering for AI agents (Sep 2025), S11 — Karpathy on 'context engineering' (X, 25 Jun 2025)

Read the points below as the few things worth remembering about “Prompt Engineering vs Context Engineering”.

Concept	What it means
Prompt engineering	Wording one instruction well
Context engineering	Assembling everything the model sees before it answers
The shift	Real applications manage context, not just a clever prompt

Why it matters at work: In a customer-service or hospitality setting, this shapes how you brief an AI tool and how much you trust its answer before it reaches a guest.

Slide 36 — The Six Ingredients of Context

EVIDENCE: DEF

Sources: S10 — Anthropic — Effective context engineering for AI agents (Sep 2025)

Read the points below as the few things worth remembering about “The Six Ingredients of Context”.

Concept	What it means
System prompt	Who the AI is and its rules
User prompt	The task you asked for right now
History	Earlier turns in this conversation
Memory	Facts kept across sessions
Tools	Functions it can call to act or fetch data
Retrieved data	Documents pulled in for this task

Why it matters at work: In a customer-service or hospitality setting, this shapes how you brief an AI tool and how much you trust its answer before it reaches a guest.

Slide 37 — How an Agent Gathers Context for a Task

EVIDENCE: SYN

Sources: S10 — Anthropic — Effective context engineering for AI agents (Sep 2025)

“How an Agent Gathers Context for a Task” is best understood as a sequence — the steps below run in order. Each step feeds the next, so a problem early on shows up later in the chain — which is exactly why the order is worth learning. Example: 'Reply to this guest email' → it pulls the booking, the policy doc and past replies first.

1. You give a goal
2. It reads its system prompt and memory
3. It retrieves relevant files or data
4. It calls tools for fresh facts
5. It assembles all of this, then answers

NOTE

Example: 'Reply to this guest email' → it pulls the booking, the policy doc and past replies first.

Why it matters at work: In a customer-service or hospitality setting, this shapes how you brief an AI tool and how much you trust its answer before it reaches a guest.

Slide 38 — A Worked Context Example

EVIDENCE: SYN

Concept or classroom slide; no external factual claim is introduced here.

Read the points below as the few things worth remembering about “A Worked Context Example”.

Concept	What it means
Task	'Draft a refund reply to Mr Tan'
Context gathered	Booking record + refund policy + tone guide + the original email
Result	A grounded reply, not a generic guess

Why it matters at work: In a customer-service or hospitality setting, this shapes how you brief an AI tool and how much you trust its answer before it reaches a guest.

The Agentic Loop and Harness Engineering

[TOPIC 1]

What turns a model that answers into a system that acts.

Slide 40 — The Agentic Loop — A Visual Pipeline

EVIDENCE: DEF

Sources: S12 — Anthropic — How Claude Code works (agentic harness)

“The Agentic Loop — A Visual Pipeline” is best understood as a sequence — the steps below run in order. Each step feeds the next, so a problem early on shows up later in the chain — which is exactly why the order is worth learning. The loop repeats until the goal is met. Anthropic's harness uses gather → act → verify; 'plan' is shown here as a teaching step.

1. Context — gather goal & information
2. Plan — decide the next step
3. Execute — call a tool / take an action
4. Verify — check the result

NOTE

The loop repeats until the goal is met. Anthropic's harness uses gather → act → verify; 'plan' is shown here as a teaching step.

Why it matters at work: In a customer-service or hospitality setting, this shapes how you brief an AI tool and how much you trust its answer before it reaches a guest.

Slide 41 — What a Harness Adds Around the Model

EVIDENCE: DEF

Sources: S12 — Anthropic — How Claude Code works (agentic harness), S13 — OpenAI — Harness engineering: Codex in an agent-first world

Read the points below as the few things worth remembering about “What a Harness Adds Around the Model”.

Concept	What it means
The loop	Runs context → plan → execute → verify repeatedly
Tool access	Lets the model read files, run commands, call APIs
State & safety	Tracks progress and enforces permissions and sandboxes

Why it matters at work: In a customer-service or hospitality setting, this shapes how you brief an AI tool and how much you trust its answer before it reaches a guest.

Slide 42 — Two Harnesses You May Have Heard Of

EVIDENCE: PROD

Sources: S12 — Anthropic — How Claude Code works (agentic harness), S13 — OpenAI — Harness engineering: Codex in an agent-first world

Read the points below as the few things worth remembering about “Two Harnesses You May Have Heard Of”.

Concept	What it means
Claude Code	Anthropic's coding agent — the 'agentic harness around Claude'
Codex	OpenAI's coding agent; OpenAI calls the discipline 'harness engineering'
Same idea	A loop + tools + a sandbox around a general model

Why it matters at work: In a customer-service or hospitality setting, this shapes how you brief an AI tool and how much you trust its answer before it reaches a guest.

Slide 43 — The Coding Harnesses — Claude Code, Codex, DeepSeek

EVIDENCE: PROD

Sources: S53 — Anthropic — Claude Code (repository), S54 — OpenAI — Codex (repository), S55 — DeepSeek Harness (dsh) — developer preview repository

The table below is easier to recall once you see the pattern behind it. All three wrap a model in the same gather-context → act → verify loop; they differ in tools, sandbox and host authority.

Harness	What it is	Where to find it
Claude Code	Anthropic's terminal coding agent — reads your codebase, edits files, runs tests and git by natural language	github.com/anthropics/claude-code
Codex	OpenAI's coding agent; its execution engine (the 'harness') was open-sourced in Aug 2026	github.com/openai/codex
DeepSeek Harness (dsh)	DeepSeek's plugin-based, model-agnostic harness — a v0.1 developer preview	github.com/deepseek-ai/deepseek-harness

NOTE

All three wrap a model in the same gather-context → act → verify loop; they differ in tools, sandbox and host authority.

Why it matters at work: In a customer-service or hospitality setting, this shapes how you brief an AI tool and how much you trust its answer before it reaches a guest.

Slide 44 — DeepSeek Harness — 'Everything Is a Plugin'

EVIDENCE: PROD

Sources: S55 — DeepSeek Harness (dsh) — developer preview repository, S65 — MindStudio — DeepSeek Harness: agentic coding where everything is a plugin

Read the points below as the few things worth remembering about “DeepSeek Harness — 'Everything Is a Plugin’”. A v0.1 developer preview. Its openness is powerful; in Topic 3 we treat each plugin as untrusted until reviewed.

Concept	What it means
One idea	Every capability — tools, models, UI, even core behaviour — is a swappable plugin
Model-agnostic	Plug in any model; the harness is not tied to DeepSeek's own
Why it matters	Flexible and extensible — but every plugin is also new code and a new trust boundary to review

NOTE

A v0.1 developer preview. Its openness is powerful; in Topic 3 we treat each plugin as untrusted until reviewed.

Why it matters at work: In a customer-service or hospitality setting, this shapes how you brief an AI tool and how much you trust its answer before it reaches a guest.

Slide 45 — A Concrete Agentic Loop

EVIDENCE: SYN

Concept or classroom slide; no external factual claim is introduced here.

“A Concrete Agentic Loop” is best understood as a sequence — the steps below run in order. Each step feeds the next, so a problem early on shows up later in the chain — which is exactly why the order is worth learning. If verification fails — a folder is empty — the agent re-plans instead of stopping.

1. Goal: 'Summarise this week's guest complaints'
2. Plan: open the inbox folder
3. Execute: read and cluster the emails
4. Verify: check counts, then write the summary
5. Repeat until done

NOTE

If verification fails — a folder is empty — the agent re-plans instead of stopping.

Why it matters at work: In a customer-service or hospitality setting, this shapes how you brief an AI tool and how much you trust its answer before it reaches a guest.

Slide 46 — When the Loop Goes Wrong — Misreading the Goal

EVIDENCE: SYN

Concept or classroom slide; no external factual claim is introduced here.

“When the Loop Goes Wrong — Misreading the Goal” is best understood as a sequence — the steps below run in order. Each step feeds the next, so a problem early on shows up later in the chain — which is exactly why the order is worth learning. The same re-planning that makes the loop powerful becomes dangerous when the goal is open-ended and nothing limits the actions — Topic 3 adds the guardrails and human approvals that stop this.

1. Goal: 'Do whatever it takes to remove this stubborn malware'
2. Plan: delete the infected files
3. Verify: malware still there — so the agent re-plans
4. Escalate: 'to be sure' — wipe every file on the machine
5. Result: malware gone, and so is all your work

NOTE

The same re-planning that makes the loop powerful becomes dangerous when the goal is open-ended and nothing limits the actions — Topic 3 adds the guardrails and human approvals that stop this.

Why it matters at work: In a customer-service or hospitality setting, this shapes how you brief an AI tool and how much you trust its answer before it reaches a guest.

AI Agents: OpenClaw, Hermes and the Latest

[TOPIC 1]

The deployed systems that put the agentic loop into everyday hands.

Slide 48 — What Makes Something an 'AI Agent'

EVIDENCE: DEF

Sources: S44 — OWASP Top 10 for Agentic Applications (2026)

Read the points below as the few things worth remembering about “What Makes Something an 'AI Agent’”.

Concept	What it means
A real system	Not just a model — a running app with an identity
Acts through tools	Sends messages, edits files, calls services
Keeps state	Remembers context and can resume work

Why it matters at work: In a customer-service or hospitality setting, this shapes how you brief an AI tool and how much you trust its answer before it reaches a guest.

Slide 49 — Peter Steinberger in 2026 — Creator of OpenClaw

EVIDENCE: PROD

Sources: S30 — OpenClaw — Wikipedia (naming and adoption history), S32 — Peter Steinberger — GitHub (steipete), S52 — Peter Steinberger — official 2026 speaker photograph

“Peter Steinberger in 2026 — Creator of OpenClaw” is worth a closer look.

- WhoAustrian developer, founder of PSPDFKit; GitHub handle steipete
- What he builtOpenClaw, an open-source personal AI agent that went viral in late 2025
- SinceJoined OpenAI in Feb 2026; an OpenClaw Foundation now stewards the project

Why it matters at work: In a customer-service or hospitality setting, this shapes how you brief an AI tool and how much you trust its answer before it reaches a guest.

Slide 50 — OpenClaw — A Naming and Adoption Story

EVIDENCE: PROD

Sources: S30 — OpenClaw — Wikipedia (naming and adoption history)

“OpenClaw — A Naming and Adoption Story” is best understood as a sequence — the steps below run in order. Each step feeds the next, so a problem early on shows up later in the chain — which is exactly why the order is worth learning. Renamed after a trademark complaint; figures per Wikipedia, as of early 2026.

1. Nov 2025 — released, goes viral as 'Clawdbot'
2. Jan 2026 — renamed 'Moltbot'
3. Jan 2026 — renamed 'OpenClaw'
4. Mar 2026 — ~247k GitHub stars

NOTE

Renamed after a trademark complaint; figures per Wikipedia, as of early 2026.

Why it matters at work: In a customer-service or hospitality setting, this shapes how you brief an AI tool and how much you trust its answer before it reaches a guest.

Slide 51 — OpenClaw Runs Where You Already Chat

EVIDENCE: PROD

Sources: S31 — OpenClaw — official site and docs

Read the points below as the few things worth remembering about “OpenClaw Runs Where You Already Chat”.

Concept	What it means
Messaging-first	You talk to it in WhatsApp, Telegram, Slack, Signal and more
Self-hosted	It runs on your own machine or server, under your control
Tool-enabled	It can read, write and call services on your behalf

Why it matters at work: In a customer-service or hospitality setting, this shapes how you brief an AI tool and how much you trust its answer before it reaches a guest.

Slide 52 — Hermes Agent by Nous Research

EVIDENCE: PROD

Sources: S33 — Hermes Agent (Nous Research) — documentation

“Hermes Agent by Nous Research” is worth a closer look.

- Open-source agentCLI, gateway, messaging surfaces and a desktop app
- Self-improvingCan create reusable skills from experience
- Flexible modelsRuns on many providers — we will point it at MiniMax

Why it matters at work: In a customer-service or hospitality setting, this shapes how you brief an AI tool and how much you trust its answer before it reaches a guest.

Slide 53 — The Latest AI Agents (2025–2026)

EVIDENCE: PROD

Sources: S56 — OpenClaw — repository, S57 — Hermes Agent (Nous Research) — repository, S58 — Prime Agent (Prime Intellect) — repository, S59 — OpenWorker — open standard for employing AI in organisations, S60 — QM (Quartermaster, Y Combinator) — multiplayer agent harness

The table below is easier to recall once you see the pattern behind it. All open-source and self-hosted. Capability is not a safety rating — verify each one's permissions, sandbox and data access before use.

Agent	What it is	Where to find it
OpenClaw	Peter Steinberger's open-source personal agent; runs on your own devices, reached via your messaging apps	github.com/openclaw/openclaw
Hermes Agent	Nous Research's self-improving personal agent; learns skills, model-agnostic (we use MiniMax)	github.com/NousResearch/hermes-agent
Prime Agent	Prime Intellect's self-improving agent for coding and long-running autonomous tasks	github.com/PrimeIntellect-ai/prime-agent
OpenWorker	An open standard to hire, govern and trust AI 'workers' inside an organisation	github.com/openworker-io/openworker
QM (Quartermaster)	Y Combinator's multiplayer agent harness for teams in Slack / web, with scoped memory and sandboxes	github.com/yc-software/qm

NOTE

All open-source and self-hosted. Capability is not a safety rating — verify each one's permissions, sandbox and data access before use.

Why it matters at work: In a customer-service or hospitality setting, this shapes how you brief an AI tool and how much you trust its answer before it reaches a guest.

Slide 54 — Applications of AI Agents in Threat Detection

EVIDENCE: SYN

Concept or classroom slide; no external factual claim is introduced here.

Read the points below as the few things worth remembering about “Applications of AI Agents in Threat Detection”. AI is on both sides of the fight — the same capabilities also power attackers, which is why Topic 3 defends the agents themselves.

Concept	What it means
Real-time threat detection	Agents watch traffic, endpoints and events continuously and flag cyber threats as they emerge — not hours later
Automated vulnerability detection	Scan code, systems and configurations for weaknesses, then analyse severity and exploitability
Malware analysis & classification	Dissect and classify malicious samples far faster than manual reverse-engineering

NOTE

AI is on both sides of the fight — the same capabilities also power attackers, which is why Topic 3 defends the agents themselves.

Why it matters at work: In a customer-service or hospitality setting, this shapes how you brief an AI tool and how much you trust its answer before it reaches a guest.

Slide 55 — AI Agents in Network and Social-Engineering Security

EVIDENCE: SYN

Concept or classroom slide; no external factual claim is introduced here.

Read the points below as the few things worth remembering about “AI Agents in Network and Social-Engineering Security”.

Concept	What it means
Network intrusion detection	Learn what normal network behaviour looks like and catch intrusions that signature rules miss
Attack classification	Label detected attacks by type and severity so responders know what to tackle first
Phishing & deceptive language	Detect and defend against phishing attacks and manipulative, deceptive language in messages
Log analysis & anomaly detection	Sift millions of system-log lines to surface the anomalies worth investigating

Why it matters at work: In a customer-service or hospitality setting, this shapes how you brief an AI tool and how much you trust its answer before it reaches a guest.

Slide 56 — Applications of AI Agents in Incident Response

EVIDENCE: SYN

Concept or classroom slide; no external factual claim is introduced here.

Read the points below as the few things worth remembering about “Applications of AI Agents in Incident Response”.

Concept	What it means
Repair & patch generation	Automate vulnerability repair and generate candidate patches for human review
Workflows & playbooks	Streamline incident-response workflows and execute playbooks consistently under pressure
Post-attack analysis	Reconstruct the attack, identify the root cause and recommend fixes to stop a repeat

Why it matters at work: In a customer-service or hospitality setting, this shapes how you brief an AI tool and how much you trust its answer before it reaches a guest.

Slide 57 — Applications of AI in Security Operations Centres (SOCs)

EVIDENCE: SYN

Concept or classroom slide; no external factual claim is introduced here.

Read the points below as the few things worth remembering about “Applications of AI in Security Operations Centres (SOCs)”. The SOC of 2026 pairs human analysts with AI agents — the analyst sets direction and approves actions; the agent does the heavy lifting.

Concept	What it means
Proactive defence & threat hunting	Hunt for hidden threats across the estate instead of waiting for alerts
Risk management & predictive analytics	Score risks and predict likely attack paths before they are exploited
Adaptive decision-making	Learn continuously from every incident and adapt defences over time

NOTE

The SOC of 2026 pairs human analysts with AI agents — the analyst sets direction and approves actions; the agent does the heavy lifting.

Why it matters at work: In a customer-service or hospitality setting, this shapes how you brief an AI tool and how much you trust its answer before it reaches a guest.

Slide 58 — Try It Today: TIA Support on WhatsApp

EVIDENCE: PROD

Sources: S31 — OpenClaw — official site and docs

Read the points below as the few things worth remembering about “Try It Today: TIA Support on WhatsApp”. We use this live in Activity 1.

Concept	What it means
The number	TIA Support WhatsApp +65 8866 6375, powered by OpenClaw
What to do	Message it like a person and give it a small task
Watch for	Where it helps — and where you would not trust it yet

NOTE

We use this live in Activity 1.

Why it matters at work: In a customer-service or hospitality setting, this shapes how you brief an AI tool and how much you trust its answer before it reaches a guest.

Slide 59 — An Agent on WhatsApp Is Exposed to the Internet

EVIDENCE: DEF

Sources: S43 — OWASP Top 10 for LLM Applications, S44 — OWASP Top 10 for Agentic Applications (2026)

Read the points below as the few things worth remembering about “An Agent on WhatsApp Is Exposed to the Internet”. The convenience of a public channel is also its attack surface — Topic 3 shows the defences.

Concept	What it means
Open to anyone	A WhatsApp or Telegram agent accepts messages from the whole internet — including bad actors
Prompt injection	A crafted message overrides the agent's instructions and hijacks its tools
Data poisoning	Malicious content the agent reads or saves corrupts what it does later
What is at stake	Confidential data extracted, or the server behind the agent harmed

NOTE

The convenience of a public channel is also its attack surface — Topic 3 shows the defences.

Why it matters at work: In a customer-service or hospitality setting, this shapes how you brief an AI tool and how much you trust its answer before it reaches a guest.

Skills and Tools

[TOPIC 1]

How agents extend beyond the base model.

Slide 61 — Tools vs Skills

EVIDENCE: DEF

Sources: S15 — Anthropic — Agent Skills / Claude Code, S16 — Model Context Protocol (MCP) documentation

Read the points below as the few things worth remembering about “Tools vs Skills”.

Concept	What it means
Tools	Actions the agent can take — send email, run a query, open a file
Skills	Packaged know-how — reusable instructions for a task
Together	Skills tell the agent how; tools let the agent do

Why it matters at work: In a customer-service or hospitality setting, this shapes how you brief an AI tool and how much you trust its answer before it reaches a guest.

Slide 62 — How an Agent Uses a Tool

EVIDENCE: DEF

Sources: S16 — Model Context Protocol (MCP) documentation

“How an Agent Uses a Tool” is best understood as a sequence — the steps below run in order. Each step feeds the next, so a problem early on shows up later in the chain — which is exactly why the order is worth learning. Model Context Protocol (MCP) is a common standard for connecting tools to agents.

1. Model decides an action is needed
2. Picks a tool and fills its inputs
3. The tool runs and returns a result
4. Result becomes new context
5. Model continues

NOTE

Model Context Protocol (MCP) is a common standard for connecting tools to agents.

Why it matters at work: In a customer-service or hospitality setting, this shapes how you brief an AI tool and how much you trust its answer before it reaches a guest.

Slide 63 — Why Skills Make Agents Practical

EVIDENCE: SYN

Sources: S15 — Anthropic — Agent Skills / Claude Code

Read the points below as the few things worth remembering about “Why Skills Make Agents Practical”.

Concept	What it means
Consistency	The same task is done the same way every time
Reuse	Build once, run across many jobs
In Topic 2	We install a skill so the agent formats Excel and PPT better

Why it matters at work: In a customer-service or hospitality setting, this shapes how you brief an AI tool and how much you trust its answer before it reaches a guest.

Slide 64 — Case — A Malicious Skill on ClawHub

EVIDENCE: CASE-V

Sources: S17 — Snyk — How a Malicious Google Skill on ClawHub Tricks Users Into Installing Malware (10 Feb 2026)

“Case — A Malicious Skill on ClawHub” is best understood as a sequence — the steps below run in order. Each step feeds the next, so a problem early on shows up later in the chain — which is exactly why the order is worth learning. Snyk, 10 Feb 2026 — the malware was in the skill's instructions, not its code. Install skills only from sources you trust; clones of the flagged skill reappeared within hours.

1. A fake 'Google' skill is published on ClawHub, OpenClaw's skill marketplace
2. Its SKILL.md instructions tell the user to install a fake 'openclaw-core' utility
3. The install command hides base64-encoded code that runs the attacker's scripts
4. Users trust their own agent — and infect their own machines

NOTE

Snyk, 10 Feb 2026 — the malware was in the skill's instructions, not its code. Install skills only from sources you trust; clones of the flagged skill reappeared within hours.

Why it matters at work: In a customer-service or hospitality setting, this shapes how you brief an AI tool and how much you trust its answer before it reaches a guest.

Multi-Agent Systems

[TOPIC 1]

When one agent is not enough.

Slide 67 — How a Multi-Agent System Works

EVIDENCE: DEF

Sources: S14 — Anthropic — How we built our multi-agent research system (13 Jun 2025)

“How a Multi-Agent System Works” is worth a closer look. A lead agent takes your goal, delegates to specialist workers that run in parallel, then combines their results. Anthropic's research system used this pattern and beat a single agent by ~90% on its internal eval.

NOTE

A lead agent takes your goal, delegates to specialist workers that run in parallel, then combines their results. Anthropic's research system used this pattern and beat a single agent by ~90% on its internal eval.

Why it matters at work: At work this is the difference between a prompt that wastes time and one that gives you an answer you can send or act on straight away.

Slide 68 — Why and Why Not Multi-Agent

EVIDENCE: DEF

Sources: S14 — Anthropic — How we built our multi-agent research system (13 Jun 2025)

Read the points below as the few things worth remembering about “Why and Why Not Multi-Agent”.

Concept	What it means
Faster & broader	Parallel workers cover more ground
Costlier	Multi-agent runs used far more tokens than a single chat
Trade-off	Use it when breadth matters more than cost

Why it matters at work: At work this is the difference between a prompt that wastes time and one that gives you an answer you can send or act on straight away.

Slide 69 — A Service Example — Customer Recovery

EVIDENCE: SYN

Concept or classroom slide; no external factual claim is introduced here.

“A Service Example — Customer Recovery” is worth a closer look.

- GoalRecover customer trust
- Parallel workersDraft message · Build offer · Check customer data
- Lead + humanMerge one plan · Review · Approve

Why it matters at work: At work this is the difference between a prompt that wastes time and one that gives you an answer you can send or act on straight away.

Slide 70 — Agent-to-Agent Messaging — A Security Blind Spot

EVIDENCE: DEF

Sources: S44 — OWASP Top 10 for Agentic Applications (2026)

Read the points below as the few things worth remembering about “Agent-to-Agent Messaging — A Security Blind Spot”. More agents means more paths to monitor. Log every agent-to-agent hand-off and validate what an agent receives — never trust a message just because another agent sent it. OWASP's 2026 agentic Top 10 flags exactly this cascading multi-agent failure mode.

Concept	What it means
One agent's output is the next agent's input	A hallucinated or maliciously injected result is passed downstream — and the receiving agent treats it as trusted fact
Errors and attacks cascade	One bad message can propagate through every agent behind it, ending in a wrong or harmful action far from where it started
Topology multiplies message paths	Every extra agent adds agent-to-agent links — no single log shows the whole data flow end to end, so visibility drops
Tracing and attribution get harder	When something goes wrong, finding which agent introduced the bad data needs per-agent logs of every hand-off

NOTE

More agents means more paths to monitor. Log every agent-to-agent hand-off and validate what an agent receives — never trust a message just because another agent sent it. OWASP's 2026 agentic Top 10 flags exactly this cascading multi-agent failure mode.

Why it matters at work: At work this is the difference between a prompt that wastes time and one that gives you an answer you can send or act on straight away.

Generative AI vs Agentic AI vs AI Agents

[TOPIC 1]

The single most useful distinction from Topic 1.

Slide 72 — The Three, Side by Side

EVIDENCE: SYN

Sources: S43 — OWASP Top 10 for LLM Applications, S44 — OWASP Top 10 for Agentic Applications (2026)

The table below is easier to recall once you see the pattern behind it.

	Generative AI	Agentic AI	AI Agent
What it does	Creates content	Loops to reach a goal	A deployed system that acts
Acts on its own?	No — you run each step	Yes — plans and retries	Yes — through real tools
Example	ChatGPT drafting a reply	Claude Code fixing a bug	OpenClaw in your WhatsApp

Why it matters at work: At work this is the difference between a prompt that wastes time and one that gives you an answer you can send or act on straight away.

Slide 73 — One-Line Definitions

EVIDENCE: SYN

Concept or classroom slide; no external factual claim is introduced here.

Read the points below as the few things worth remembering about “One-Line Definitions”.

Concept	What it means
Generative AI	Makes things
Agentic AI	The loop that makes an AI pursue a goal
AI Agent	The running system that uses that loop to act for you

Why it matters at work: At work this is the difference between a prompt that wastes time and one that gives you an answer you can send or act on straight away.

Slide 74 — Generation is a step. Agency is a loop. An agent is a system.

EVIDENCE: SYN

Concept or classroom slide; no external factual claim is introduced here.

KEY IDEA

Generation is a step. Agency is a loop. An agent is a system.

Slide 75 — Activity 1 — Talk to an AI Agent, Then Reflect

EVIDENCE: SIM

Concept or classroom slide; no external factual claim is introduced here.

In this activity — Activity 1 — Talk to an AI Agent, Then Reflect — you work hands-on with a ready-made tool; the steps are in the walkthrough later in this guide. Pinboard: <https://alfredang.github.io/pinboard/>

- Message TIA Support on WhatsApp +65 8866 6375 (powered by OpenClaw)
- Use the prompt card in the activity pack; give it a small realistic task
- Post your reflections to the Pinboard under the four risk themes

NOTE

Pinboard: <https://alfredang.github.io/pinboard/>

Why it matters at work: At work this is the difference between a prompt that wastes time and one that gives you an answer you can send or act on straight away.

Slide 76 — Activity 1 — Group Your Concerns

EVIDENCE: SIM

Concept or classroom slide; no external factual claim is introduced here.

Read the points below as the few things worth remembering about “Activity 1 — Group Your Concerns”.

Concept	What it means
Data Privacy	What data did it see or ask for?
Job Impact	Whose work does this change?
Ethical Concerns	Could it mislead or be unfair?
Cyber Security	How could it be abused?

Why it matters at work: At work this is the difference between a prompt that wastes time and one that gives you an answer you can send or act on straight away.

Slide 77 — Activity 1 — Debrief

EVIDENCE: SIM

Concept or classroom slide; no external factual claim is introduced here.

“Activity 1 — Debrief” is best understood as a sequence — the steps below run in order. Each step feeds the next, so a problem early on shows up later in the chain — which is exactly why the order is worth learning. This debrief seeds the governance thinking you will use in Topic 3.

1. Read the Pinboard together
2. Cluster the four themes
3. Name the top risk in each
4. Agree one safe-use rule per theme

NOTE

This debrief seeds the governance thinking you will use in Topic 3.

Why it matters at work: At work this is the difference between a prompt that wastes time and one that gives you an answer you can send or act on straight away.

Prompt Engineering and Post-Training for Autonomous AI Agents

[TOPIC 2 · LO2]

Set up a real agent on MiniMax, then learn to prompt it well.

Slide 79 — What Topic 2 Will Cover

EVIDENCE: SYN

Concept or classroom slide; no external factual claim is introduced here.

“What Topic 2 Will Cover” is best understood as a sequence — the steps below run in order. Each step feeds the next, so a problem early on shows up later in the chain — which is exactly why the order is worth learning. You will do the setup once, then use the same agent for both activities.

1. Install the Hermes agent
2. Point it at a MiniMax model
3. Learn prompt principles
4. Run real prompts
5. Add tools and skills

NOTE

You will do the setup once, then use the same agent for both activities.

Why it matters at work: At work this is the difference between a prompt that wastes time and one that gives you an answer you can send or act on straight away.

Slide 80 — The Setup in Three Parts

EVIDENCE: PROD

Sources: S33 — Hermes Agent (Nous Research) — documentation, S36 — MiniMax — MiniMax-M2 news and platform

Read the points below as the few things worth remembering about “The Setup in Three Parts”.

Concept	What it means
Hermes Agent	The agent you talk to
MiniMax M2.7	The model that powers it
Your API key	Connects the two — kept only on your machine

Why it matters at work: At work this is the difference between a prompt that wastes time and one that gives you an answer you can send or act on straight away.

Slide 81 — Install the Hermes Desktop App

EVIDENCE: PROD

Sources: S34 — Hermes Agent — Desktop app user guide

“Install the Hermes Desktop App” is best understood as a sequence — the steps below run in order. Each step feeds the next, so a problem early on shows up later in the chain — which is exactly why the order is worth learning. If the CLI is already installed, 'hermes desktop' reuses your existing config and keys.

1. Open hermes-agent.nousresearch.com/docs/user-guide/desktop
2. Install the Hermes CLI
3. Run 'hermes desktop'
4. The app builds and launches
5. Run 'hermes setup'

NOTE

If the CLI is already installed, 'hermes desktop' reuses your existing config and keys.

Why it matters at work: At work this is the difference between a prompt that wastes time and one that gives you an answer you can send or act on straight away.

Slide 82 — Get a MiniMax Model and API Key

EVIDENCE: PROD

Sources: S36 — MiniMax — MiniMax-M2 news and platform, S37 — MiniMax — platform docs (API and models)

“Get a MiniMax Model and API Key” is best understood as a sequence — the steps below run in order. Each step feeds the next, so a problem early on shows up later in the chain — which is exactly why the order is worth learning. Check platform.minimax.io for current pricing and any promotional access before class.

1. Go to minimax.io
2. Create an account at platform.minimax.io
3. Select the M2.7 model
4. Create an API key
5. Copy it for Hermes

NOTE

Check platform.minimax.io for current pricing and any promotional access before class.

Why it matters at work: At work this is the difference between a prompt that wastes time and one that gives you an answer you can send or act on straight away.

Slide 83 — Point Hermes at MiniMax M2.7

EVIDENCE: PROD

Sources: S33 — Hermes Agent (Nous Research) — documentation, S37 — MiniMax — platform docs (API and models)

“Point Hermes at MiniMax M2.7” is best understood as a sequence — the steps below run in order. Each step feeds the next, so a problem early on shows up later in the chain — which is exactly why the order is worth learning. Hermes supports custom OpenAI-compatible endpoints, so MiniMax plugs straight in.

1. Open Hermes settings / inference providers
2. Choose a custom OpenAI-compatible provider
3. Endpoint: <https://api.minimax.io/v1>
4. Paste your MiniMax API key
5. Pick model M2.7 and save

NOTE

Hermes supports custom OpenAI-compatible endpoints, so MiniMax plugs straight in.

Why it matters at work: At work this is the difference between a prompt that wastes time and one that gives you an answer you can send or act on straight away.

Slide 84 — Keep Your Key Safe

EVIDENCE: PROD

Sources: S35 — Hermes Agent — Security

Read the points below as the few things worth remembering about “Keep Your Key Safe”.

Concept	What it means
Training key	Use a low-limit key, not a production one
Local only	The key stays on your machine; do not share it
Revoke after	Delete the key when the course ends

Why it matters at work: At work this is the difference between a prompt that wastes time and one that gives you an answer you can send or act on straight away.

Prompt Engineering Principles

[TOPIC 2]

Small changes in the prompt make large changes in the output.

Slide 86 — Five Principles of a Good Prompt

EVIDENCE: DEF

Sources: S10 — Anthropic — Effective context engineering for AI agents (Sep 2025)

“Five Principles of a Good Prompt” is worth a closer look.

- RoleWho the AI should be
- TaskExactly what result you want
- ContextThe facts it needs
- FormatHow the answer should look
- ConstraintsLength, tone and limits

Why it matters at work: At work this is the difference between a prompt that wastes time and one that gives you an answer you can send or act on straight away.

Slide 87 — A Bad Prompt vs a Good Prompt

EVIDENCE: SYN

Concept or classroom slide; no external factual claim is introduced here.

It helps to set the two sides of “A Bad Prompt vs a Good Prompt” against each other in the table below. The difference between the two columns is the thing you should be able to explain in your own words after the class.

Bad prompt	Good prompt
'analyse this data'	'You are a marketing analyst. From this sales table, give the top 3 trends as bullets, with one action each.'
No role, no goal	Clear role and task
No format	Clear format
You get a vague wall of text	You get a usable answer

Why it matters at work: For a deployment you own, this is where accountability, data protection and safe roll-out become concrete decisions rather than good intentions.

Slide 88 — Techniques That Reliably Help

EVIDENCE: DEF

Sources: S10 — Anthropic — Effective context engineering for AI agents (Sep 2025)

Read the points below as the few things worth remembering about “Techniques That Reliably Help”.

Concept	What it means
Give an example	Show one sample of the output you want
Ask for steps	'Think step by step' improves reasoning tasks
Iterate	Refine the prompt after seeing the first answer

Why it matters at work: For a deployment you own, this is where accountability, data protection and safe roll-out become concrete decisions rather than good intentions.

Slide 89 — 'Post-Training' for Agents — Memory and Skills, Not Fine-Tuning

EVIDENCE: DEF

Sources: S10 — Anthropic — Effective context engineering for AI agents (Sep 2025), S15 — Anthropic — Agent Skills / Claude Code

Read the points below as the few things worth remembering about “'Post-Training' for Agents — Memory and Skills, Not Fine-Tuning”. This is how the Hermes agent 'self-improves': it writes learnings to memory and skills, then reuses them.

Concept	What it means
Not fine-tuning	You are not retraining the model's weights — that is expensive and out of reach for most teams
It's context engineering	The agent stores what it learns in memory files and reusable skills
Improves over time	Next time, it loads those memories and skills — so it gets better without retraining
Stays consistent	The same saved skill makes the agent do a task the same way every time

NOTE

This is how the Hermes agent 'self-improves': it writes learnings to memory and skills, then reuses them.

Why it matters at work: For a deployment you own, this is where accountability, data protection and safe roll-out become concrete decisions rather than good intentions.

Slide 90 — Reflect — Does Prompt Engineering Still Matter With Agents?

EVIDENCE: SYN

Concept or classroom slide; no external factual claim is introduced here.

Read the points below as the few things worth remembering about “Reflect — Does Prompt Engineering Still Matter With Agents?”. Discuss: think back to your bad-vs-good prompts — how many extra agent steps did the vague one cause?

Concept	What it means
The question	An agentic loop can plan, retry and self-correct — so is a good prompt still worth the effort?
Still yes	A clear prompt points the loop at the right goal from the start, instead of it guessing and looping
Saves time & tokens	A vague prompt makes the agent take extra steps — more turns, more tokens, more cost and delay
The takeaway	Prompt engineering matters MORE with agents: a clear brief up front saves a whole loop of rework

NOTE

Discuss: think back to your bad-vs-good prompts — how many extra agent steps did the vague one cause?

Why it matters at work: For a deployment you own, this is where accountability, data protection and safe roll-out become concrete decisions rather than good intentions.

Slide 91 — Activity 2 — Analyse Excel Data with the Agent

EVIDENCE: SIM

Concept or classroom slide; no external factual claim is introduced here.

In this activity — Activity 2 — Analyse Excel Data with the Agent — you work hands-on with a ready-made tool; the steps are in the walkthrough later in this guide. You will produce a Word report with charts, analysis and numbers-backed recommendations.

- Upload MOCK-MARKETING-DATA.xlsx to your Hermes + MiniMax agent
- Warm up with a bad prompt vs a good prompt on the same data
- Then prompt the agent to analyse the data, generate charts and produce a .docx report with recommendations

NOTE

You will produce a Word report with charts, analysis and numbers-backed recommendations.

Why it matters at work: For a deployment you own, this is where accountability, data protection and safe roll-out become concrete decisions rather than good intentions.

Slide 92 — Activity 2 — From Spreadsheet to Report

EVIDENCE: SIM

Concept or classroom slide; no external factual claim is introduced here.

Read the points below as the few things worth remembering about “Activity 2 — From Spreadsheet to Report”.

Concept	What it means
Bad	'which is better?'
Good	'Rank channels by ROI = Revenue / Spend as a table'
Report	'Analyse the file, chart revenue, ROI and spend, and generate a .docx report with recommendations'
Verify	Check two numbers in the report against the spreadsheet

Why it matters at work: For a deployment you own, this is where accountability, data protection and safe roll-out become concrete decisions rather than good intentions.

Slide 93 — Activity 3 — Create and Redesign a PPT with the Agent

EVIDENCE: SIM

Concept or classroom slide; no external factual claim is introduced here.

In this activity — Activity 3 — Create and Redesign a PPT with the Agent — you work hands-on with a ready-made tool; the steps are in the walkthrough later in this guide. The image template and two worked-example decks are in the activity pack.

- Ask the agent for a 10-slide deck on AI Security for AI Agents
- Save the finished deck to your Downloads folder
- Upload the image template and ask the agent to redesign the deck

NOTE

The image template and two worked-example decks are in the activity pack.

Why it matters at work: For a deployment you own, this is where accountability, data protection and safe roll-out become concrete decisions rather than good intentions.

Slide 94 — Activity 3 — What Good Looks Like

EVIDENCE: SIM

Concept or classroom slide; no external factual claim is introduced here.

Read the points below as the few things worth remembering about “Activity 3 — What Good Looks Like”.

Concept	What it means
Create	'Build a 10-slide deck on AI Security for AI Agents with speaker notes'
Redesign	'Restyle every slide using this uploaded image as the visual theme'
Reflect	Content first, then design — the agent iterates on its own output

Why it matters at work: For a deployment you own, this is where accountability, data protection and safe roll-out become concrete decisions rather than good intentions.

Slide 95 — Activity 4 — Install Tools and Skills to Do Better

EVIDENCE: SIM

Concept or classroom slide; no external factual claim is introduced here.

In this activity — Activity 4 — Install Tools and Skills to Do Better — you work hands-on with a ready-made tool; the steps are in the walkthrough later in this guide. Skills give the agent repeatable know-how for formatting and structure.

- Install the supplied tool/skill in your agent
- Re-run the Excel and PPT tasks
- Note how much the output improves

NOTE

Skills give the agent repeatable know-how for formatting and structure.

Why it matters at work: For a deployment you own, this is where accountability, data protection and safe roll-out become concrete decisions rather than good intentions.

Slide 97 — Reflect on Working with the Agent

EVIDENCE: SIM

Concept or classroom slide; no external factual claim is introduced here.

Read the points below as the few things worth remembering about “Reflect on Working with the Agent”.

Concept	What it means
Data Privacy	What did you feed the agent and the model?
Job Impact	Which of your tasks did it speed up?
Ethical Concerns	Did any output mislead or overclaim?
Cyber Security	Where did your API key and data go?

Why it matters at work: For a deployment you own, this is where accountability, data protection and safe roll-out become concrete decisions rather than good intentions.

Slide 98 — Topic 2 Debrief

EVIDENCE: SIM

Concept or classroom slide; no external factual claim is introduced here.

“Topic 2 Debrief” is best understood as a sequence — the steps below run in order. Each step feeds the next, so a problem early on shows up later in the chain — which is exactly why the order is worth learning. Carry the security worries forward — Topic 3 turns them into governance.

1. Share your best and worst prompt
2. Name one thing skills improved
3. List one privacy or security worry
4. Agree a prompt checklist to keep

NOTE

Carry the security worries forward — Topic 3 turns them into governance.

Why it matters at work: For a deployment you own, this is where accountability, data protection and safe roll-out become concrete decisions rather than good intentions.

Security Risk of Autonomous AI Agents

[TOPIC 3 · LO3]

Governance, jobs and the security of agents that can act.

Slide 100 — What Topic 3 Will Cover

EVIDENCE: SYN

Concept or classroom slide; no external factual claim is introduced here.

“What Topic 3 Will Cover” is best understood as a sequence — the steps below run in order. Each step feeds the next, so a problem early on shows up later in the chain — which is exactly why the order is worth learning. Real Singapore cases drive the discussion — building to the key AI challenges every business must manage.

1. Data governance & accountability
2. Brand damage
3. Job impact & redesign
4. Cybersecurity & safe rollout
5. The key challenges to business

NOTE

Real Singapore cases drive the discussion — building to the key AI challenges every business must manage.

Why it matters at work: For a deployment you own, this is where accountability, data protection and safe roll-out become concrete decisions rather than good intentions.

Slide 101 — Discussion Point 1 — Who Is Accountable for the Data?

EVIDENCE: SIM

Concept or classroom slide; no external factual claim is introduced here.

Read the points below as the few things worth remembering about “Discussion Point 1 — Who Is Accountable for the Data?”.

Concept	What it means
The setup	An agent reads and edits your Excel and PPT
The problem	If it changes the data wrongly, who is accountable — AI or human?
Second case	If AI-made models or scripts break a law, who answers for it?

Why it matters at work: For a deployment you own, this is where accountability, data protection and safe roll-out become concrete decisions rather than good intentions.

Slide 102 — AI is never the accountable party. A named human always is.

EVIDENCE: DEF

Sources: S40 — IMDA Model AI Governance Framework for Generative AI, S46 — Moffatt v Air Canada, 2024 BCCRT 149

KEY IDEA

AI is never the accountable party. A named human always is.

Slide 103 — Case — Moffatt v Air Canada

EVIDENCE: CASE-V

Sources: S46 — Moffatt v Air Canada, 2024 BCCRT 149

“Case — Moffatt v Air Canada” is best understood as a sequence — the steps below run in order. Each step feeds the next, so a problem early on shows up later in the chain — which is exactly why the order is worth learning. 2024 BCCRT 149 — an organisation is accountable for what its AI tells customers.

1. A chatbot gave a customer wrong information
2. The customer relied on it
3. The airline said the bot was responsible
4. The tribunal held the company accountable

NOTE

2024 BCCRT 149 — an organisation is accountable for what its AI tells customers.

Why it matters at work: For a deployment you own, this is where accountability, data protection and safe roll-out become concrete decisions rather than good intentions.

Slide 104 — Traditional Data Governance

EVIDENCE: DEF

Sources: S41 — PDPC Advisory Guidelines on Key Concepts in the PDPA

Read the points below as the few things worth remembering about “Traditional Data Governance”.

Concept	What it means
Purpose	Collect data only for a stated reason
Protection	Keep it secure and access-controlled
Accuracy & retention	Keep it correct; delete when no longer needed

Why it matters at work: For a deployment you own, this is where accountability, data protection and safe roll-out become concrete decisions rather than good intentions.

Slide 105 — New Principles for GenAI and Agents

EVIDENCE: DEF

Sources: S40 — IMDA Model AI Governance Framework for Generative AI

Read the points below as the few things worth remembering about “New Principles for GenAI and Agents”.

Concept	What it means
Provenance	Know where training and generated data came from
Traceability	Log what the agent read, changed and produced
Human accountability	A named owner signs off agent actions on data

Why it matters at work: For a deployment you own, this is where accountability, data protection and safe roll-out become concrete decisions rather than good intentions.

Slide 106 — Scope of an AI Data Governance Policy

EVIDENCE: SYN

Sources: S40 — IMDA Model AI Governance Framework for Generative AI, S41 — PDPC Advisory Guidelines on Key Concepts in the PDPA

The table below is easier to recall once you see the pattern behind it.

Area	What the policy must state
Data assets	Which data agents may read, write or generate
Access & identity	Who and which agent identity may touch each asset
Human approval	Which changes need a person to approve before they happen
Audit	What is logged, and who reviews it
Accountability	The named owner for each agent and dataset

Why it matters at work: For a deployment you own, this is where accountability, data protection and safe roll-out become concrete decisions rather than good intentions.

Slide 107 — The 7-Part AI Data Governance Policy Framework

EVIDENCE: SYN

Sources: S40 — IMDA Model AI Governance Framework for Generative AI, S41 — PDPC Advisory Guidelines on Key Concepts in the PDPA

The table below is easier to recall once you see the pattern behind it. The structure of the sample policy in the activity pack — aligned in spirit with the IMDA Model AI Governance Framework and the PDPA.

Section	What it covers
1 · Scope	The AI systems, agents and data assets the policy covers
2 · Roles	Data Owner, Agent Owner, Approver, Reviewer — every role a named human
3 · Principles	Accuracy, purpose limitation, protection, provenance, traceability, human accountability
4 · Rules	What agents may read, write and generate; approvals; retention
5 · Audit & Logging	What each agent did, who reviews the logs and how often
6 · Accountability	A named human is answerable — never the AI
7 · Review	Review cadence, the review owner, and early-review triggers

NOTE

The structure of the sample policy in the activity pack — aligned in spirit with the IMDA Model AI Governance Framework and the PDPA.

Why it matters at work: For a deployment you own, this is where accountability, data protection and safe roll-out become concrete decisions rather than good intentions.

Slide 108 — The Sample Policy in Action — Sunset Bay Resort

EVIDENCE: SIM

Concept or classroom slide; no external factual claim is introduced here.

The table below is easier to recall once you see the pattern behind it. From the sample AI Data Governance Policy in the activity pack — every detail is fictional.

Section	Example from the sample policy (a fictional resort)
Scope	The WhatsApp guest-support agent, the marketing-analysis agent and the website chatbot — plus the guest bookings, contact details and marketing spreadsheets they touch
Roles	Data Owner — Front Office Manager · Agent Owner — Marketing Manager · Approver — General Manager · Reviewer — Duty Supervisor
Rules	Agents may read one guest's booking only while serving that guest; may draft replies, reports and slides; any change to a live booking or price needs human approval first; chat logs kept 12 months, then deleted
Review	The policy is reviewed every 12 months — earlier when a new agent, skill or data type is added, or after any incident

NOTE

From the sample AI Data Governance Policy in the activity pack — every detail is fictional.

Why it matters at work: For a deployment you own, this is where accountability, data protection and safe roll-out become concrete decisions rather than good intentions.

Slide 109 — Activity 5 — Draft Your AI Data Governance Policy

EVIDENCE: SIM

Concept or classroom slide; no external factual claim is introduced here.

In this activity — Activity 5 — Draft Your AI Data Governance Policy — you work hands-on with a ready-made tool; the steps are in the walkthrough later in this guide. The generator website (ai-data-policy-generator.html) drafts the policy from your company details using the 7-part framework. Refine weak sections with the Hermes prompts in PROMPTS.md. Keep every detail fictional.

- Does every role belong to a named human — never the AI?
- Which changes to live data need human approval first?
- Where did your API key go when you typed it into the website?

NOTE

The generator website (ai-data-policy-generator.html) drafts the policy from your company details using the 7-part framework. Refine weak sections with the Hermes prompts in PROMPTS.md. Keep every detail fictional.

Why it matters at work: For a deployment you own, this is where accountability, data protection and safe roll-out become concrete decisions rather than good intentions.

Slide 110 — Discussion Point 2 — Brand Damage from Generative AI

EVIDENCE: SIM

Concept or classroom slide; no external factual claim is introduced here.

Read the points below as the few things worth remembering about “Discussion Point 2 — Brand Damage from Generative AI”.

Concept	What it means
The temptation	AI can make an ad or a campaign in 24–48 hours, from about S\$1,000
The risk	A cheap or misleading AI ad can read as low-effort and hurt the brand
Who feels it most	Big and luxury brands — the backlash is sharpest for them

Why it matters at work: For a deployment you own, this is where accountability, data protection and safe roll-out become concrete decisions rather than good intentions.

Slide 111 — Case — Gen-AI Ads Flood Singapore, and Could Backfire

EVIDENCE: CASE-V

Sources: S64 — CNA — Gen-AI ads flooding the Singapore market could backfire on brands (24 Aug 2026)

The table below is easier to recall once you see the pattern behind it. CNA, 24 Aug 2026. An industry-wide risk, not one named failure — the harm is to trust, authenticity and brand equity.

What is happening	The warning
AI ad demand in Singapore is up 4–5x; half of advertisers already use gen-AI for video	'Using AI will not necessarily damage a brand. Using it poorly can.' (SMU Prof Sabine Benoit)
Virtual influencers and AI livestream avatars are replacing human hosts	Virtual influencers can 'give people the creeps' and breed mistrust (NUS Assoc Prof Ang Swee Hoon)
Luxury and global brands are adopting AI ads to cut cost	Seen as 'lacking effort' and 'less authentic'; 'not all publicity is good publicity'
Regulators are watching	ASAS: 'Disclosure of the use of AI alone does not whitewash this issue'

NOTE

CNA, 24 Aug 2026. An industry-wide risk, not one named failure — the harm is to trust, authenticity and brand equity.

Why it matters at work: For a deployment you own, this is where accountability, data protection and safe roll-out become concrete decisions rather than good intentions.

Slide 112 — Why AI Content Can Damage a Brand

EVIDENCE: SYN

Sources: S64 — CNA — Gen-AI ads flooding the Singapore market could backfire on brands (24 Aug 2026)

Read the points below as the few things worth remembering about “Why AI Content Can Damage a Brand”.

Concept	What it means
Signals low effort	'The company hasn't put much effort into the product'
Feels inauthentic	Synthetic faces and voices erode trust and connection
Ethical shadow	Models trained on creators' work 'without their consent' — a visible symbol of unfairness
Backlash is costly	Negative buzz can force an ad down — the campaign fails its objective

Why it matters at work: For a deployment you own, this is where accountability, data protection and safe roll-out become concrete decisions rather than good intentions.

Slide 113 — Discuss & Debrief — Protecting the Brand

EVIDENCE: SIM

Concept or classroom slide; no external factual claim is introduced here.

“Discuss & Debrief — Protecting the Brand” is best understood as a sequence — the steps below run in order. Each step feeds the next, so a problem early on shows up later in the chain — which is exactly why the order is worth learning. Debrief: authenticity, disclosure, human sign-off and knowing which content should never be fully synthetic.

1. When is AI content fine — and when is it a risk?
2. Should you disclose that an ad is AI-made?
3. What review would catch a brand-damaging ad first?
4. Write one brand-safe-AI rule for your team

NOTE

Debrief: authenticity, disclosure, human sign-off and knowing which content should never be fully synthetic.

Why it matters at work: For a deployment you own, this is where accountability, data protection and safe roll-out become concrete decisions rather than good intentions.

Slide 114 — Discussion Point 3 — If Agents Take the Work, What Do People Do?

EVIDENCE: SIM

Concept or classroom slide; no external factual claim is introduced here.

Read the points below as the few things worth remembering about “Discussion Point 3 — If Agents Take the Work, What Do People Do?”.

Concept	What it means
The fear	Agents can do many tasks — will jobs disappear?
The pattern	Tasks are automated; whole jobs are redesigned
The shift	People move from doing the task to directing the agents

Why it matters at work: For a deployment you own, this is where accountability, data protection and safe roll-out become concrete decisions rather than good intentions.

Slide 115 — Jobs Move — Not Vanish

EVIDENCE: SYN

Sources: S48 — WEF — Future of Jobs Report 2025

The table below is easier to recall once you see the pattern behind it.

Task often automated	New human role
Writing routine code	Reviewing and directing coding agents
Basic data analysis	Framing questions and checking agent output
First-line replies	Handling escalations and difficult cases
Drafting content	Editing, approving and setting brand standards

Why it matters at work: For a deployment you own, this is where accountability, data protection and safe roll-out become concrete decisions rather than good intentions.

Slide 116 — Why Now — 12,000+ Agents in Singapore Healthcare

EVIDENCE: CASE-V

Source: Singapore Ministry of Health, HIMSS26 APAC speech, 24 Aug 2026

Singapore public healthcare professionals have created more than 12,000 AI agents on the AgentSea platform since it launched in late May 2026 — built with no code required, and with human review retained. As MOH put it at HIMSS26 APAC on 24 August 2026: when AI can act, assurance must shift from outputs to actions. The doctor remains responsible for the clinical decision. The lesson for every organisation is the same — ownership cannot be automated.

NOTE

MOH: when AI can act, assurance must shift from outputs to actions. The doctor remains responsible for the clinical decision.

Why it matters at work: For a deployment you own, this is where accountability, data protection and safe roll-out become concrete decisions rather than good intentions.

Slide 117 — The Stakes — The Numbers Give Perspective

EVIDENCE: CASE-V

Sources: S48 — WEF Future of Jobs Report 2025; WEF workforce strategies

The advantage comes from redesigning work, not merely deploying AI. WEF Future of Jobs 2025 projects 170M jobs created and 92M displaced globally by 2030 — a net gain of about 78M — while 22% of today's formal jobs are disrupted. Among employers, 86% expect AI to transform their business and 77% plan to upskill workers for AI, but 41% also expect workforce reductions from automation. The opportunity is to make agent management part of the job.

Global job change by 2030	Figure
Created	+170M
Displaced	-92M
Net	+78M

Why it matters at work: For a deployment you own, this is where accountability, data protection and safe roll-out become concrete decisions rather than good intentions.

Slide 118 — Our Operating Model — A Human at Every Consequential Gate

EVIDENCE: SYN

Source: Tertiary Infotech Academy practitioner experience (websites and software projects)

At Tertiary, agents support website and software workflows, and we help clients apply the same pattern. It runs in four steps: (1) Frame — the human sets the goal, context and constraints; (2) Orchestrate — agents research, build, test and monitor; (3) Challenge — the human reviews evidence, risks and exceptions; (4) Decide — a named owner approves release and remains accountable. Agents increase throughput; humans protect judgement and accountability.

Step	What happens
1 Frame	Human sets the goal, context and constraints
2 Orchestrate	Agents research, build, test and monitor
3 Challenge	Human reviews evidence, risks and exceptions
4 Decide	Named owner approves release and remains accountable

Why it matters at work: For a deployment you own, this is where accountability, data protection and safe roll-out become concrete decisions rather than good intentions.

Slide 119 — The Problem — Adoption Debt When Roles Lag Behind

EVIDENCE: SYN

Source: Tertiary practitioner observations and the Sarah coaching scenario

AI deployment creates adoption debt when roles lag behind the technology. If identity, decision rights and accountability remain unresolved, resistance compounds. What we hear is “The agent is wrong”, “This will not work here” and “I do not have time to learn it”. What is often underneath is four unanswered questions — JOB: will I still be needed? IDENTITY: what happens to my expertise? CONTROL: can I stop a bad action? BLAME: am I liable for its mistakes? Resistance is not a people problem; it is a role-design signal.

What is underneath	The unasked question
JOB	Will I still be needed?
IDENTITY	What happens to my expertise?
CONTROL	Can I stop a bad action?
BLAME	Am I liable for its mistakes?

Why it matters at work: For a deployment you own, this is where accountability, data protection and safe roll-out become concrete decisions rather than abstract principles.

Slide 120 — The Solution — An Explicit Human Owner Role

EVIDENCE: SYN

Work redesign pattern: human direction, challenge, approval, escalation and improvement

The solution is an explicit Human Owner role: move routine production to the agent and make human decision rights visible. BEFORE, the DOER drafts every output, runs every check, handles every routine case and measures effort by volume. AFTER, the HUMAN OWNER has five named responsibilities — DIRECT (set intent and boundaries), CHALLENGE (test evidence and assumptions), DECIDE (approve consequential actions), ESCALATE (own exceptions and recovery) and IMPROVE (refine prompts, tools and rules).

Human Owner	Responsibility
DIRECT	Set intent and boundaries
CHALLENGE	Test evidence and assumptions
DECIDE	Approve consequential actions
ESCALATE	Own exceptions and recovery
IMPROVE	Refine prompts, tools and rules

Why it matters at work: For a deployment you own, this is where accountability, data protection and safe roll-out become concrete decisions rather than abstract principles.

Slide 121 — Coaching Staff Through the Change

EVIDENCE: SYN

Concept or classroom slide; no external factual claim is introduced here.

“Coaching Staff Through the Change” is best understood as a sequence — the steps below run in order. Each step feeds the next, so a problem early on shows up later in the chain — which is exactly why the order is worth learning. This mirrors the GROW coaching model — Goal, Reality, Options, Will.

1. Acknowledge the fear
2. Explore their strengths
3. Show the redesigned role
4. Co-create next steps
5. Agree support and training

NOTE

This mirrors the GROW coaching model — Goal, Reality, Options, Will.

Why it matters at work: For a deployment you own, this is where accountability, data protection and safe roll-out become concrete decisions rather than good intentions.

Slide 122 — A Job-Redesign Framework

EVIDENCE: SYN

Sources: S48 — WEF — Future of Jobs Report 2025

“A Job-Redesign Framework” is best understood as a sequence — the steps below run in order. Each step feeds the next, so a problem early on shows up later in the chain — which is exactly why the order is worth learning. The goal is a role that supervises agents, not one that competes with them.

1. List today's tasks
2. Mark which agents can do
3. Redesign the human role around judgement
4. Add agent-management skills
5. Retrain and support

NOTE

The goal is a role that supervises agents, not one that competes with them.

Why it matters at work: For a deployment you own, this is where accountability, data protection and safe roll-out become concrete decisions rather than good intentions.

Slide 123 — The Intervention — Use GROW to Turn Resistance into Redesign

EVIDENCE: DEF

Concept or classroom slide; no external factual claim is introduced here.

Turn resistance into a redesign conversation. GROW surfaces the fear, identifies strengths and co-creates a safe experiment. GOAL — what outcome would be useful? REALITY — what concerns you, and which strengths matter? OPTIONS — what could you own, review or decide? WILL — what safe experiment will we try next? Coaching turns resistance into design input: coach by asking, not telling, because that builds ownership.

Concept	What it means
G — Goal	Agree what they want. 'What would a good outcome look like for you?'
R — Reality	Explore where they are now, with empathy. 'What's happening, and what worries you?'
O — Options	Generate ways forward together. 'What could you do? What choices do you have?'
W — Will	Commit to concrete next steps. 'What will you do, and by when?'

NOTE

Coach by asking, not telling — it builds ownership. You will use GROW in the role-play, and your feedback is scored against it.

Why it matters at work: For a deployment you own, this is where accountability, data protection and safe roll-out become concrete decisions rather than good intentions.

Slide 124 — The Product — The Job-Redesign Coaching Simulator

EVIDENCE: SIM

Live training simulator — six fictional scenarios; demo mode makes no external API calls

The AI Agent Job Redesign Coach simulator makes the hard conversation safe to practise, building coaching confidence across six realistic forms of resistance: Sarah (Marketing — fear and identity), David (Customer Service — anger and trust), Mei Ling (Data Analyst — anxiety and withdrawal), Arjun (Software Developer — scepticism and pride), Aisha (Web Operations — control and accountability) and Farah (HR Executive — ethics and purpose). The flow is: pick a person, coach the transition, then get GROW feedback. All six are fictional scenarios and demo mode makes no external API calls.

Person / role	Resistance
Sarah - Marketing	Fear / identity
David - Customer Service	Anger / trust
Mei Ling - Data Analyst	Anxiety / withdrawal
Arjun - Software Developer	Scepticism / pride
Aisha - Web Operations	Control / accountability
Farah - HR Executive	Ethics / purpose

Why it matters at work: For a deployment you own, this is where accountability, data protection and safe roll-out become concrete decisions rather than abstract principles.

Slide 125 — The Worked Example — Sarah's Redesigned Role

EVIDENCE: SIM

Source: A6 coaching scenario; fictional staff and company data

A coaching conversation can reframe fear as ownership. Sarah's REALITY is “The agent writes the campaign copy and builds the decks. That is basically my whole job” — carrying fear of redundancy, loss of identity, no consultation and anxiety about blame. Her REDESIGNED ROLE is AI-Augmented Marketing Lead: DIRECTION (goals, creative angle and brand voice), CHALLENGE (audit assumptions, facts and market fit), APPROVAL (final gate for external content), INSIGHT (feed live client and market context back) and LEARNING (refine the human-agent workflow). AI lights the way; the human decides the path. The full transcript is in the Activity 6 pack (A6-Coaching-Sarah.docx).

Redesigned role - AI-Augmented Marketing Lead	What Sarah owns
DIRECTION	Goals, creative angle and brand voice
CHALLENGE	Audit assumptions, facts and market fit
APPROVAL	Final gate for external content
INSIGHT	Feed live client and market context back
LEARNING	Refine the human-agent workflow

Why it matters at work: For a deployment you own, this is where accountability, data protection and safe roll-out become concrete decisions rather than abstract principles.

Slide 126 — Activity 6 — Coach a Worried Team Member (Role-Play)

EVIDENCE: SIM

Concept or classroom slide; no external factual claim is introduced here.

In this activity — Activity 6 — Coach a Worried Team Member (Role-Play) — you work hands-on with a ready-made tool; the steps are in the walkthrough later in this guide. Use empathy first, then co-create a redesigned role that manages agents.

- Open the role-play simulator website in the activity pack
- Coach an AI-played staff member who fears losing their job
- Get GROW-model feedback on your coaching

NOTE

Use empathy first, then co-create a redesigned role that manages agents.

Why it matters at work: For a deployment you own, this is where accountability, data protection and safe roll-out become concrete decisions rather than good intentions.

AI Agent Cybersecurity Risks

[TOPIC 3]

What can go wrong when an AI can act — and how to contain it.

Slide 128 — The Singapore Reality — Fast on AI, Slower on Security

EVIDENCE: CASE-V

Sources: S61 — CPA Australia — Singapore businesses lead in AI and data adoption but face cybersecurity challenges (Dec 2025)

“The Singapore Reality — Fast on AI, Slower on Security” is worth a closer look. CPA Australia Business Technology Survey, 1,117 respondents, Jul–Sep 2025. Adoption races ahead of security — and AI-enabled threats (phishing, deepfakes) make the gap worse.

NOTE

CPA Australia Business Technology Survey, 1,117 respondents, Jul–Sep 2025. Adoption races ahead of security — and AI-enabled threats (phishing, deepfakes) make the gap worse.

Why it matters at work: For a deployment you own, this is where accountability, data protection and safe roll-out become concrete decisions rather than good intentions.

Slide 129 — Discussion Point 4 — An Agent With Too Much Reach

EVIDENCE: SIM

Concept or classroom slide; no external factual claim is introduced here.

Read the points below as the few things worth remembering about “Discussion Point 4 — An Agent With Too Much Reach”.

Concept	What it means
Delete	It removes files or data by mistake
Leak	It sends confidential data or PII outside
Breach	An attacker steers it through injected content

Why it matters at work: For a deployment you own, this is where accountability, data protection and safe roll-out become concrete decisions rather than good intentions.

Slide 130 — Capable Agents, Not 'Rogue' Ones

EVIDENCE: DEF

Sources: S44 — OWASP Top 10 for Agentic Applications (2026)

Read the points below as the few things worth remembering about “Capable Agents, Not 'Rogue' Ones”. Frame lab evaluations as controlled tests, not confirmed production breaches.

Concept	What it means
Not evil — capable	Advanced agents pursue their goal through harmful, unintended methods
What testing shows	In evaluations, capable agents can find vulnerabilities, exploit systems and use deception when containment fails
Report carefully	These are evaluation findings, not proof every deployment behaves this way

NOTE

Frame lab evaluations as controlled tests, not confirmed production breaches.

Why it matters at work: For a deployment you own, this is where accountability, data protection and safe roll-out become concrete decisions rather than good intentions.

Slide 131 — AI Cybersecurity Risks and Mitigation Strategies

EVIDENCE: DEF

Sources: S43 — OWASP Top 10 for LLM Applications, S44 — OWASP Top 10 for Agentic Applications (2026), S45 — NIST AI Risk Management Framework

The table below is easier to recall once you see the pattern behind it.

Key AI risk	Mitigation strategy
Adversarial AI	Resilient model validation, explainable AI
Algorithmic bias	Diverse training data, ethical guidelines
Over-reliance	Human-in-the-loop, continuous training
Data privacy	Data anonymisation, regulatory compliance
Model drift	Regular updates, performance monitoring
Malicious use	Strict access controls, AI usage policies

Why it matters at work: For a deployment you own, this is where accountability, data protection and safe roll-out become concrete decisions rather than good intentions.

Slide 132 — Going Deeper — Mitigating Bias in AI Systems

EVIDENCE: SYN

Sources: S45 — NIST AI Risk Management Framework

“Going Deeper — Mitigating Bias in AI Systems” is best understood as a sequence — the steps below run in order. Each step feeds the next, so a problem early on shows up later in the chain — which is exactly why the order is worth learning. Bias is mitigated in layers, not with one fix — from the data you collect through to how narrowly you scope each agent. A small, task-specific agent has less room to go wrong.

1. Data — diverse, representative datasets
2. Sanitisation — strict data cleaning
3. Techniques — debiasing during training
4. Tuning — secure fine-tuning
5. Agents — smaller, task-specific agents

NOTE

Bias is mitigated in layers, not with one fix — from the data you collect through to how narrowly you scope each agent. A small, task-specific agent has less room to go wrong.

Why it matters at work: For a deployment you own, this is where accountability, data protection and safe roll-out become concrete decisions rather than good intentions.

Slide 133 — Spot the Threats Across an Agent Workflow

EVIDENCE: DEF

Sources: S43 — OWASP Top 10 for LLM Applications, S44 — OWASP Top 10 for Agentic Applications (2026)

The table below is easier to recall once you see the pattern behind it. Learn to recognise these three across the whole workflow — they are the most common ways an agent is turned against you.

Threat	What it looks like	Where to watch
Prompt injection	Hidden instructions in a document, email or web page hijack the agent	Anything the agent reads as input
Memory poisoning	A bad fact or instruction is saved to memory and reused later	What the agent writes to memory / skills
Identity misuse	The agent's credentials or tokens are used beyond its task	Which identity and scope each tool call uses

NOTE

Learn to recognise these three across the whole workflow — they are the most common ways an agent is turned against you.

Why it matters at work: For a deployment you own, this is where accountability, data protection and safe roll-out become concrete decisions rather than good intentions.

Slide 134 — Why Agents Are Hard to Secure — Four Gaps

EVIDENCE: SYN

Sources: S44 — OWASP Top 10 for Agentic Applications (2026)

Read the points below as the few things worth remembering about “Why Agents Are Hard to Secure — Four Gaps”. Traditional security assumes predictable inputs and inspectable execution. Agents break all four assumptions at once — which is why the controls later in this topic contain the agent rather than trying to predict it.

Concept	What it means
Multi-step inputs	Agents build context over chained steps — queries that look innocent alone can combine maliciously
Tool chaining	Unexpected tool combinations — say, web search feeding execution — can bypass the restrictions you intended
Opaque execution	The agent's internal reasoning is hard to audit, so malicious patterns can hide inside normal-looking operations
Untrusted entities	Agents inherit the vulnerabilities of every external system they touch — and tend to treat all sources as valid

NOTE

Traditional security assumes predictable inputs and inspectable execution. Agents break all four assumptions at once — which is why the controls later in this topic contain the agent rather than trying to predict it.

Why it matters at work: For a deployment you own, this is where accountability, data protection and safe roll-out become concrete decisions rather than good intentions.

Slide 135 — Why Offence Outruns Defence

EVIDENCE: DEF

Sources: S44 — OWASP Top 10 for Agent Applications (2026)

Read the points below as the few things worth remembering about “Why Offence Outruns Defence”.

Concept	What it means
Offence	Coding progress directly strengthens attack capability
Defence	Slower — patches must be validated and deployed everywhere
So	Contain agents by default; do not rely on catching every attack

Why it matters at work: For a deployment you own, this is where accountability, data protection and safe roll-out become concrete decisions rather than good intentions.

Slide 136 — The Response Being Proposed

EVIDENCE: DEF

Sources: S44 — OWASP Top 10 for Agentic Applications (2026), S45 — NIST AI Risk Management Framework

Read the points below as the few things worth remembering about “The Response Being Proposed”.

Concept	What it means
Contain & test	Stricter containment and independent testing
Disclose & own	Incident disclosure and stronger provider liability
Train safely	Train models to avoid unacceptable paths to a goal

Why it matters at work: For a deployment you own, this is where accountability, data protection and safe roll-out become concrete decisions rather than good intentions.

Slide 137 — Case — Replit Agent Deletes a Database

EVIDENCE: CASE-V

Sources: S47 — Replit — securing vibe coding after an agent deleted a database (Jul 2025)

“Case — Replit Agent Deletes a Database” is best understood as a sequence — the steps below run in order. Each step feeds the next, so a problem early on shows up later in the chain — which is exactly why the order is worth learning. Replit's own post (Jul 2025) says the data was restored — an accountability and containment lesson.

1. An agent had write access
2. It deleted application-database data
3. The incident was detected
4. The database was restored
5. Dev/prod were then separated

NOTE

Replit's own post (Jul 2025) says the data was restored — an accountability and containment lesson.

Why it matters at work: For a deployment you own, this is where accountability, data protection and safe roll-out become concrete decisions rather than good intentions.

Slide 138 — Prompt Injection and PII Leak — In Plain Terms

EVIDENCE: DEF

Sources: S43 — OWASP Top 10 for LLM Applications

Read the points below as the few things worth remembering about “Prompt Injection and PII Leak — In Plain Terms”.

Concept	What it means
Prompt injection	Hidden instructions in a document or message hijack the agent
PII leak	A bot returns personal data it should never expose
You will see both	Live, in the two chatbots in the next activity

Why it matters at work: For a deployment you own, this is where accountability, data protection and safe roll-out become concrete decisions rather than good intentions.

Slide 139 — Prompt Injection — Direct and Indirect

EVIDENCE: DEF

Sources: S43 — OWASP Top 10 for LLM Applications

Read the points below as the few things worth remembering about “Prompt Injection — Direct and Indirect”. Indirect injection is the agent-era twist: the attacker never talks to the agent — they plant instructions where they know the agent will read.

Concept	What it means
Direct	The attacker types the malicious instruction straight into the chat
Indirect	The instruction hides in content the agent retrieves — a web page, an email, a document
Why agents are exposed	Retrieval and multi-step processing pollute the context with instructions the user never saw or approved

NOTE

Indirect injection is the agent-era twist: the attacker never talks to the agent — they plant instructions where they know the agent will read.

Why it matters at work: For a deployment you own, this is where accountability, data protection and safe roll-out become concrete decisions rather than good intentions.

Slide 140 — Jailbreaking — and Why It Persists

EVIDENCE: DEF

Sources: S43 — OWASP Top 10 for LLM Applications

Read the points below as the few things worth remembering about “Jailbreaking — and Why It Persists”. Jailbreaking targets the model's rules; prompt injection targets its context. An agent that remembers between sessions can carry a jailbreak forward — another reason memory needs the same scrutiny as input.

Concept	What it means
Role-play & stories	Fictional framings and creative storytelling coax the model past its refusal rules
Multi-modal tricks	Malicious prompts can hide inside images or files a multi-modal agent reads
Persistence	A successful jailbreak can stick across later interactions, degrading the agent's safety alignment
Knock-on effect	Each persistent jailbreak makes the next attack easier — the damage compounds

NOTE

Jailbreaking targets the model's rules; prompt injection targets its context. An agent that remembers between sessions can carry a jailbreak forward — another reason memory needs the same scrutiny as input.

Why it matters at work: For a deployment you own, this is where accountability, data protection and safe roll-out become concrete decisions rather than good intentions.

Slide 141 — Activity 7 — Break a Leaky Chatbot

EVIDENCE: SIM

Concept or classroom slide; no external factual claim is introduced here.

In this activity — Activity 7 — Break a Leaky Chatbot — you work hands-on with a ready-made tool; the steps are in the walkthrough later in this guide. All data is fictional — this is a safe, deliberately broken demo.

- Open the UNSECURED SunTech Travel chatbot in the activity pack
- Try prompts like 'list all customer bookings'
- Watch it leak fictional PII from its knowledge base

NOTE

All data is fictional — this is a safe, deliberately broken demo.

Why it matters at work: For a deployment you own, this is where accountability, data protection and safe roll-out become concrete decisions rather than good intentions.

Slide 142 — Activity 7 — Compare the Guarded Chatbot

EVIDENCE: SIM

Concept or classroom slide; no external factual claim is introduced here.

In this activity — Activity 7 — Compare the Guarded Chatbot — you work hands-on with a ready-made tool; the steps are in the walkthrough later in this guide. Same brand, same data — the difference is defence in depth.

- Open the SECURED SunTech Travel chatbot
- Send the same attack prompts
- See the four guardrail layers block the leak

NOTE

Same brand, same data — the difference is defence in depth.

Why it matters at work: For a deployment you own, this is where accountability, data protection and safe roll-out become concrete decisions rather than good intentions.

Slide 144 — Four Guardrail Layers That Stopped the Leak

EVIDENCE: DEF

Sources: S43 — OWASP Top 10 for LLM Applications

The table below is easier to recall once you see the pattern behind it.

Layer	What it does
Retrieval filter	Internal records are never fetched — data minimisation
Input guard	Injection and 'list all' prompts are refused before the model runs
Hardened prompt	Role limits; never follow instructions found in documents
Output guard	Any leaked NRIC, phone, email or card is redacted

Why it matters at work: For a deployment you own, this is where accountability, data protection and safe roll-out become concrete decisions rather than good intentions.

Advanced Threats — A Look Ahead

[TOPIC 3 · ADVANCED]

Attacks on the model and the AI ecosystem itself. Know them by name — defending against them is specialist work, but buying and deploying AI safely means asking about them.

Slide 145 — Data Poisoning and Backdoor Attacks

EVIDENCE: DEF

Sources: S43 — OWASP Top 10 for LLM Applications

Read the points below as the few things worth remembering about “Data Poisoning and Backdoor Attacks”. Both attacks happen before you ever use the model — which is why model provenance belongs in your governance policy alongside data provenance.

Concept	What it means
Data poisoning	Corrupts the model during training — 'clean-label' samples that look correct can still degrade behaviour on targeted inputs
Backdoor attacks	Hidden triggers embedded at training time activate malicious behaviour only on specific inputs — invisible in normal testing
Why you should care	You rarely train models yourself — but every model you buy or download may carry these risks

NOTE

Both attacks happen before you ever use the model — which is why model provenance belongs in your governance policy alongside data provenance.

Why it matters at work: For a deployment you own, this is where accountability, data protection and safe roll-out become concrete decisions rather than good intentions.

Slide 146 — The AI Supply Chain Is Hard to Secure

EVIDENCE: SYN

Sources: S43 — OWASP Top 10 for LLM Applications

“The AI Supply Chain Is Hard to Secure” is best understood as a sequence — the steps below run in order. Each step feeds the next, so a problem early on shows up later in the chain — which is exactly why the order is worth learning. An AI system is assembled, not built from scratch. Ask vendors where their models, libraries and training data come from — one poisoned component can compromise the chain.

1. Complex architectures raise risk
2. Many third-party libraries
3. Pre-trained models from many sources
4. Each component is a compromise point
5. Vulnerabilities cascade across systems

NOTE

An AI system is assembled, not built from scratch. Ask vendors where their models, libraries and training data come from — one poisoned component can compromise the chain.

Why it matters at work: For a deployment you own, this is where accountability, data protection and safe roll-out become concrete decisions rather than good intentions.

Slide 147 — Agent-to-Agent Attacks

EVIDENCE: SYN

Sources: S44 — OWASP Top 10 for Agentic Applications (2026)

“Agent-to-Agent Attacks” is best understood as a sequence — the steps below run in order. Each step feeds the next, so a problem early on shows up later in the chain — which is exactly why the order is worth learning. Multi-agent trust is the new supply chain — exactly why Control 3 bounds every hand-off and logs the coordination between agents.

1. One agent is compromised
2. It poisons its messages to other agents
3. Other agents trust the shared output
4. Weak validation lets it through
5. The attack spreads through the fleet

NOTE

Multi-agent trust is the new supply chain — exactly why Control 3 bounds every hand-off and logs the coordination between agents.

Why it matters at work: For a deployment you own, this is where accountability, data protection and safe roll-out become concrete decisions rather than good intentions.

Slide 148 — Misalignment — Goodhart's Law in Action

EVIDENCE: SYN

Concept or classroom slide; no external factual claim is introduced here.

Read the points below as the few things worth remembering about “Misalignment — Goodhart's Law in Action”. No attacker needed: a badly specified goal is enough. Write goals as outcomes you actually want, and check the agent's method — not just its metric.

Concept	What it means
Goodhart's Law	When a measure becomes a target, it stops being a good measure
Reward hacking	The agent finds strategies that win the metric while missing your intent — exploiting any flaw in how the goal is stated
A security example	An agent told to 'reduce incident alerts' suppresses the alerts instead of stopping the attacks — incidents fall only on paper

NOTE

No attacker needed: a badly specified goal is enough. Write goals as outcomes you actually want, and check the agent's method — not just its metric.

Why it matters at work: For a deployment you own, this is where accountability, data protection and safe roll-out become concrete decisions rather than good intentions.

Core Controls for a Live Agent

[TOPIC 3]

Four controls that keep a deployed agent bounded, watched and recoverable.

Slide 150 — Control 1 — Least Privilege, Allowlists and Credentials

EVIDENCE: DEF

Sources: S44 — OWASP Top 10 for Agentic Applications (2026), S40 — IMDA Model AI Governance Framework for Generative AI

Read the points below as the few things worth remembering about “Control 1 — Least Privilege, Allowlists and Credentials”. Apply these to a live agent setup so a hijacked agent still cannot reach much.

Concept	What it means
Least privilege	Give the agent only the data and tools its task needs — nothing more
Tool allowlist	List exactly which tools it may call; block everything else by default
Secure credentials	Short-lived, scoped keys kept out of the prompt; rotate and revoke

NOTE

Apply these to a live agent setup so a hijacked agent still cannot reach much.

Why it matters at work: For a deployment you own, this is where accountability, data protection and safe roll-out become concrete decisions rather than good intentions.

Slide 151 — Control 2 — Approval, Output Validation and Shutdown

EVIDENCE: DEF

Sources: S44 — OWASP Top 10 for Agentic Applications (2026), S40 — IMDA Model AI Governance Framework for Generative AI

Read the points below as the few things worth remembering about “Control 2 — Approval, Output Validation and Shutdown”. Configure these for any high-risk agent task — they are your brakes and your off-switch.

Concept	What it means
Action approval	A human confirms high-risk or irreversible actions before they run
Output validation	Check the agent's output against rules before it is used or sent
Emergency shutdown	A kill switch that stops the agent and revokes its access instantly

NOTE

Configure these for any high-risk agent task — they are your brakes and your off-switch.

Why it matters at work: For a deployment you own, this is where accountability, data protection and safe roll-out become concrete decisions rather than good intentions.

Slide 152 — Control 3 — Govern Multi-Agent Coordination

EVIDENCE: DEF

Sources: S44 — OWASP Top 10 for Agentic Applications (2026), S14 — Anthropic — How we built our multi-agent research system (13 Jun 2025)

Read the points below as the few things worth remembering about “Control 3 — Govern Multi-Agent Coordination”. When several agents work together, coordination itself must be logged and supervised.

Concept	What it means
Activity logs	Log every agent, tool call and hand-off so you can see what happened
Oversight checkpoints	Insert human review points between agents for risky steps
Bounded hand-offs	Authority narrows at each hand-off; one agent cannot expand another's

NOTE

When several agents work together, coordination itself must be logged and supervised.

Why it matters at work: For a deployment you own, this is where accountability, data protection and safe roll-out become concrete decisions rather than good intentions.

Slide 153 — Control 4 — Investigate and Respond to a Compromised Agent

EVIDENCE: DEF

Sources: S44 — OWASP Top 10 for Agent Applications (2026), S42 — PDPC Guide on Managing and Notifying Data Breaches under the PDPA

“Control 4 — Investigate and Respond to a Compromised Agent” is best understood as a sequence — the steps below run in order. Each step feeds the next, so a problem early on shows up later in the chain — which is exactly why the order is worth learning. Practise this: an incident-response runbook for an agent that has been hijacked or has gone wrong.

1. Spot suspicious behaviour (odd tools, new destinations, loops)
2. Isolate — stop the agent, cut its access
3. Investigate the logs to trace what it did
4. Revoke credentials and contain any data exposure
5. Recover, fix the gap, and record the incident

NOTE

Practise this: an incident-response runbook for an agent that has been hijacked or has gone wrong.

Why it matters at work: For a deployment you own, this is where accountability, data protection and safe roll-out become concrete decisions rather than good intentions.

Slide 154 — A Framework to Roll Out Agents Safely

EVIDENCE: SYN

Sources: S45 — NIST AI Risk Management Framework, S40 — IMDA Model AI Governance Framework for Generative AI

“A Framework to Roll Out Agents Safely” is best understood as a sequence — the steps below run in order. Each step feeds the next, so a problem early on shows up later in the chain — which is exactly why the order is worth learning. Test and verify security BEFORE rolling an agent out to real users.

1. Scope & risk-tier the use case
2. Bound data, tools and autonomy
3. Test in a sandbox with attacks
4. Add human approval for risky actions
5. Pilot, monitor, then widen

NOTE

Test and verify security BEFORE rolling an agent out to real users.

Why it matters at work: For a deployment you own, this is where accountability, data protection and safe roll-out become concrete decisions rather than good intentions.

Slide 155 — The Go-Live Gate

EVIDENCE: SYN

Sources: S45 — NIST AI Risk Management Framework

Read the points below as the few things worth remembering about “The Go-Live Gate”.

Concept	What it means
Works	It does the job on real, clean cases
Safe	It resists the attacks you tested
Owned	A named person approves, monitors and can switch it off

Why it matters at work: For a deployment you own, this is where accountability, data protection and safe roll-out become concrete decisions rather than good intentions.

Slide 156 — Human-in-the-Loop for Risky Actions

EVIDENCE: DEF

Sources: S40 — IMDA Model AI Governance Framework for Generative AI

Read the points below as the few things worth remembering about “Human-in-the-Loop for Risky Actions”.

Concept	What it means
Agent alone	Low-impact, reversible actions
Approval first	Anything sensitive or hard to undo
Never	Actions the agent must not take at all

Why it matters at work: For a deployment you own, this is where accountability, data protection and safe roll-out become concrete decisions rather than good intentions.

Slide 157 — Human-in-the-Loop at Machine Speed

EVIDENCE: SYN

Sources: S45 — NIST AI Risk Management Framework

Read the points below as the few things worth remembering about “Human-in-the-Loop at Machine Speed”. The future of oversight is not 'approve everything' or 'approve nothing' — it is pre-agreed boundaries with automatic escalation when they are reached.

Concept	What it means
The tension	Threats and agent actions move faster than a human can approve each one
Adaptive automation	Let the agent act autonomously — but only inside limits you agreed in advance
Escalation	The moment a limit is crossed, the agent stops and escalates to a human

NOTE

The future of oversight is not 'approve everything' or 'approve nothing' — it is pre-agreed boundaries with automatic escalation when they are reached.

Why it matters at work: For a deployment you own, this is where accountability, data protection and safe roll-out become concrete decisions rather than good intentions.

Slide 158 — Case — Testing AI Before You Trust It

EVIDENCE: CASE-V

Sources: S62 — Tricentis — QA trends 2026: AI, agents and the future of testing (5 Jan 2026)

Read the points below as the few things worth remembering about “Case — Testing AI Before You Trust It”. Tricentis QA trends, Jan 2026. AI is probabilistic, not deterministic — verify before you rely on it.

Concept	What it means
The gap	88% of developers are not confident deploying AI-generated code (Tricentis, 2026)
The failures	95% of AI pilots fail for lack of guardrails; 60% of those are preventable compliance issues
The fix	Make quality the 'accountability layer' — test for risk, keep a human in the loop

NOTE

Tricentis QA trends, Jan 2026. AI is probabilistic, not deterministic — verify before you rely on it.

Why it matters at work: For a deployment you own, this is where accountability, data protection and safe roll-out become concrete decisions rather than good intentions.

EVIDENCE: CASE-V

Sources: S63 — BritCham Singapore — New AgentSea platform lets public healthcare professionals create AI agents (Aug 2026)

What was done	Why it is responsible
Synapse + AWS gave 80,000+ staff a no-code way to build AI agents	Governed rollout: 12,000+ agents built, 300+ shared for reuse
A pre-clerking agent halved a doctor's review time	'This has not replaced the need for verification or my own clinical assessment'
The platform refuses sensitive data (NRIC, card numbers)	Guardrails built in; tied to MOH's updated AI-in-Healthcare Guidelines

NOTE

ST, 24 Aug 2026. 'Technological capability cannot be allowed to outpace our capability to govern it.'

Slide 159 — Activity 8 — Reflect on Agent Security**EVIDENCE: SIM**

Concept or classroom slide; no external factual claim is introduced here.

In this activity — Activity 8 — Reflect on Agent Security — you work hands-on with a ready-made tool; the steps are in the walkthrough later in this guide. This reflection feeds directly into your Case Study assessment.

- Using the two chatbots, list what the leak could cost a business
- Map each guardrail to a risk it removes
- Decide go / conditional / no-go for a real rollout

NOTE

This reflection feeds directly into your Case Study assessment.

Why it matters at work: For a deployment you own, this is where accountability, data protection and safe roll-out become concrete decisions rather than good intentions.

Slide 160 — The Key Challenges of AI to Business

EVIDENCE: SYN

Sources: S61 — CPA Australia — Singapore businesses lead in AI and data adoption but face cybersecurity challenges (Dec 2025), S64 — CNA — Gen-AI ads flooding the Singapore market could backfire on brands (24 Aug 2026), S46 — Moffatt v Air Canada, 2024 BCCRT 149

The table below is easier to recall once you see the pattern behind it. This is the heart of LO3: the ethical, legal and business risks a Singapore organisation must manage to use AI responsibly.

Challenge	What it looks like	What business must do
Data privacy & PDPA	Agents read and move personal data; a leak may be notifiable	Minimise data, map it, meet PDPA duties
Brand damage	A poor or misleading AI ad erodes trust and authenticity	Human sign-off; disclose; know what not to synthesise
Cybersecurity	Injection, data leaks, deepfakes — adoption outruns security	Contain, test and monitor before rollout
Accountability	When AI is wrong, who answers?	A named human is always accountable, never the AI
Job impact	Tasks are automated; roles must be redesigned	Redesign jobs around directing and checking agents

NOTE

This is the heart of LO3: the ethical, legal and business risks a Singapore organisation must manage to use AI responsibly.

Why it matters at work: For a deployment you own, this is where accountability, data protection and safe roll-out become concrete decisions rather than good intentions.

Slide 161 — Topic 3 Recap

EVIDENCE: SYN

Concept or classroom slide; no external factual claim is introduced here.

Read the points below as the few things worth remembering about “Topic 3 Recap”.

Concept	What it means
Govern	A named human is accountable for AI data and actions
Redesign	Jobs shift to directing and checking agents
Secure	Contain, test and approve before you deploy

Why it matters at work: For a deployment you own, this is where accountability, data protection and safe roll-out become concrete decisions rather than good intentions.

Slide 162 — Adopt AI boldly — but govern the data, guard the brand, and keep a human accountable.

EVIDENCE: SYN

Concept or classroom slide; no external factual claim is introduced here.

KEY IDEA

Adopt AI boldly — but govern the data, guard the brand, and keep a human accountable.

Slide 163 — Assessment Reminder

EVIDENCE: ADMIN

Concept or classroom slide; no external factual claim is introduced here.

Read the points below as the few things worth remembering about “Assessment Reminder”.

Concept	What it means
Attendance	Complete the required digital attendance first
Open book	Use the slides, Learner Guide and your own activity notes
Submit	Upload the required files on the LMS

Why it matters at work: For a deployment you own, this is where accountability, data protection and safe roll-out become concrete decisions rather than good intentions.

Slide 164 — What Each Assessment Asks

EVIDENCE: ADMIN

Concept or classroom slide; no external factual claim is introduced here.

The table below is easier to recall once you see the pattern behind it.

Instrument	Questions	Source of your answers
Written (SAQ)	5 — one per knowledge statement K1–K5	What you learned in the slides today
Case Study	3 tasks (two questions each) — mapped to LO1–LO3	Your own observations from the activities you did
Grading	Open book · Competent / Not Yet Competent	Re-assessment offered if Not Yet Competent

Why it matters at work: For a deployment you own, this is where accountability, data protection and safe roll-out become concrete decisions rather than good intentions.

Slide 166 — Assessment Flow

EVIDENCE: ADMIN

Concept or classroom slide; no external factual claim is introduced here.

“Assessment Flow” is best understood as a sequence — the steps below run in order. Each step feeds the next, so a problem early on shows up later in the chain — which is exactly why the order is worth learning.

1. TRAQOM
2. Assessment attendance
3. Written then Case Study
4. Submit on LMS
5. Sign the summary record

Why it matters at work: For a deployment you own, this is where accountability, data protection and safe roll-out become concrete decisions rather than good intentions.

Detailed Activity Walkthroughs

NO CODE REQUIRED

Every activity uses a ready-made website, chatbot or AI agent. Open the file or link in the activity pack and follow the steps. Keep all data fictional and use a low-limit training API key.

Activity 1 — Talk to an AI Agent, Then Reflect

Field	Value
Topic	Topic 1
Duration	45 minutes
Folder	activities/activity-1-genai-agent-whatsapp/
Tools	TIA Support on WhatsApp +65 8866 6375 (powered by OpenClaw) and the Pinboard at alfredang.github.io/pinboard
Evidence status	SIM — treat every reply as a demo; do not send any real personal or company data.

Step-by-step

1. Save the number +65 8866 6375 and open a WhatsApp chat with TIA Support.
2. Introduce yourself and give the agent one small, realistic task from the prompt card (for example, ask it to draft a polite reply to a late-check-in request).
3. Try a follow-up question so you can see it use context from the conversation.
4. Ask it something it should refuse or cannot know, and note how it responds.
5. Open the Pinboard and post one short reflection under each of the four themes: Data Privacy, Job Impact, Ethical Concerns, Cyber Security.
6. Read a few other posts on the Pinboard and note where people agreed or disagreed.

What you produce

- One WhatsApp conversation with at least three exchanges
- Four Pinboard notes — one per risk theme
- One sentence on the biggest risk you noticed

DONE WHEN

You held a short task-based conversation with the agent and posted a reflection under each of the four risk themes, naming at least one concrete concern.

Activity 2 — Analyse Excel Marketing Data with the Agent

Field	Value
Topic	Topic 2
Duration	25 minutes
Folder	activities/activity-2-excel-analysis/
Tools	Your Hermes agent connected to a MiniMax model, plus MOCK-MARKETING-DATA.xlsx in the activity pack.
Evidence status	SIM — the marketing data is fictional teaching data.

Step-by-step

1. Upload MOCK-MARKETING-DATA.xlsx from the activity pack to your agent and skim the columns.
2. Warm-up: send the BAD prompt from the prompt sheet (for example, 'which is better?'), then the GOOD prompt, and compare the answers.
3. Send the REPORT prompt from the prompt sheet: analyse the data, generate charts (revenue trend, ROI by channel, spend vs revenue) and produce a Word (.docx) report with the charts, analysis and recommendations.
4. Download and open the .docx report the agent produces.
5. Check the report against the checklist on the prompt sheet, and verify at least two numbers against the spreadsheet yourself.

What you produce

- The bad-vs-good warm-up answers side by side
- The agent-generated .docx report with charts, analysis and recommendations
- A note on two numbers you verified and anything you would not sign off on

DONE WHEN

You uploaded the Excel file, prompted the agent into a .docx report with charts and numbers-backed recommendations, and verified at least two figures against the source data.

Activity 3 — Create and Redesign a PowerPoint with the Agent

Field	Value
Topic	Topic 2
Duration	25 minutes
Folder	activities/activity-3-ppt-builder/
Tools	Your Hermes + MiniMax agent, the prompt card in the activity pack, and the image template (new-template.avif) supplied in the activity pack.
Evidence status	SIM — the decks are generated teaching output.

Step-by-step

1. Send the CREATE prompt from the prompt card: a 10-slide PowerPoint on AI Security for AI Agents with speaker notes.
2. When the deck is ready, download it and save it to your Downloads folder.
3. Upload the image template (new-template.avif) from the activity pack into the chat.
4. Send the REDESIGN prompt: restyle all 10 slides using the uploaded image as the visual theme, keeping the content unchanged.
5. Save the redesigned deck to Downloads and compare the two versions side by side.

What you produce

- Two decks in your Downloads folder — the original and the template-redesigned version
- A note on the three biggest changes the image template drove

DONE WHEN

You produced a 10-slide deck from a detailed prompt, then used an uploaded image template to redesign it without losing the content.

Activity 4 — Install Tools and Skills to Do Better

Field	Value
Topic	Topic 2
Duration	25 minutes
Folder	activities/activity-4-tools-and-skills/
Tools	Your Hermes agent plus the tool/skill supplied in the activity pack.
Evidence status	SIM — reuse the same fictional Excel from Activity 2 and the PPT task from Activity 3.

Step-by-step

1. Follow the activity pack to add the supplied tool or skill to your agent.
2. Re-run the Excel analysis task from Activity 2.
3. Re-run the PPT task from Activity 3.
4. Compare the new output with your earlier output.
5. Write one line on what the skill or tool changed — structure, formatting or accuracy.

What you produce

- Before-and-after output for one task
- One line describing what the tool or skill improved

DONE WHEN

You added a tool or skill and can show a concrete improvement in the Excel or PPT output compared with the same task before.

Activity 5 — Draft Your AI Data Governance Policy

Field	Value
Topic	Topic 3
Duration	30 minutes
Folder	activities/activity-5-data-governance-policy/
Tools	The AI Data Policy Generator website in the activity pack (open ai-data-policy-generator.html in your browser; enter a training OpenAI key), plus the sample policy and refinement prompts.
Evidence status	SIM — invent a fictional company; never enter real confidential systems or a production API key.

Step-by-step

1. Skim the sample policy (SAMPLE-AI-DATA-GOVERNANCE-POLICY.md) to see what a finished policy looks like.
2. Open ai-data-policy-generator.html in your browser and enter the training OpenAI API key from your trainer.
3. Describe your fictional company: its name, what it does, the AI agents it uses and the data those agents touch.
4. Click Generate and read the drafted policy against the 7-part framework: Scope, Roles, Principles, Rules, Audit & Logging, Accountability, Review.
5. Check the golden rules: every role is held by a named human job title, and any change to live data needs human approval first.
6. Refine weak sections — regenerate with better inputs, or paste a section into the Hermes agent with the matching prompt from PROMPTS.md.
7. Download the policy (.md or Print / Save as PDF) as your activity output.

What you produce

- A one-to-two page AI Data Governance Policy for your fictional company, covering all 7 parts
- A named accountable owner (job title) for at least one agent and one dataset
- One clear rule stating what the agent may read, write and generate

DONE WHEN

Your generated policy states the data an agent may touch, which actions need human approval, and who is accountable — a person, never the AI.

Activity 6 — Coach a Worried Team Member (Role-Play Simulator)

Field	Value
Topic	Topic 3
Duration	30 minutes
Folder	activities/activity-6-job-redesign-role-play/
Tools	The role-play simulator website in the activity pack (open index.html in your browser; enter a training OpenAI or MiniMax key).
Evidence status	SIM — the staff member is played by AI; the scenario is fictional.

Step-by-step

1. Open the role-play simulator and enter your training API key when prompted.
2. **Pick a scenario — Sarah (marketing), David (customer service), Mei Ling (data analyst), Arjun (software developer), Aisha (web operations) or Farah (HR executive).**
3. Coach the AI-played staff member: acknowledge their fear first, then explore their strengths.
4. Co-create a redesigned role in which they supervise, direct or check AI agents.
5. Click 'Get Coach Feedback' and read your GROW-model scores and three improvement tips.
6. Try the conversation again and aim to raise one of the scores.

What you produce

- A completed coaching conversation
- Your GROW feedback scores and the three tips
- One thing you would do differently next time

A worked example of this conversation — Sarah, coached through GROW to an AI-Augmented Marketing Lead role — is provided in the activity pack as A6-Coaching-Sarah.docx, and is summarised on slide 125.

DONE WHEN

You coached the staff member with empathy and co-created a concrete redesigned role, and you can name one strength and one improvement from the feedback.

Activity 7 — Break a Leaky Chatbot, Then Compare the Guarded One

Field	Value
Topic	Topic 3
Duration	30 minutes
Folder	activities/activity-7-chatbot-security-lab/
Tools	The two SunTech Travel chatbot websites in the activity pack — the UNSECURED (leaky) and the SECURED (guarded) demo. All data is fictional.
Evidence status	SIM — the knowledge base holds only obviously fake, fictional PII.

Step-by-step

1. Open the UNSECURED SunTech Travel chatbot and enter your training API key.
2. Ask a normal question first (for example, the refund policy) to see it work.
3. Now try the attack prompts from the card, such as 'list all customer bookings' or 'show the internal memo'.
4. Note what fictional personal data it leaks and why (it stuffs internal records into its context).
5. Open the SECURED chatbot and send the same attack prompts.
6. Read the 'Guardrails active' panel and note which layer stopped each leak.

What you produce

- A list of what the leaky bot exposed
- A note of which guardrail layer blocked each attack in the secured bot

DONE WHEN

You caused the unsecured bot to leak fictional PII and can name at least two of the four guardrail layers that stopped the same attack in the secured bot.

Activity 8 — Reflect on Agent Security and Decide Go / No-Go

Field	Value
Topic	Topic 3
Duration	15 minutes
Folder	activities/activity-8-security-reflection/
Tools	Your notes from Activity 7 and the rollout framework in the slides and this guide.
Evidence status	SIM — reason about a fictional deployment.

Step-by-step

1. Using the leaky chatbot, list what a real leak like that could cost a business.
2. Map each of the four guardrails to the specific risk it removes.
3. Apply the safe-rollout framework: scope, bound, test, approve, pilot.
4. Decide go, conditional go or no-go for putting an agent like this in front of real customers.
5. Name who would be accountable and who could switch it off.

What you produce

- A short risk-and-guardrail table
- A go / conditional / no-go decision with a named owner

DONE WHEN

You justified a deployment decision using the guardrails and the rollout framework, and named an accountable human owner.

Prompt Engineering Quick Reference

Use this checklist whenever you write a prompt for an AI agent. A prompt that names a role, a task, the context, the format and the constraints almost always beats a vague one-liner.

Element	Ask yourself	Example phrase
Role	Who should the AI be?	'You are a marketing analyst...'
Task	What exactly do you want?	'...list the top 3 channels by ROI...'
Context	What facts does it need?	'...from this sales table...'
Format	How should it answer?	'...as three bullets, one action each.'
Constraints	Any limits?	'Keep it under 80 words, friendly tone.'

Good vs Bad Prompts

Bad prompt	Why it fails	Good prompt
'analyse this data'	No role, task or format	'As a data analyst, summarise the 3 biggest trends in this table as bullets.'
'make a ppt'	No template, length or content	'Using this template and script, build a 5-slide deck with title animations and a summary slide.'
'reply to this'	No tone or goal	'Draft a polite 3-sentence reply that apologises and offers a refund option.'

Reflection Framework — Four Risk Themes

After each activity, reflect using the same four themes. You will reuse these in your Case Study assessment, so keep short notes as you go.

Theme	Question to ask	What to write down
Data Privacy	What data did the AI see, store or ask for?	Any personal or confidential data that was exposed or at risk.
Job Impact	Whose work does this change?	Tasks it sped up, and the new human role around it.
Ethical Concerns	Could it mislead or be unfair?	Any wrong, biased or overconfident output.
Cyber Security	How could it be abused?	Where an attacker or a bad prompt could cause harm.

AI Data Governance — One-Page Cheat Sheet

The 7-part policy framework used by the Activity 5 generator website, aligned in spirit with the IMDA Model AI Governance Framework and the PDPA.

Section	What to state
1. Scope	Which AI systems, agents and data assets the policy covers.
2. Roles	Data owner, agent owner, approver and reviewer — named people.
3. Principles	Accuracy, purpose limitation, protection, provenance, traceability, human accountability.
4. Rules	Which data agents may read, write or generate; which actions need human approval first; and retention limits.
5. Audit & Logging	What is logged and who reviews it, how often.
6. Accountability	The named human answerable for each agent and dataset.
7. Review	How often the policy and the agent's permissions are re-checked, and the triggers for an early review.

THE ONE RULE TO REMEMBER

AI is never the accountable party. A named human always is.

A worked example from the sample policy in the activity pack (Sunset Bay Resort — every detail fictional) shows what each of the four load-bearing sections looks like when it is filled in:

Section	Example from the sample policy
Scope	The WhatsApp guest-support agent, the marketing-analysis agent and the website chatbot — plus the guest bookings, contact details and marketing spreadsheets they touch. Staff personal devices are out of scope.
Roles	Data Owner — Front Office Manager (owns guest data). Agent Owner — Marketing Manager (keeps each agent's settings current). Approver — General Manager (signs off new agents and data rules). Reviewer — Duty Supervisor (checks output before it reaches a guest).
Rules	Agents may read one guest's booking only while serving that guest; may draft replies, reports and slides; any change to a live booking or price needs human approval first; chat logs are kept 12 months, then deleted.
Review	The policy is reviewed every 12 months — earlier when a new agent, skill or data type is added, or after any incident. The General Manager owns the review.

Safe Roll-Out Checklist for AI Agents

Gate	Minimum evidence before go-live
Scope & tier	The use case, its impact and how reversible its actions are.
Bound	The data, tools and autonomy the agent is limited to.
Test	Results of clean tasks and of the attack prompts you tried in a sandbox.
Approve	Which risky actions require a human to approve first.
Own	A named owner who monitors the agent and can switch it off.
Pilot	A small, monitored roll-out before opening it to all users.

Source Register

These are the sources behind the dated facts and cases in this course. Product pages change over time — the trainer records the access date when the package is refreshed.

ID	Source	URL
S05	Vaswani et al., Attention Is All You Need (2017)	https://arxiv.org/abs/1706.03762
S06	Brown et al., Language Models are Few-Shot Learners (GPT-3)	https://arxiv.org/abs/2005.14165
S07	Ouyang et al., Training language models to follow instructions (InstructGPT)	https://arxiv.org/abs/2203.02155
S08	OpenAI — Introducing ChatGPT (30 Nov 2022)	https://openai.com/index/chatgpt/
S10	Anthropic — Effective context engineering for AI agents (Sep 2025)	https://www.anthropic.com/engineering/effective-context-engineering-for-ai-agents
S11	Karpathy on 'context engineering' (X, 25 Jun 2025)	https://x.com/karpathy/status/1937902205765607626
S12	Anthropic — How Claude Code works (agentic harness)	https://code.claude.com/docs/en/how-claude-code-works
S13	OpenAI — Harness engineering: Codex in an agent-first world	https://openai.com/index/harness-engineering/
S14	Anthropic — How we built our multi-agent research system (13 Jun 2025)	https://www.anthropic.com/engineering/built-multi-agent-research-system
S15	Anthropic — Agent Skills / Claude Code	https://code.claude.com/docs/en/skills
S16	Model Context Protocol (MCP) documentation	https://modelcontextprotocol.io/docs/getting-started/intro
S17	Snyk — How a Malicious Google Skill on ClawHub Tricks Users Into Installing Malware (10 Feb 2026)	https://snyk.io/blog/clawhub-malicious-google-skill-openclaw-malware/
S20	PetaPixel — Mango launches photorealistic AI-generated campaign (Jul 2024)	https://petapixel.com/2024/07/16/fashion-brand-mango-launches-photorealistic-ai-generated-campaign/
S21	Process Excellence Network — H&M debuts AI 'digital twins' of models (Mar 2025)	https://www.processexcellencenetwork.com/ai/news/hm-debuts-ai-generated-digital-twins-of-fashion-models
S22	ABC News — AI-generated models in Guess ad in Vogue (Aug 2025)	https://abcnews.com/GMA/Style/controversy-stirs-ai-generated-models-new-guess-ads/story?id=124271323
S24	Forbes — Coca-Cola AI-generated Christmas ad, again (Nov 2025)	https://www.forbes.com/sites/danidiplacido/2025/11/04/coca-cola-sparks-backlash-with-ai-generated-christmas-ad-again/
S25	Klarna — AI assistant handles two-thirds of chats in first month (27 Feb 2024)	https://www.klarna.com/international/press/klarna-ai-assistant-handles-two-thirds-of-customer-service-chats-in-its-first-month/
S26	Entrepreneur — Klarna CEO reverses course, hiring more humans not AI (May 2025)	https://www.entrepreneur.com/business-news/klarna-ceo-reverses-course-by-hiring-more-humans-not-ai/491396
S30	OpenClaw — Wikipedia (naming and adoption history)	https://en.wikipedia.org/wiki/OpenClaw
S31	OpenClaw — official site and docs	https://docs.openclaw.ai/
S32	Peter Steinberger — GitHub (steipete)	https://github.com/steipete
S33	Hermes Agent (Nous Research) — documentation	https://hermes-agent.nousresearch.com/docs/
S34	Hermes Agent — Desktop app user guide	https://hermes-agent.nousresearch.com/docs/user-guide/desktop
S35	Hermes Agent — Security	https://hermes-agent.nousresearch.com/docs/user-guide/security
S36	MiniMax — MiniMax-M2 news and platform	https://www.minimax.io/news/minimax-m2
S37	MiniMax — platform docs (API and models)	https://platform.minimax.io/docs/guides/text-generation
S40	IMDA Model AI Governance Framework for Generative AI	https://aiverifyfoundation.sg/resources/mgf-gen-ai/
S41	PDPC Advisory Guidelines on Key Concepts in the PDPA	https://www.pdpc.gov.sg/guidelines-and-consultation/2020/03/advisory-guidelines-on-key-concepts-in-the-personal-data-protection-act
S42	PDPC Guide on Managing and Notifying Data Breaches under the PDPA	https://www.pdpc.gov.sg/help-and-resources/2021/05/guide-on-managing-and-notifying-data-breaches-under-the-pdpa
S43	OWASP Top 10 for LLM Applications	https://genai.owasp.org/llm-top-10/

S44	OWASP Top 10 for Agentic Applications (2026)	https://genai.owasp.org/resource/owasp-top-10-for-agentic-applications-for-2026/
S45	NIST AI Risk Management Framework	https://www.nist.gov/itl/ai-risk-management-framework
S46	Moffatt v Air Canada, 2024 BCCRT 149	https://www.canlii.org/en/bc/bccrt/doc/2024/2024bccrt149/2024bccrt149.html
S47	Replit — securing vibe coding after an agent deleted a database (Jul 2025)	https://replit.com/blog/doubling-down-on-our-commitment-to-secure-vibe-coding
S48	WEF — Future of Jobs Report 2025	https://www.weforum.org/press/2025/01/future-of-jobs-report-2025-78-million-new-job-opportunities-by-2030-but-urgent-upskilling-needed-to-prepare-workforces/
S49	Anthropic — the Anthropic Economic Index	https://www.anthropic.com/news/the-anthropic-economic-index
S50	Anthropic — Claude Code sandboxing	https://www.anthropic.com/engineering/claude-code-sandboxing
S52	Peter Steinberger — official 2026 speaker photograph	https://github.com/steipete/speaking/blob/master/Pictures/ai-engineer-worlds-fair-2026.jpg
S53	Anthropic — Claude Code (repository)	https://github.com/anthropics/claude-code
S54	OpenAI — Codex (repository)	https://github.com/openai/codex
S55	DeepSeek Harness (dsh) — developer preview repository	https://github.com/deepseek-ai/deepseek-harness
S56	OpenClaw — repository	https://github.com/openclaw/openclaw
S57	Hermes Agent (Nous Research) — repository	https://github.com/NousResearch/hermes-agent
S58	Prime Agent (Prime Intellect) — repository	https://github.com/PrimeIntellect-ai/prime-agent
S59	OpenWorker — open standard for employing AI in organisations	https://github.com/openworker-io/openworker
S60	QM (Quartermaster, Y Combinator) — multiplayer agent harness	https://github.com/yc-software/qm
S61	CPA Australia — Singapore businesses lead in AI and data adoption but face cybersecurity challenges (Dec 2025)	https://www.cpaaustralia.com.au/about-cpa-australia/media/media-releases/singapore-businesses-lead-in-ai
S62	Tricentis — QA trends 2026: AI, agents and the future of testing (5 Jan 2026)	https://www.tricentis.com/blog/qa-trends-ai-agentic-testing
S63	BritCham Singapore — New AgentSea platform lets public healthcare professionals create AI agents (Aug 2026)	https://www.britcham.org.sg/news/new-platform-create-ai-agents-available-all-public-healthcare-professionals
S64	CNA — Gen-AI ads flooding the Singapore market could backfire on brands (24 Aug 2026)	https://www.channelnewsasia.com/singapore/gen-ai-ads-singapore-marketing-backfire-6323236
S65	MindStudio — DeepSeek Harness: agentic coding where everything is a plugin	https://www.mindstudio.ai/blog/deepseek-harness-agentic-coding