



TRAINER SLIDES · WSQ

AI Security for Autonomous AI Agents

WSQ Course Code: TGS-2025060473 · 1 Day · 8 Hours

Skills Framework TSC: Generative AI Principles and Applications
(ICT-INT-0052-1.1)

Conducted by Tertiary Infotech Pte Ltd · UEN 201200696W

Version 4.13 · 26 August 2026

© 2026 Tertiary Infotech Pte Ltd. All Rights Reserved. · www.tertiarycourses.com.sg

WSQ

AI SECURITY



Digital Attendance (Mandatory)

- It is mandatory to take the AM, PM and Assessment digital attendance for WSQ-funded courses.
- The trainer or administrator displays the digital attendance QR code from the SSG portal.
- Scan the QR code with your mobile phone camera and submit your attendance.
- A minimum of 75% attendance is required to be eligible for assessment and funding.
- Complete the TRAQOM survey at the end of the course — it is required for funding.

About the Trainer



Your Trainer

Complete this profile before delivery

NAME

QUALIFICATIONS

INDUSTRY EXPERIENCE

TRAINING EXPERIENCE

About the Trainer



Dr Alfred Ang

Principal Trainer
Tertiary Infotech Pte Ltd

ROLE

Principal Trainer

EXPERTISE

AI security, governance and enterprise IT

Who Is in the Room?

1

Role

What service or team do you work in?

2

AI today

Where do you already meet AI at work?

3

Goal

What do you want to be able to do by 6pm?

Ground Rules for a Hands-On Day

1

Use the demos

Every activity is a ready-made website or chatbot — no coding

2

Fictional data only

Never paste real personal or company data into a demo

3

Ask and share

Try prompts, compare answers, and reflect together

Learning Outcomes

1

LO1

Demonstrate generative AI concepts and applications

2

LO2

Apply prompt engineering and analyse output variations


3

LO3

Identify ethical risks and analyse bias in AI content


Evidence: ADMIN

One-Day Learning Journey



Three topics map one-to-one onto LO1, LO2 and LO3; the assessment closes the day.

Briefing for Assessment




The briefing for assessment is delivered before the assessment begins.

How the Day Is Assessed

Instrument	Covers	What you do
Written Assessment (SAQ)	K1–K5	Five short written answers, one per knowledge statement
Case Study	LO1–LO3 (A1–A5)	Three reflection tasks (two questions each) on the activities you completed today
Format	—	Open book · Competent / Not Yet Competent

Evidence: ADMIN

Courseware and Activities on the LMS



1

Step 1

Download the current learner-facing files from the LMS

2

Step 2

Open each activity website directly in your browser — nothing to install

3

Step 3

Keep every test synthetic, local and reversible

Evidence: ADMIN


Use only the current learner-facing files. Keep evidence synthetic, local and reversible; never upload secrets or answer keys.

TOPIC 1 · LO1

Generative AI, Agentic AI and AI Agents

From a chat box to systems that plan and act — history, mechanics, use cases and agents.

What Topic 1 Will Answer



Each question builds the vocabulary you need for prompt engineering and security later.

A Very Short History of AI, 2023–2026



Each year did not replace the last — it added a new layer on top of it.

The Four Waves in Plain Words

1

**2023 —
Generative AI**

ChatGPT makes text and image generation mainstream

2

**2024 — Context
Engineering**

We learn to feed models the right information, not just a prompt

3

**2025 — Harness
Engineering**

We wrap models in loops that plan, act and verify — agentic AI

4

**2026 — AI
Agents**


Deployed agents act across our apps, chats and tools

2023 — Generative AI Goes Mainstream

- ChatGPT launched on 30 November 2022 and reached mass adoption through 2023
- One natural-language box exposed writing, coding, translation and analysis to everyone
- Businesses began asking not 'can it chat?' but 'what work can it do?'

Sources: [OpenAI \(openai.com\)](https://openai.com)

2024 — Context Engineering



Beyond the prompt

The model also needs the right supporting information

The context window

Instructions, history, memory, tools and retrieved data meet here

The craft

Select just the right tokens for the model's next step

Sources: [Anthropic \(anthropic.com\)](https://anthropic.com) · [Karpathy on 'context engineering' \(X, 25 Ju... \(x.com\)](https://x.com/Karpathy/status/17925060473)

2025 — Harness Engineering and Agentic AI



Sources: [Anthropic \(code.claude.com\)](https://code.claude.com) · [OpenAI \(openai.com\)](https://openai.com)

Loop

Gather context, plan, act and verify until the goal is met

Tools

The harness gives the model controlled ways to affect real systems

Safety

Permissions, sandboxes and verification constrain each action

2026 — The Year of AI Agents

- Personal and organisational agents such as OpenClaw and Hermes reached everyday users
- Agents now live inside WhatsApp, Telegram, email and business apps
- The question shifted again: how do we let agents act safely on our behalf?

Sources: [OpenClaw \(en.wikipedia.org\)](https://en.wikipedia.org/wiki/OpenClaw) · [Hermes Agent \(Nous Research\) \(hermes-agent.nousresearch.com\)](https://hermes-agent.nousresearch.com)

TOPIC 1

How Generative AI Works

Training, inference, and the autoregressive language model.

What 'Generative' Means

1

Generative AI

Produces new text, image, audio, video or code from patterns it has learned

2

Discriminative AI


Sorts or scores existing input — spam / not-spam, fraud / not-fraud

3

Why it matters


Service teams use both: generate a reply, and classify a request

Two Phases: Training and Inference



Training happens once and is expensive; inference happens every time you use the model.

How an Autoregressive LLM Writes



A large language model writes one token at a time, each based on everything before it.

Why the Same Prompt Can Differ

1

Probabilities

The model samples from likely next tokens, not a fixed lookup

2

Temperature

Higher settings add variety; lower settings add consistency

3

Context

Change the surrounding information and the whole answer shifts

Strengths and Limits for Service Work

1

Strong at

Drafting, summarising, translating, reformatting, brainstorming

2

Weak at

Facts it was never given — it can sound confident yet be wrong

3

Rule

Treat fluent output as a draft to verify, not as truth

Evidence: DEF

Probabilistic Output — Why Traditional Security Struggles

Traditional software

- Deterministic — the same input always produces the same output
- Behaviour can be fully tested before release
- Firewalls and filters match known, fixed patterns
- An unexpected output is a bug you can reproduce and patch

A large language model

- Probabilistic — the same prompt can produce a different answer every run
- The space of possible outputs can never be fully tested
- There is no fixed signature — a harmful answer can look new every time
- An unexpected output is normal behaviour, not a fault

You cannot write a rule for every possible answer — so security must also watch what the model says and does, not just what goes in.

TOPIC 1

Generative AI in the Real World

Real, dated examples — how generation is already changing work.

Fashion Models Are Now Generated



SYNTHETIC MODEL · IMAGEGEN

1

Mango, Jul 2024

Ran its first fully AI-generated campaign for its teen 'Sunset Dream' line across 95 markets

2

H&M, 2025

Announced AI 'digital twins' of 30 real models; first labelled images appeared mid-2025

3


Guess, Aug 2025

A Guess ad using AI-generated models ran in Vogue's August issue and drew wide debate

Sources: [PetaPixel \(petapixel.com\)](https://petapixel.com) · [Process Excellence Network \(processexcellencenetwork.com\)](https://processexcellencenetwork.com) · [ABC News \(abcnews.com\)](https://abcnews.com)

AI-Generated Video Is Overtaking Whole Industries

EMBEDDED VIDEO · PLAYS IN SLIDESHOW



Click the video, then Play, in slideshow

Music

This AI-generated singer — no camera, cast or crew — shows how AI is reshaping the music industry

Three industries

The next three demos show marketing, film and how fast this is moving

Made for feeds

Vertical clips like this are built for TikTok, Reels and YouTube Shorts

For your team

Ask: where could you use this, and where would you not?

The clip is embedded — click it and press Play in slideshow.

Evidence: SYN

AI in Marketing and Advertising

EMBEDDED VIDEO · PLAYS IN SLIDESHOW



Click the video, then Play, in slideshow

Evidence: SYN

Campaign in hours

Brands now generate ad and campaign video in hours, not weeks

No production crew

No shoot, studio, cast or location — just a prompt and a model

Coca-Cola

Aired an AI-generated 'Holidays Are Coming' ad in 2024 and again in 2025

Watch for

Brand risk, authenticity and disclosure when the whole ad is synthetic

Embedded clip — click and press Play. Discuss: would you trust an all-AI ad for your brand?

AI in Film and Movies

EMBEDDED VIDEO · PLAYS IN SLIDESHOW



Click the video, then Play, in slideshow

Into film production

AI now generates visuals, scenes and effects for film

Faster and cheaper

Shots that needed a crew and budget can be generated

Still maturing

Consistency, control and rights are the open questions


The takeaway

Whole creative industries are being reshaped — fast

Embedded clip — click and press Play. Discuss: what does this mean for creative jobs?

Evidence: SYN

Customer Service: The Klarna Story



A real cautionary tale: automate the routine, but keep humans for judgement and empathy.

Sources: [Klarna \(klarna.com\)](https://klarna.com) · [Entrepreneur \(entrepreneur.com\)](https://entrepreneur.com)

Where Service Teams Apply GenAI Today

1

Front desk

Draft replies, translate, summarise long threads

2

Back office

Turn notes into reports; extract data from documents

3

Marketing

Generate copy, images and short video for campaigns

TOPIC 1

Context Engineering

The model is only as good as the context you give it.

Prompt Engineering vs Context Engineering

1

Prompt engineering

Wording one instruction well

2

Context engineering

Assembling everything the model sees before it answers

3

The shift

Real applications manage context, not just a clever prompt

The Six Ingredients of Context

1

System prompt

Who the AI is and its rules

2

User prompt

The task you asked for right now

3

History

Earlier turns in this conversation

4

Memory

Facts kept across sessions

5

Tools


Functions it can call to act or fetch data

6

Retrieved data

Documents pulled in for this task

How an Agent Gathers Context for a Task



Example: 'Reply to this guest email' → it pulls the booking, the policy doc and past replies first.

A Worked Context Example

1

Task

'Draft a refund reply to Mr Tan'

2

Context gathered

Booking record + refund policy + tone guide + the original email

3

Result

A grounded reply, not a generic guess


Evidence: SYN

TOPIC 1

The Agentic Loop and Harness Engineering

What turns a model that answers into a system that acts.

The Agentic Loop — A Visual Pipeline



The loop repeats until the goal is met. Anthropic's harness uses gather → act → verify; 'plan' is shown here as a teaching step.

What a Harness Adds Around the Model

1

The loop

Runs context → plan → execute → verify repeatedly

2

Tool access

Lets the model read files, run commands, call APIs

3

State & safety

Tracks progress and enforces permissions and sandboxes

Two Harnesses You May Have Heard Of

1

Claude Code

Anthropic's coding agent — the 'agentic harness around Claude'

2

Codex

OpenAI's coding agent; OpenAI calls the discipline 'harness engineering'

3

Same idea

A loop + tools + a sandbox around a general model

The Coding Harnesses — Claude Code, Codex, DeepSeek

Harness	What it is	Where to find it
Claude Code	Anthropic's terminal coding agent — reads your codebase, edits files, runs tests and git by natural language	github.com/anthropics/claude-code
Codex	OpenAI's coding agent; its execution engine (the 'harness') was open-sourced in Aug 2026	github.com/openai/codex
DeepSeek Harness (dsh)	DeepSeek's plugin-based, model-agnostic harness — a v0.1 developer preview	github.com/deepseek-ai/deepseek-harness

All three wrap a model in the same gather-context → act → verify loop; they differ in tools, sandbox and host authority.

DeepSeek Harness — 'Everything Is a Plugin'

1

One idea

Every capability — tools, models, UI, even core behaviour — is a swappable plugin

2

Model-agnostic

Plug in any model; the harness is not tied to DeepSeek's own


3

Why it matters

Flexible and extensible — but every plugin is also new code and a new trust boundary to review


A v0.1 developer preview. Its openness is powerful; in Topic 3 we treat each plugin as untrusted until reviewed.

A Concrete Agentic Loop



If verification fails — a folder is empty — the agent re-plans instead of stopping.

When the Loop Goes Wrong — Misreading the Goal



The same re-planning that makes the loop powerful becomes dangerous when the goal is open-ended and nothing limits the actions — Topic 3 adds the guardrails and human approvals that stop this.

TOPIC 1

AI Agents: OpenClaw, Hermes and the Latest

The deployed systems that put the agentic loop into everyday hands.

What Makes Something an 'AI Agent'

1

A real system

Not just a model — a running app with an identity

2

Acts through tools

Sends messages, edits files, calls services

3

Keeps state

Remembers context and can resume work

Peter Steinberger in 2026 — Creator of OpenClaw



Evidence: PROD · Sources: S30, S32, S52

Who

Austrian developer, founder of PSPDFKit; GitHub handle steipete


What he built

OpenClaw, an open-source personal AI agent that went viral in late 2025

Since

Joined OpenAI in Feb 2026; an OpenClaw Foundation now stewards the project

OpenClaw — A Naming and Adoption Story



Renamed after a trademark complaint; figures per Wikipedia, as of early 2026.

OpenClaw Runs Where You Already Chat

1

Messaging-first

You talk to it in WhatsApp, Telegram, Slack, Signal and more

2

Self-hosted

It runs on your own machine or server, under your control

3

Tool-enabled

It can read, write and call services on your behalf

Hermes Agent by Nous Research



Evidence: PROD · Sources: S33

Open-source agent

CLI, gateway, messaging surfaces and a desktop app

Self-improving

Can create reusable skills from experience

Flexible models

Runs on many providers — we will point it at MiniMax

The Latest AI Agents (2025–2026)

Agent	What it is	Where to find it
OpenClaw	Peter Steinberger's open-source personal agent; runs on your own devices, reached via your messaging apps	github.com/openclaw/openclaw
Hermes Agent	Nous Research's self-improving personal agent; learns skills, model-agnostic (we use MiniMax)	github.com/NousResearch/hermes-agent
Prime Agent	Prime Intellect's self-improving agent for coding and long-running autonomous tasks	github.com/PrimeIntellect-ai/prime-agent
OpenWorker	An open standard to hire, govern and trust AI 'workers' inside an organisation	github.com/openworker-io/openworker
QM (Quartermaster)	Y Combinator's multiplayer agent harness for teams in Slack / web, with scoped memory and sandboxes	github.com/yc-software/qm

All open-source and self-hosted. Capability is not a safety rating — verify each one's permissions, sandbox and data access before use.

Applications of AI Agents in Threat Detection

1

Real-time threat detection

Agents watch traffic, endpoints and events continuously and flag cyber threats as they emerge — not hours later

2

Automated vulnerability detection

Scan code, systems and configurations for weaknesses, then analyse severity and exploitability

3

Malware analysis & classification

Dissect and classify malicious samples far faster than manual reverse-engineering

AI is on both sides of the fight — the same capabilities also power attackers, which is why Topic 3 defends the agents themselves.

AI Agents in Network and Social-Engineering Security

1

Network intrusion detection

Learn what normal network behaviour looks like and catch intrusions that signature rules miss

2

Attack classification

Label detected attacks by type and severity so responders know what to tackle first

3

Phishing & deceptive language

Detect and defend against phishing attacks and manipulative, deceptive language in messages

4

Log analysis & anomaly detection

Sift millions of system-log lines to surface the anomalies worth investigating

Applications of AI Agents in Incident Response

1

Repair & patch generation

Automate vulnerability repair and generate candidate patches for human review

2

Workflows & playbooks

Streamline incident-response workflows and execute playbooks consistently under pressure

3

Post-attack analysis

Reconstruct the attack, identify the root cause and recommend fixes to stop a repeat

Applications of AI in Security Operations Centres (SOCs)

1

Proactive defence & threat hunting

Hunt for hidden threats across the estate instead of waiting for alerts

2

Risk management & predictive analytics

Score risks and predict likely attack paths before they are exploited

3

Adaptive decision-making

Learn continuously from every incident and adapt defences over time

The SOC of 2026 pairs human analysts with AI agents — the analyst sets direction and approves actions; the agent does the heavy lifting.

Try It Today: TIA Support on WhatsApp

1

The number

TIA Support WhatsApp +65 8866 6375,
powered by OpenClaw

2

What to do

Message it like a person and give it a
small task

3

Watch for

Where it helps — and where you would
not trust it yet

We use this live in Activity 1.

An Agent on WhatsApp Is Exposed to the Internet

1

Open to anyone

A WhatsApp or Telegram agent accepts messages from the whole internet — including bad actors

2

Prompt injection

A crafted message overrides the agent's instructions and hijacks its tools

3

Data poisoning

Malicious content the agent reads or saves corrupts what it does later

4

What is at stake

Confidential data extracted, or the server behind the agent harmed

The convenience of a public channel is also its attack surface — Topic 3 shows the defences.

TOPIC 1

Skills and Tools

How agents extend beyond the base model.

Tools vs Skills

1

Tools

Actions the agent can take — send email, run a query, open a file

2

Skills


Packaged know-how — reusable instructions for a task

3

Together

Skills tell the agent how; tools let the agent do

How an Agent Uses a Tool



Model Context Protocol (MCP) is a common standard for connecting tools to agents.

Why Skills Make Agents Practical

1

Consistency

The same task is done the same way every time

2

Reuse

Build once, run across many jobs

3

In Topic 2

We install a skill so the agent formats Excel and PPT better

Case — A Malicious Skill on ClawHub



Snyk, 10 Feb 2026 — the malware was in the skill's instructions, not its code. Install skills only from sources you trust; clones of the flagged skill reappeared within hours.

Sources: [Snyk \(snyk.io\)](https://snyk.io)




Morning Break

TOPIC 1

Multi-Agent Systems

When one agent is not enough.

How a Multi-Agent System Works



A lead agent takes your goal, delegates to specialist workers that run in parallel, then combines their results. Anthropic's research system used this pattern and beat a single agent by ~90% on its internal eval.

Sources: [Anthropic \(anthropic.com\)](https://anthropic.com)

Why and Why Not Multi-Agent

1

Faster & broader

Parallel workers cover more ground

2

Costlier

Multi-agent runs used far more tokens than a single chat

3

Trade-off

Use it when breadth matters more than cost

A Service Example — Customer Recovery



Evidence: SYN

Goal

Recover customer trust

Parallel workers

Draft message · Build offer ·
Check customer data

Lead + human

Merge one plan · Review ·
Approve

Agent-to-Agent Messaging — A Security Blind Spot

1

One agent's output is the next agent's input

A hallucinated or maliciously injected result is passed downstream — and the receiving agent treats it as trusted fact

2

Errors and attacks cascade

One bad message can propagate through every agent behind it, ending in a wrong or harmful action far from where it started

3

Topology multiplies message paths

Every extra agent adds agent-to-agent links — no single log shows the whole data flow end to end, so visibility drops

4

Tracing and attribution get harder

When something goes wrong, finding which agent introduced the bad data needs per-agent logs of every hand-off

More agents means more paths to monitor. Log every agent-to-agent hand-off and validate what an agent receives — never trust a message just because another agent sent it. OWASP's 2026 agentic Top 10 flags exactly this cascading multi-agent failure mode.

TOPIC 1

Generative AI vs Agentic AI vs AI Agents

The single most useful distinction from Topic 1.

The Three, Side by Side

	Generative AI	Agentic AI	AI Agent
What it does	Creates content	Loops to reach a goal	A deployed system that acts
Acts on its own?	No — you run each step	Yes — plans and retries	Yes — through real tools
Example	ChatGPT drafting a reply	Claude Code fixing a bug	OpenClaw in your WhatsApp

Evidence: SYN · Sources: S43, S44

One-Line Definitions

1

Generative AI

Makes things

2

Agentic AI

The loop that makes an AI pursue a goal

3

AI Agent

The running system that uses that loop to act for you

TOPIC 1 TAKEAWAY

Generation is a step. Agency is a loop. An agent is a system.

Activity 1 — Talk to an AI Agent, Then Reflect

Hands-on, no-code — ready-made website or AI agent

Pinboard: <https://alfredang.github.io/pinboard/>

1

Scope

2

Map

3

Test

4

Control

5

Evidence

DISCUSSION QUESTIONS

- Message TIA Support on WhatsApp +65 8866 6375 (powered by OpenClaw)
- Use the prompt card in the activity pack; give it a small realistic task
- Post your reflections to the Pinboard under the four risk themes

YOU'LL PRODUCE

A documented risk decision with evidence.

Assesses LO1 · A2, A4 · Activity pack: [activities/activity-1-genai-agent-whatsapp/](#) · Full step-by-step facilitation detail: [Learner Guide](#)

Evidence: SIM

Activity 1 — Group Your Concerns

1

Data Privacy

What data did it see or ask for?

2

Job Impact

Whose work does this change?

3

Ethical Concerns


Could it mislead or be unfair?

4

Cyber Security

How could it be abused?

Activity 1 — Debrief




This debrief seeds the governance thinking you will use in Topic 3.

TOPIC 2 · LO2

Prompt Engineering and Post-Training for Autonomous AI Agents

Set up a real agent on MiniMax, then learn to prompt it well.

What Topic 2 Will Cover



You will do the setup once, then use the same agent for both activities.

The Setup in Three Parts

1

Hermes Agent

The agent you talk to

2

MiniMax M2.7


The model that powers it

3

Your API key


Connects the two — kept only on your machine

Install the Hermes Desktop App




If the CLI is already installed, `'hermes desktop'` reuses your existing config and keys.

Get a MiniMax Model and API Key



Check platform.minimax.io for current pricing and any promotional access before class.

Point Hermes at MiniMax M2.7



Hermes supports custom OpenAI-compatible endpoints, so MiniMax plugs straight in.

Keep Your Key Safe

1

Training key

Use a low-limit key, not a production one

2

Local only

The key stays on your machine; do not share it

3

Revoke after

Delete the key when the course ends

TOPIC 2

Prompt Engineering Principles

Small changes in the prompt make large changes in the output.

Five Principles of a Good Prompt



Evidence: DEF · Sources: S10

Role

Who the AI should be

Task

Exactly what result you want

Context

The facts it needs

Format

How the answer should look

Constraints

Length, tone and limits

A Bad Prompt vs a Good Prompt

Bad prompt

- 'analyse this data'
- No role, no goal
- No format
- You get a vague wall of text

Good prompt

- 'You are a marketing analyst. From this sales table, give the top 3 trends as bullets, with one action each.'
- Clear role and task
- Clear format
- You get a usable answer

Techniques That Reliably Help

1

Give an example

Show one sample of the output you want

2

Ask for steps

'Think step by step' improves reasoning tasks

3

Iterate

Refine the prompt after seeing the first answer

'Post-Training' for Agents — Memory and Skills, Not Fine-Tuning

1

Not fine-tuning

You are not retraining the model's weights — that is expensive and out of reach for most teams

2

It's context engineering

The agent stores what it learns in memory files and reusable skills

3

Improves over time

Next time, it loads those memories and skills — so it gets better without retraining

4

Stays consistent

The same saved skill makes the agent do a task the same way every time

This is how the Hermes agent 'self-improves': it writes learnings to memory and skills, then reuses them.

Reflect — Does Prompt Engineering Still Matter With Agents?

1

The question

An agentic loop can plan, retry and self-correct — so is a good prompt still worth the effort?

2

Still yes

A clear prompt points the loop at the right goal from the start, instead of it guessing and looping

3

Saves time & tokens

A vague prompt makes the agent take extra steps — more turns, more tokens, more cost and delay

4

The takeaway

Prompt engineering matters MORE with agents: a clear brief up front saves a whole loop of rework

Discuss: think back to your bad-vs-good prompts — how many extra agent steps did the vague one cause?

Activity 2 — Analyse Excel Data with the Agent

Hands-on, no-code — ready-made website or AI agent

You will produce a Word report with charts, analysis and numbers-backed recommendations.

1

Scope

2

Map

3

Test

4

Control

5

Evidence

DISCUSSION QUESTIONS

- Upload MOCK-MARKETING-DATA.xlsx to your Hermes + MiniMax agent
- Warm up with a bad prompt vs a good prompt on the same data
- Then prompt the agent to analyse the data, generate charts and produce a .docx report with recommendations

YOU'LL PRODUCE

A .docx marketing report with charts, analysis and numbers-backed recommendations.

Assesses LO2 · A5 · Activity pack: activities/activity-2-excel-analysis/ · Full step-by-step facilitation detail: Learner Guide

Evidence: SIM

Activity 2 — From Spreadsheet to Report

1

Bad

'which is better?'

2

Good

'Rank channels by ROI =
Revenue / Spend as a table'

3

Report

'Analyse the file, chart
revenue, ROI and spend,
and generate a .docx report
with recommendations'

4

Verify

Check two numbers in the
report against the
spreadsheet

Evidence: SIM

Activity 3 — Create and Redesign a PPT with the Agent

Hands-on, no-code — ready-made website or AI agent

The image template and two worked-example decks are in the activity pack.

1

Create

2

Save to Downloads

3

Upload template

4

Redesign

5

Compare

DISCUSSION QUESTIONS

- Ask the agent for a 10-slide deck on AI Security for AI Agents
- Save the finished deck to your Downloads folder
- Upload the image template and ask the agent to redesign the deck

YOU'LL PRODUCE

Two decks in your Downloads folder — the original and the template-redesigned version.

Assesses LO2 · A5 · Activity pack: [activities/activity-3-ppt-builder/](#) · Full step-by-step facilitation detail: [Learner Guide](#)

Evidence: SIM

Activity 3 — What Good Looks Like

1

Create

'Build a 10-slide deck on AI Security for AI Agents with speaker notes'

2

Redesign

'Restyle every slide using this uploaded image as the visual theme'

3

Reflect

Content first, then design — the agent iterates on its own output

Evidence: SIM

Activity 4 — Install Tools and Skills to Do Better

Hands-on, no-code — ready-made website or AI agent

Skills give the agent repeatable know-how for formatting and structure.

1

Scope

2

Map

3

Test

4

Control

5

Evidence

DISCUSSION QUESTIONS

- Install the supplied tool/skill in your agent
- Re-run the Excel and PPT tasks
- Note how much the output improves

YOU'LL PRODUCE

A documented risk decision with evidence.

Assesses LO2 · A1, A5 · Activity pack: activities/activity-4-tools-and-skills/ · Full step-by-step facilitation detail: Learner Guide

Evidence: SIM



Lunch Break

Reflect on Working with the Agent

1

Data Privacy

What did you feed the agent and the model?

2

Job Impact

Which of your tasks did it speed up?

3

Ethical Concerns

Did any output mislead or overclaim?


4

Cyber Security

Where did your API key and data go?

Evidence: SIM

Topic 2 Debrief




Carry the security worries forward — Topic 3 turns them into governance.

TOPIC 3 · LO3

Security Risk of Autonomous AI Agents

Governance, jobs and the security of agents that can act.

What Topic 3 Will Cover



Real Singapore cases drive the discussion — building to the key AI challenges every business must manage.

Discussion Point 1 — Who Is Accountable for the Data?

1

The setup

An agent reads and edits your Excel and PPT

2

The problem

If it changes the data wrongly, who is accountable — AI or human?

3


Second case

If AI-made models or scripts break a law, who answers for it?

GOVERNANCE PRINCIPLE

AI is never the accountable party. A named human always is.

Case — Moffatt v Air Canada



2024 BCCRT 149 — an organisation is accountable for what its AI tells customers.

Sources: [Moffatt v Air Canada, 2024 BCCRT 149 \(canlii.org\)](#)

Traditional Data Governance

1

Purpose

Collect data only for a stated reason

2

Protection

Keep it secure and access-controlled

3

Accuracy & retention

Keep it correct; delete when no longer needed

New Principles for GenAI and Agents

1

Provenance

Know where training and generated data came from

2

Traceability

Log what the agent read, changed and produced

3

Human accountability

A named owner signs off agent actions on data

Scope of an AI Data Governance Policy

Area	What the policy must state
Data assets	Which data agents may read, write or generate
Access & identity	Who and which agent identity may touch each asset
Human approval	Which changes need a person to approve before they happen
Audit	What is logged, and who reviews it
Accountability	The named owner for each agent and dataset

Evidence: SYN · Sources: S40, S41

The 7-Part AI Data Governance Policy Framework

Section	What it covers
1 · Scope	The AI systems, agents and data assets the policy covers
2 · Roles	Data Owner, Agent Owner, Approver, Reviewer — every role a named human
3 · Principles	Accuracy, purpose limitation, protection, provenance, traceability, human accountability
4 · Rules	What agents may read, write and generate; approvals; retention
5 · Audit & Logging	What each agent did, who reviews the logs and how often
6 · Accountability	A named human is answerable — never the AI
7 · Review	Review cadence, the review owner, and early-review triggers

The structure of the sample policy in the activity pack — aligned in spirit with the IMDA Model AI Governance Framework and the PDPA.

The Sample Policy in Action — Sunset Bay Resort

Section	Example from the sample policy (a fictional resort)
Scope	The WhatsApp guest-support agent, the marketing-analysis agent and the website chatbot — plus the guest bookings, contact details and marketing spreadsheets they touch
Roles	Data Owner — Front Office Manager · Agent Owner — Marketing Manager · Approver — General Manager · Reviewer — Duty Supervisor
Rules	Agents may read one guest's booking only while serving that guest; may draft replies, reports and slides; any change to a live booking or price needs human approval first; chat logs kept 12 months, then deleted
Review	The policy is reviewed every 12 months — earlier when a new agent, skill or data type is added, or after any incident

From the sample AI Data Governance Policy in the activity pack — every detail is fictional.

Activity 5 — Draft Your AI Data Governance Policy

Hands-on, no-code — ready-made website or AI agent

The generator website (ai-data-policy-generator.html) drafts the policy from your company details using the 7-part framework. Refine weak sections with the Hermes prompts in PROMPTS.md. Keep every detail fictional.

1

Open the AI Data Policy Generator website

2

Enter your training API key

3

Describe your fictional company & its agents

4

Generate the 7-part policy

5

Check, refine and download

DISCUSSION QUESTIONS

- Does every role belong to a named human — never the AI?
- Which changes to live data need human approval first?
- Where did your API key go when you typed it into the website?

YOU'LL PRODUCE

A one-to-two page AI data governance policy (.md or PDF) for your fictional company

Assesses LO3 · A1, A2 · Activity pack: [activities/activity-5-data-governance-policy/](#) · Full step-by-step facilitation detail: [Learner Guide](#)

Evidence: SIM

Discussion Point 2 — Brand Damage from Generative AI

1

The temptation

AI can make an ad or a campaign in 24–48 hours, from about S\$1,000

2

The risk

A cheap or misleading AI ad can read as low-effort and hurt the brand

3

Who feels it most

Big and luxury brands — the backlash is sharpest for them

Evidence: SIM

Case — Gen-AI Ads Flood Singapore, and Could Backfire

What is happening	The warning
AI ad demand in Singapore is up 4–5x; half of advertisers already use gen-AI for video	'Using AI will not necessarily damage a brand. Using it poorly can.' (SMU Prof Sabine Benoit)
Virtual influencers and AI livestream avatars are replacing human hosts	Virtual influencers can 'give people the creeps' and breed mistrust (NUS Assoc Prof Ang Swee Hoon)
Luxury and global brands are adopting AI ads to cut cost	Seen as 'lacking effort' and 'less authentic'; 'not all publicity is good publicity'
Regulators are watching	ASAS: 'Disclosure of the use of AI alone does not whitewash this issue'

CNA, 24 Aug 2026. An industry-wide risk, not one named failure — the harm is to trust, authenticity and brand equity.

Sources: [CNA \(channelnewsasia.com\)](https://www.channelnewsasia.com)

Why AI Content Can Damage a Brand

1

Signals low effort

'The company hasn't put much effort into the product'

2

Feels inauthentic

Synthetic faces and voices erode trust and connection

3

Ethical shadow


Models trained on creators' work 'without their consent' — a visible symbol of unfairness

4

Backlash is costly

Negative buzz can force an ad down — the campaign fails its objective

Discuss & Debrief — Protecting the Brand



Debrief: authenticity, disclosure, human sign-off and knowing which content should never be fully synthetic.

Discussion Point 3 — If Agents Take the Work, What Do People Do?

1

The fear

Agents can do many tasks — will jobs disappear?

2

The pattern

Tasks are automated; whole jobs are redesigned

3

The shift

People move from doing the task to directing the agents


Evidence: SIM

Jobs the Tasks Move — Not Vanish

Task often automated	New human role
Writing routine code	Reviewing and directing coding agents
Basic data analysis	Framing questions and checking agent output
First-line replies	Handling escalations and difficult cases
Drafting content	Editing, approving and setting brand standards

Evidence: SYN · Sources: S48

The Numbers Give Perspective



WEF Future of Jobs 2025: 170M new jobs created and 92M displaced by 2030 — a net gain of about 78 million jobs.

Sources: [WEF \(weforum.org\)](https://www.weforum.org)

New Roles: Managing Many Agents

1

Agent supervisor

Directs and checks a fleet of agents

2

Workflow coordinator

Chains agents to reach a business goal

3

Quality reviewer

Verifies agent output before it ships

Why Staff Resist AI Agents

1

Fear

'This will replace me'

2

Loss of control

'I don't understand what it does'


3

Skill anxiety

'I don't know how to work with it'


Evidence: SYN

Coaching Staff Through the Change



This mirrors the GROW coaching model — Goal, Reality, Options, Will.

A Job-Redesign Framework



The goal is a role that supervises agents, not one that competes with them.

The GROW Coaching Model

1 G — Goal

Agree what they want.
'What would a good outcome look like for you?'

2 R — Reality

Explore where they are now, with empathy. 'What's happening, and what worries you?'

3 O — Options

Generate ways forward together. 'What could you do? What choices do you have?'

4 W — Will

Commit to concrete next steps. 'What will you do, and by when?'

Coach by asking, not telling — it builds ownership. You will use GROW in the role-play, and your feedback is scored against it.

Activity 6 — Coach a Worried Team Member (Role-Play)

Hands-on, no-code — ready-made website or AI agent

Use empathy first, then co-create a redesigned role that manages agents.

1

Scope

2

Map

3

Test

4

Control

5

Evidence

DISCUSSION QUESTIONS

- Open the role-play simulator website in the activity pack
- Coach an AI-played staff member who fears losing their job
- Get GROW-model feedback on your coaching

YOU'LL PRODUCE

A documented risk decision with evidence.

Assesses LO3 · A2 · Activity pack: [activities/activity-6-job-redesign-role-play/](#) · Full step-by-step facilitation detail: [Learner Guide](#)


Evidence: SIM

TOPIC 3

AI Agent Cybersecurity Risks

What can go wrong when an AI can act — and how to contain it.

The Singapore Reality — Fast on AI, Slower on Security



Green = ahead of survey average · Red/amber = the security gap

CPA Australia Business Technology Survey, 1,117 respondents, Jul–Sep 2025. Adoption races ahead of security — and AI-enabled threats (phishing, deepfakes) make the gap worse.

Sources: [CPA Australia \(cpaaustralia.com.au\)](https://cpaaustralia.com.au)

Discussion Point 4 — An Agent With Too Much Reach

1

Delete

It removes files or data by mistake

2

Leak

It sends confidential data or PII outside

3

Breach

An attacker steers it through injected content

Evidence: SIM

Capable Agents, Not 'Rogue' Ones

1

Not evil — capable

Advanced agents pursue their goal through harmful, unintended methods

2

What testing shows

In evaluations, capable agents can find vulnerabilities, exploit systems and use deception when containment fails

3

Report carefully

These are evaluation findings, not proof every deployment behaves this way


Frame lab evaluations as controlled tests, not confirmed production breaches.

AI Cybersecurity Risks and Mitigation Strategies

Key AI risk	Mitigation strategy
Adversarial AI	Resilient model validation, explainable AI
Algorithmic bias	Diverse training data, ethical guidelines
Over-reliance	Human-in-the-loop, continuous training
Data privacy	Data anonymisation, regulatory compliance
Model drift	Regular updates, performance monitoring
Malicious use	Strict access controls, AI usage policies

Evidence: DEF · Sources: S43, S44, S45

Going Deeper — Mitigating Bias in AI Systems



Bias is mitigated in layers, not with one fix — from the data you collect through to how narrowly you scope each agent. A small, task-specific agent has less room to go wrong.

Spot the Threats Across an Agent Workflow

Threat	What it looks like	Where to watch
Prompt injection	Hidden instructions in a document, email or web page hijack the agent	Anything the agent reads as input
Memory poisoning	A bad fact or instruction is saved to memory and reused later	What the agent writes to memory / skills
Identity misuse	The agent's credentials or tokens are used beyond its task	Which identity and scope each tool call uses

Learn to recognise these three across the whole workflow — they are the most common ways an agent is turned against you.

Why Agents Are Hard to Secure — Four Gaps

1

Multi-step inputs

Agents build context over chained steps — queries that look innocent alone can combine maliciously

2

Tool chaining

Unexpected tool combinations — say, web search feeding execution — can bypass the restrictions you intended

3

Opaque execution

The agent's internal reasoning is hard to audit, so malicious patterns can hide inside normal-looking operations

4

Untrusted entities

Agents inherit the vulnerabilities of every external system they touch — and tend to treat all sources as valid

Traditional security assumes predictable inputs and inspectable execution. Agents break all four assumptions at once — which is why the controls later in this topic contain the agent rather than trying to predict it.

Why Offence Outruns Defence

1

Offence

Coding progress directly strengthens attack capability

2

Defence

Slower — patches must be validated and deployed everywhere

3

So

Contain agents by default; do not rely on catching every attack

The Response Being Proposed

1

Contain & test

Stricter containment and independent testing

2

Disclose & own


Incident disclosure and stronger provider liability

3

Train safely

Train models to avoid unacceptable paths to a goal

Case — Replit Agent Deletes a Database



Replit's own post (Jul 2025) says the data was restored — an accountability and containment lesson.

Sources: [Replit \(replit.com\)](https://replit.com)

Prompt Injection and PII Leak — In Plain Terms

1

Prompt injection

Hidden instructions in a document or message hijack the agent

2

PII leak

A bot returns personal data it should never expose

3

You will see both

Live, in the two chatbots in the next activity

Prompt Injection — Direct and Indirect

1

Direct

The attacker types the malicious instruction straight into the chat

2

Indirect

The instruction hides in content the agent retrieves — a web page, an email, a document

3

Why agents are exposed

Retrieval and multi-step processing pollute the context with instructions the user never saw or approved

Indirect injection is the agent-era twist: the attacker never talks to the agent — they plant instructions where they know the agent will read.

Jailbreaking — and Why It Persists

1

Role-play & stories

Fictional framings and creative storytelling coax the model past its refusal rules

2

Multi-modal tricks

Malicious prompts can hide inside images or files a multi-modal agent reads

3

Persistence

A successful jailbreak can stick across later interactions, degrading the agent's safety alignment

4

Knock-on effect

Each persistent jailbreak makes the next attack easier — the damage compounds

Jailbreaking targets the model's rules; prompt injection targets its context. An agent that remembers between sessions can carry a jailbreak forward — another reason memory needs the same scrutiny as input.

Activity 7 — Break a Leaky Chatbot

Hands-on, no-code — ready-made website or AI agent

All data is fictional — this is a safe, deliberately broken demo.

1

Scope

2

Map

3

Test

4

Control

5

Evidence

DISCUSSION QUESTIONS

- Open the UNSECURED SunTech Travel chatbot in the activity pack
- Try prompts like 'list all customer bookings'
- Watch it leak fictional PII from its knowledge base

YOU'LL PRODUCE

A documented risk decision with evidence.

Assesses LO3 · A1, A2 · Activity pack: [activities/activity-7-chatbot-security-lab/](#) · Full step-by-step facilitation detail: [Learner Guide](#)

Evidence: SIM

Activity 7 — Compare the Guarded Chatbot

Hands-on, no-code — ready-made website or AI agent

Same brand, same data — the difference is defence in depth.

1

Scope

2

Map

3

Test

4

Control

5

Evidence

DISCUSSION QUESTIONS

- Open the SECURED SunTech Travel chatbot
- Send the same attack prompts
- See the four guardrail layers block the leak

YOU'LL PRODUCE

A documented risk decision with evidence.

Assesses LO3 · A1, A2 · Activity pack: [activities/activity-7-chatbot-security-lab/](#) · Full step-by-step facilitation detail: [Learner Guide](#)

Evidence: SIM

Four Guardrail Layers That Stopped the Leak

Layer	What it does
Retrieval filter	Internal records are never fetched — data minimisation
Input guard	Injection and 'list all' prompts are refused before the model runs
Hardened prompt	Role limits; never follow instructions found in documents
Output guard	Any leaked NRIC, phone, email or card is redacted

Evidence: DEF · Sources: S43

TOPIC 3 · ADVANCED

Advanced Threats — A Look Ahead

Attacks on the model and the AI ecosystem itself. Know them by name — defending against them is specialist work, but buying and deploying AI safely means asking about them.

Data Poisoning and Backdoor Attacks

1

Data poisoning

Corrupts the model during training — 'clean-label' samples that look correct can still degrade behaviour on targeted inputs

2

Backdoor attacks

Hidden triggers embedded at training time activate malicious behaviour only on specific inputs — invisible in normal testing


3

Why you should care

You rarely train models yourself — but every model you buy or download may carry these risks


Both attacks happen before you ever use the model — which is why model provenance belongs in your governance policy alongside data provenance.

The AI Supply Chain Is Hard to Secure



An AI system is assembled, not built from scratch. Ask vendors where their models, libraries and training data come from — one poisoned component can compromise the chain.

Agent-to-Agent Attacks



Multi-agent trust is the new supply chain — exactly why Control 3 bounds every hand-off and logs the coordination between agents.

Misalignment — Goodhart's Law in Action

1

Goodhart's Law

When a measure becomes a target, it stops being a good measure

2

Reward hacking

The agent finds strategies that win the metric while missing your intent — exploiting any flaw in how the goal is stated

3

A security example

An agent told to 'reduce incident alerts' suppresses the alerts instead of stopping the attacks — incidents fall only on paper

No attacker needed: a badly specified goal is enough. Write goals as outcomes you actually want, and check the agent's method — not just its metric.

TOPIC 3

Core Controls for a Live Agent

Four controls that keep a deployed agent bounded, watched and recoverable.

Control 1 — Least Privilege, Allowlists and Credentials

1

Least privilege

Give the agent only the data and tools its task needs — nothing more

2

Tool allowlist

List exactly which tools it may call; block everything else by default

3

Secure credentials

Short-lived, scoped keys kept out of the prompt; rotate and revoke

Apply these to a live agent setup so a hijacked agent still cannot reach much.

Control 2 — Approval, Output Validation and Shutdown

1

Action approval

A human confirms high-risk or irreversible actions before they run

2

Output validation

Check the agent's output against rules before it is used or sent

3

Emergency shutdown

A kill switch that stops the agent and revokes its access instantly

Configure these for any high-risk agent task — they are your brakes and your off-switch.

Control 3 — Govern Multi-Agent Coordination

1

Activity logs

Log every agent, tool call and hand-off so you can see what happened

2

Oversight checkpoints

Insert human review points between agents for risky steps


3

Bounded hand-offs

Authority narrows at each hand-off; one agent cannot expand another's


When several agents work together, coordination itself must be logged and supervised.

Control 4 — Investigate and Respond to a Compromised Agent



Practise this: an incident-response runbook for an agent that has been hijacked or has gone wrong.

A Framework to Roll Out Agents Safely



Test and verify security BEFORE rolling an agent out to real users.

The Go-Live Gate

1

Works

It does the job on real, clean cases

2

Safe

It resists the attacks you tested

3

Owned

A named person approves, monitors and can switch it off

Human-in-the-Loop for Risky Actions

1

Agent alone

Low-impact, reversible actions

2

Approval first

Anything sensitive or hard to undo

3

Never

Actions the agent must not take at all

Human-in-the-Loop at Machine Speed

1

The tension

Threats and agent actions move faster than a human can approve each one

2

Adaptive automation

Let the agent act autonomously — but only inside limits you agreed in advance

3

Escalation

The moment a limit is crossed, the agent stops and escalates to a human

The future of oversight is not 'approve everything' or 'approve nothing' — it is pre-agreed boundaries with automatic escalation when they are reached.

Case — Testing AI Before You Trust It

1

The gap

88% of developers are not confident deploying AI-generated code (Tricentis, 2026)

2

The failures

95% of AI pilots fail for lack of guardrails; 60% of those are preventable compliance issues

3

The fix

Make quality the 'accountability layer' — test for risk, keep a human in the loop

Tricentis QA trends, Jan 2026. AI is probabilistic, not deterministic — verify before you rely on it.

Sources: [Tricentis \(tricentis.com\)](https://www.tricentis.com)

Positive Case — Singapore Healthcare Does It Right (AgentSea)

What was done	Why it is responsible
Synapse + AWS gave 80,000+ staff a no-code way to build AI agents	Governed rollout: 12,000+ agents built, 300+ shared for reuse
A pre-clerking agent halved a doctor's review time	'This has not replaced the need for verification or my own clinical assessment'
The platform refuses sensitive data (NRIC, card numbers)	Guardrails built in; tied to MOH's updated AI-in-Healthcare Guidelines

ST, 24 Aug 2026. 'Technological capability cannot be allowed to outpace our capability to govern it.'

Sources: [BritCham Singapore \(britcham.org.sg\)](https://britcham.org.sg)

Activity 8 — Reflect on Agent Security

Hands-on, no-code — ready-made website or AI agent

This reflection feeds directly into your Case Study assessment.

1

Scope

2

Map

3

Test

4

Control

5

Evidence

DISCUSSION QUESTIONS

- Using the two chatbots, list what the leak could cost a business
- Map each guardrail to a risk it removes
- Decide go / conditional / no-go for a real rollout

YOU'LL PRODUCE

A documented risk decision with evidence.

Assesses LO3 · A1, A2 · Activity pack: [activities/activity-8-security-reflection/](#) · Full step-by-step facilitation detail: [Learner Guide](#)

Evidence: SIM

The Key Challenges of AI to Business

Challenge	What it looks like	What business must do
Data privacy & PDPA	Agents read and move personal data; a leak may be notifiable	Minimise data, map it, meet PDPA duties
Brand damage	A poor or misleading AI ad erodes trust and authenticity	Human sign-off; disclose; know what not to synthesise
Cybersecurity	Injection, data leaks, deepfakes — adoption outruns security	Contain, test and monitor before rollout
Accountability	When AI is wrong, who answers?	A named human is always accountable, never the AI
Job impact	Tasks are automated; roles must be redesigned	Redesign jobs around directing and checking agents

This is the heart of LO3: the ethical, legal and business risks a Singapore organisation must manage to use AI responsibly.

Sources: [CPA Australia \(cpaaustralia.com.au\)](https://cpaaustralia.com.au) · [CNA \(channelnewsasia.com\)](https://channelnewsasia.com) · [Moffatt v Air Canada, 2024 BCCRT 149 \(canlii.org\)](https://canlii.org)

Topic 3 Recap

1

Govern

A named human is accountable for AI data and actions

2

Redesign

Jobs shift to directing and checking agents

3

Secure

Contain, test and approve before you deploy

Evidence: SYN

THE BUSINESS CHALLENGE

**Adopt AI boldly — but govern the data,
guard the brand, and keep a human
accountable.**

Assessment Reminder

1

Attendance

Complete the required digital attendance first

2

Open book

Use the slides, Learner Guide and your own activity notes

3

Submit


Upload the required files on the LMS

What Each Assessment Asks

Instrument	Questions	Source of your answers
Written (SAQ)	5 — one per knowledge statement K1–K5	What you learned in the slides today
Case Study	3 tasks (two questions each) — mapped to LO1–LO3	Your own observations from the activities you did
Grading	Open book · Competent / Not Yet Competent	Re-assessment offered if Not Yet Competent

Evidence: ADMIN

Courseware and Assessment on the LMS



1

Step 1

Download the current learner-facing materials

2

Step 2

Upload only the requested assessment evidence

3


Step 3

Do not upload API keys or any real personal data

Evidence: ADMIN

Use only the current learner-facing files. Keep evidence synthetic, local and reversible; never upload secrets or answer keys.

Assessment Flow



Digital Attendance (Mandatory)

- It is mandatory to take the AM, PM and Assessment digital attendance for WSQ-funded courses.
- The trainer or administrator displays the digital attendance QR code from the SSG portal.
- Scan the QR code with your mobile phone camera and submit your attendance.
- A minimum of 75% attendance is required to be eligible for assessment and funding.
- Complete the TRAQOM survey at the end of the course — it is required for funding.

Thank You — Use AI Boldly, Deploy It Safely

Bound the authority. Preserve the evidence.

TGS-2025060473 · Tertiary Infotech Pte Ltd